



IPTC  
Information Processing and  
Telecommunications Center

# IPTC Seminar Introduction to Deep Learning and Keras

July 15, 16 and 17, 2019  
ETSIT UPM, Madrid



Information Processing and  
Telecommunications Center



## Introduction to Deep Learning & Keras Day 3

**iptc**

Information Processing and Telecommunications Center

Prof. Luis A. Hernández Gómez  
[luisalfonso.hernandez@upm.es](mailto:luisalfonso.hernandez@upm.es)

E.T.S. Ingenieros de Telecomunicación  
Universidad Politécnica de Madrid



Information Processing and  
Telecommunications Center

Day 1 (July 15) : Intro to Deep Learning & Keras.  
Feed Forward models

Day 2 (July 16) : Backpropagation. Convolutional Networks  
Transfer Learning. Data Augmentation.

Day 3 (July 17) : Recurrent Networks. Advanced Architectures.  
GANs. Applications. Final discussion on AI

IPTC Seminar

Introduction to Deep Learning and Keras

July 15, 16 and 17, 2019  
ETSIT UPM, Madrid

# IPTC Summer Seminar MATERIALS

Please go to our MSTC GitHub:

[https://github.com/MasterMSTC/IPTC\\_DeepLearning](https://github.com/MasterMSTC/IPTC_DeepLearning)

The screenshot shows the GitHub repository page for 'MasterMSTC / IPTC\_DeepLearning'. The repository has 26 commits, 1 branch, 0 releases, and 1 contributor. It contains notebooks, presentations, and README files. The latest commit was made 11 minutes ago.

Materials for IPTC Summer Seminar: Practical Introduction to Deep Learning & Keras

Manage topics

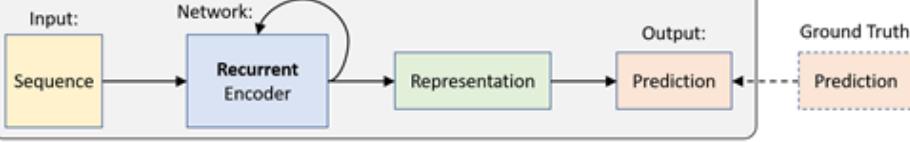
26 commits | 1 branch | 0 releases | 1 contributor

Branch: master ▾ | New pull request | Create new file | Upload files | Find File | Clone or download ▾

File	Commit Message	Time
MasterMSTC_Delete IPTC_GoogleColab_Intro.ipynb	Latest commit f7a7cf3 11 minutes ago	
Notebooks	Delete readme.txt	11 minutes ago
PRESENTATIONS	Create readme.txt	8 hours ago
README.md	Update README.md	8 hours ago
README.md		

# Deep learning basics

## 3. Recurrent Neural Networks

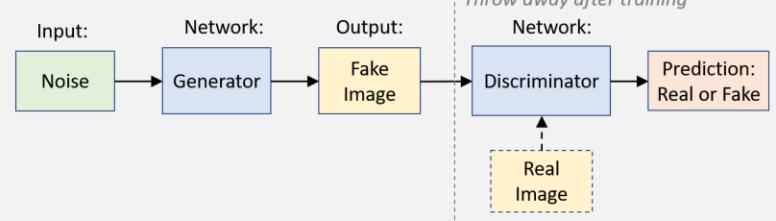


## Unsupervised Learning

### 5. Autoencoder

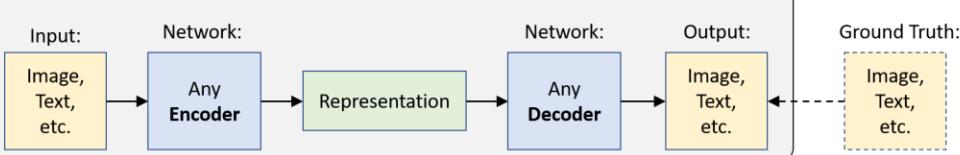


### 6. Generative Adversarial Networks

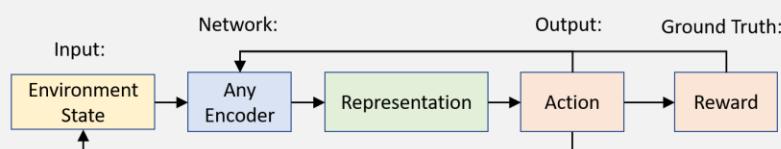


## Reinforcement Learning

### 4. Encoder-Decoder Architectures



### 7. Networks for Learning Actions, Values, and Policies

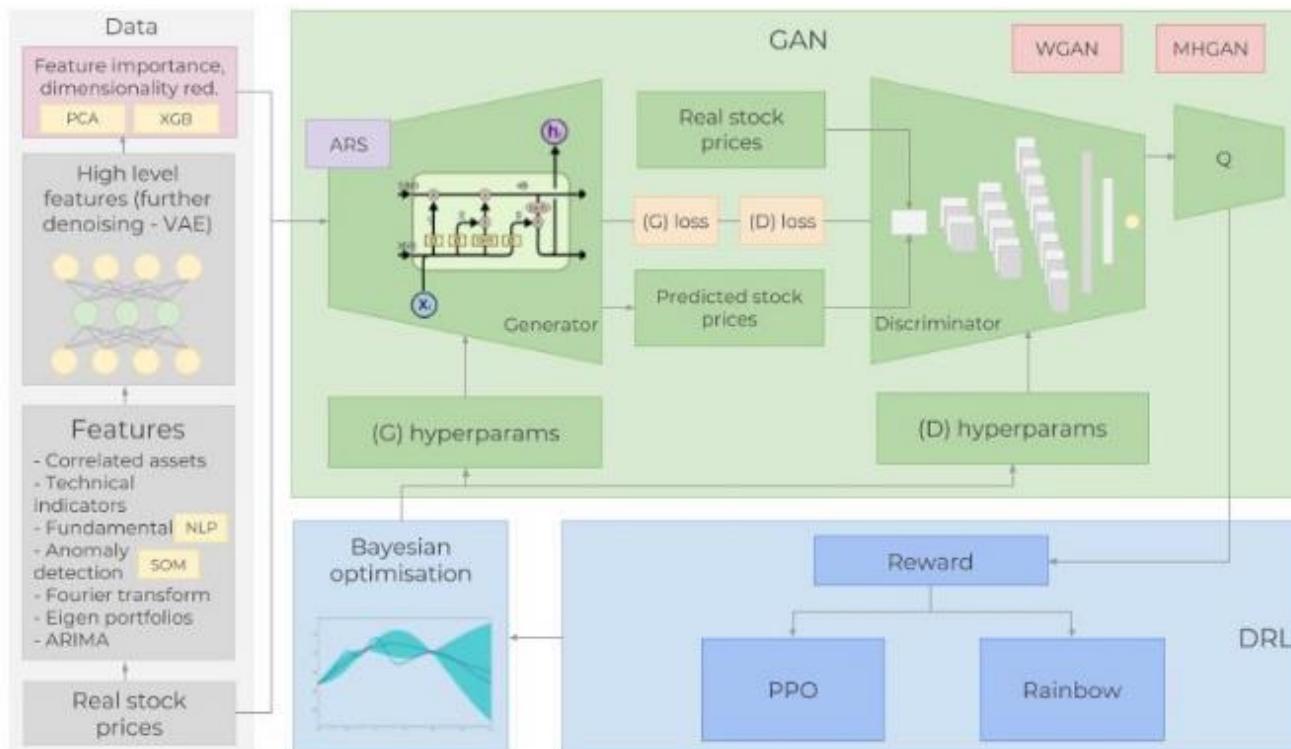


[https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial\\_deep\\_learning\\_basics/deep\\_learning\\_basics.ipynb](https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial_deep_learning_basics/deep_learning_basics.ipynb)

# Using the latest advancements in deep learning to predict stock price movements



Boris B [Follow](#)  
Jan 10 · 34 min read



Overview of the complete architecture.

## 2.8. Extracting high-level features with Stacked Autoencoders

Before we proceed to the autoencoders, we'll explore an alternative activation function.

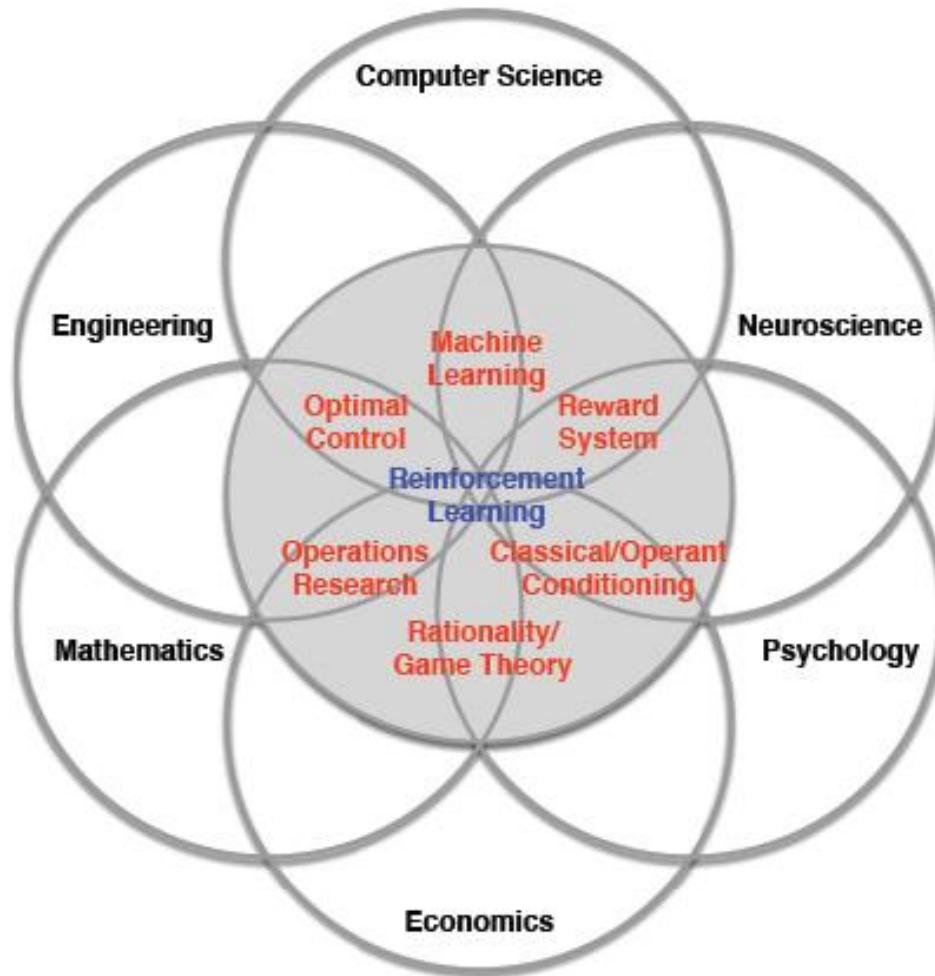
### 2.8.1. Activation function — GELU (Gaussian Error)

GELU — Gaussian Error Linear Unites was recently proposed — [link](#). In the paper the authors show several instances in which neural networks using GELU outperform networks using ReLU as an activation. `gelu` is also used in BERT, the NLP approach we used for news sentiment analysis.

We will use GELU for the autoencoders.

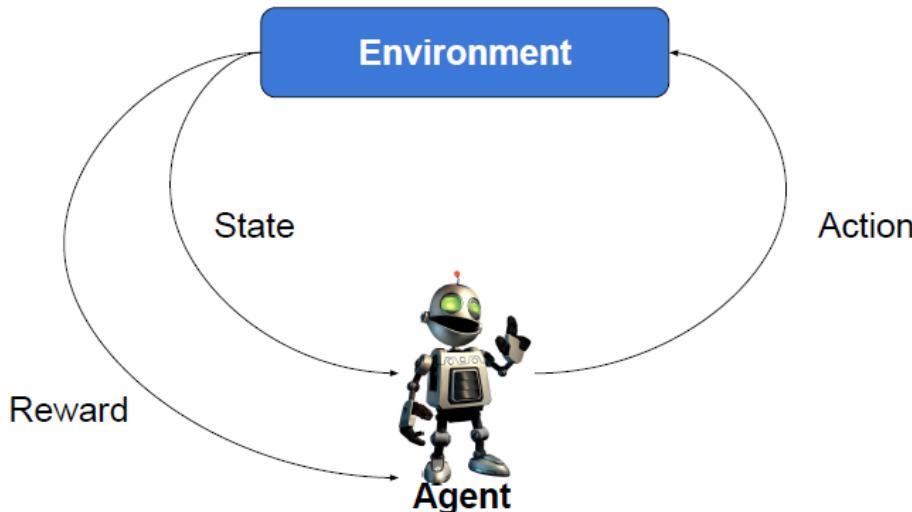
<https://towardsdatascience.com/aifortrading-2edd6fac689d>

# Many Faces of Reinforcement Learning

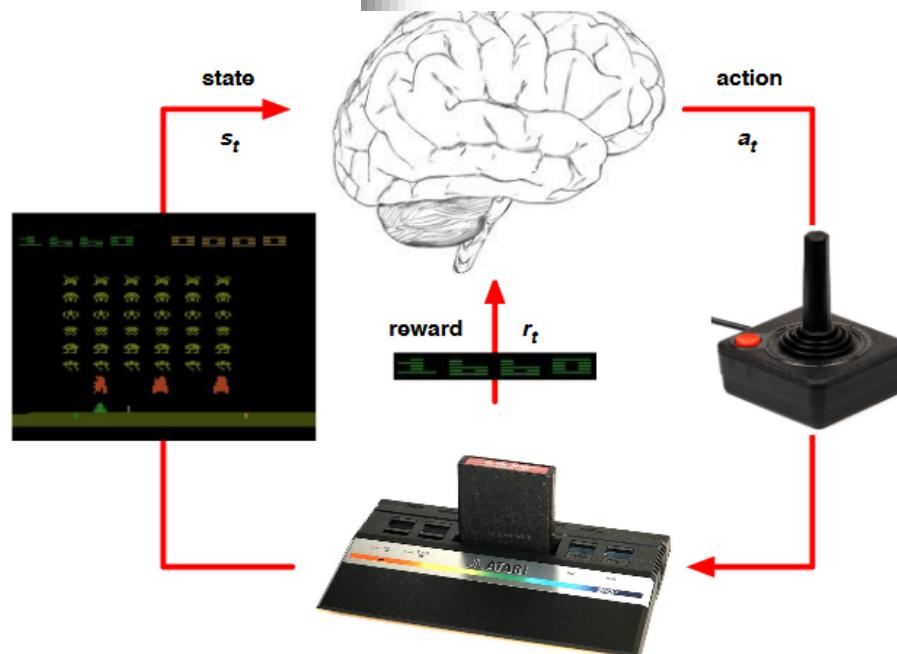
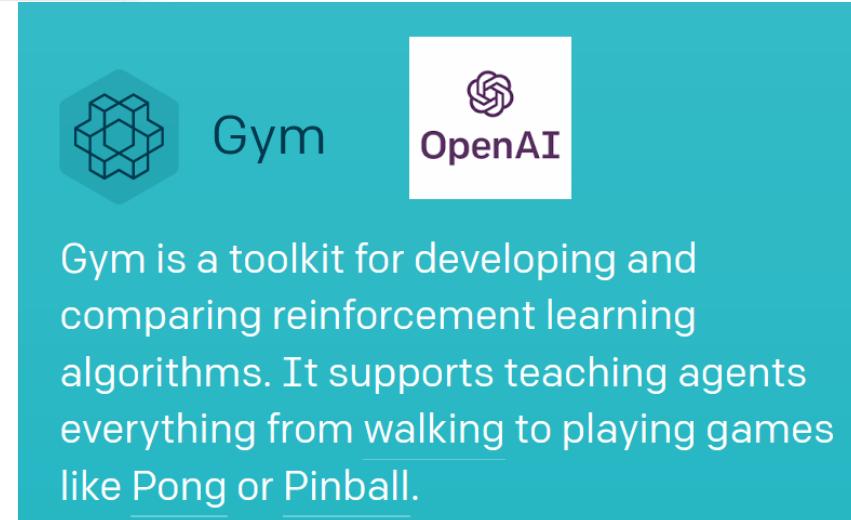


Tutorial: Deep Reinforcement Learning  
David Silver, Google DeepMind

# What's Reinforcement Learning?



- Agent interacts with an environment and learns by maximizing a reward.
- No labels or any other supervision is given.



# Major Components of an RL Agent

- ▶ An RL agent may include one or more of these components:
  - ▶ **Policy**: agent's behaviour function
  - ▶ **Value function**: how good is each state and/or action
  - ▶ **Model**: agent's representation of the environment

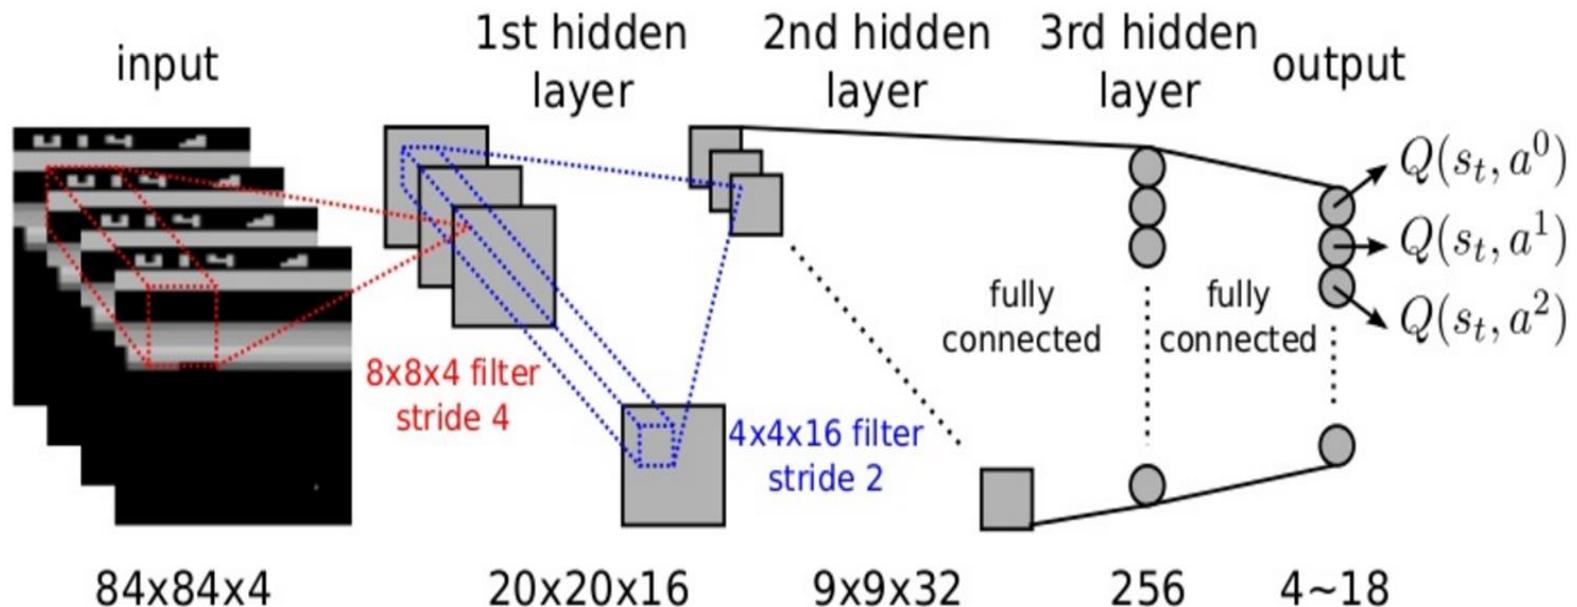
- **Policy  $\pi$**  is a behavior function selecting actions given states

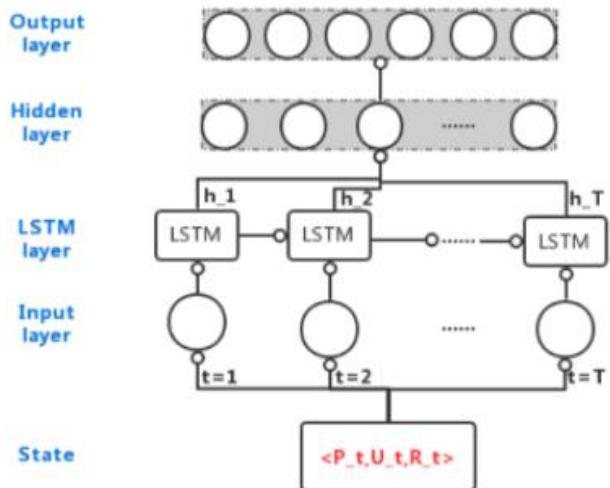
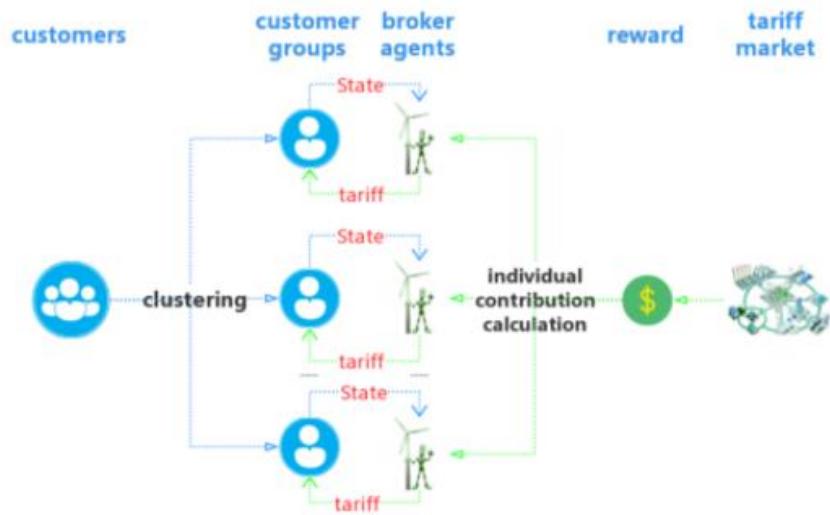
$$a = \pi(s)$$

- **Value Function  $Q^\pi(s, a)$**  is expected total reward from state  $s$  and action  $a$  under policy  $\pi$

$$Q^\pi(s, a) = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s, a]$$

# Example: Deep Q Learning (in Atari)

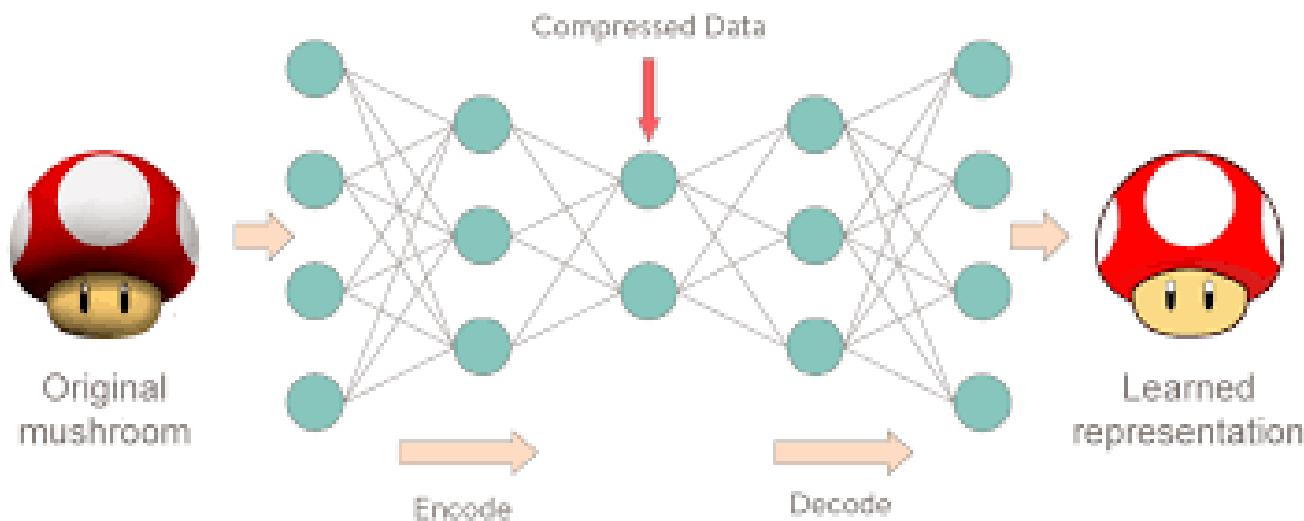




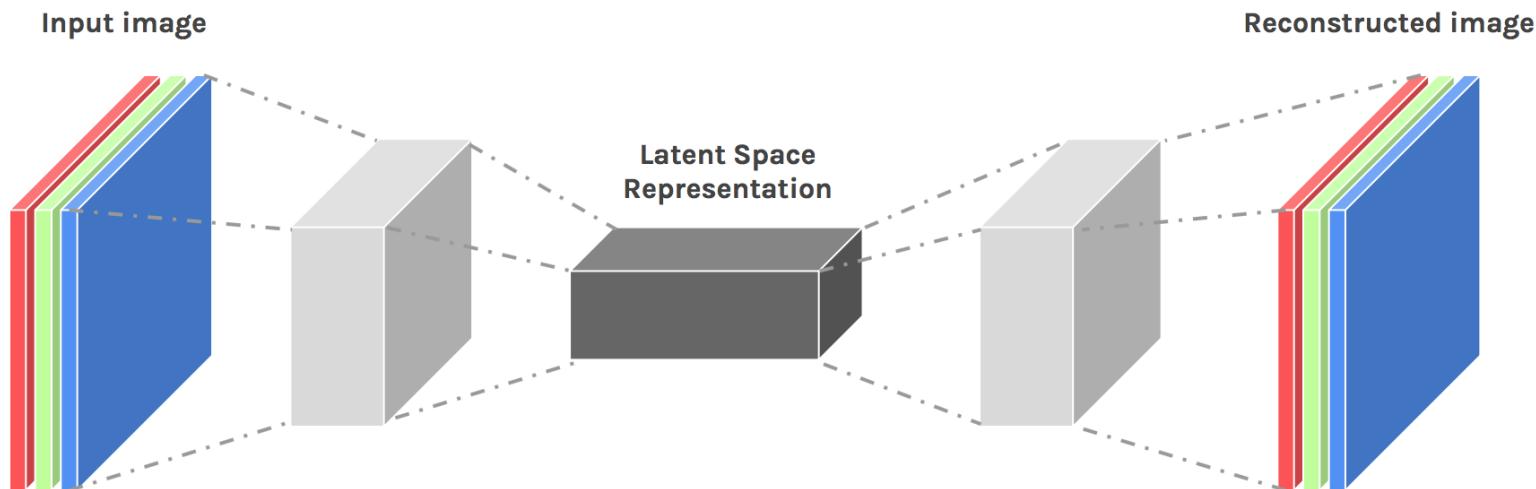
Deep Reinforcement Learning in Retail Pricing Strategy, Source :  
<https://www.ijcai.org/proceedings/2018/0079.pdf>

<https://towardsdatascience.com/deep-learning-vs-deep-reinforcement-learning-algorithms-in-retail-industry-ii-9c17c83ecf2f>

# Autoencoders



You can use different Architectures,  
for example Convolutional Autoencoder



To capture the distribution of  
the data

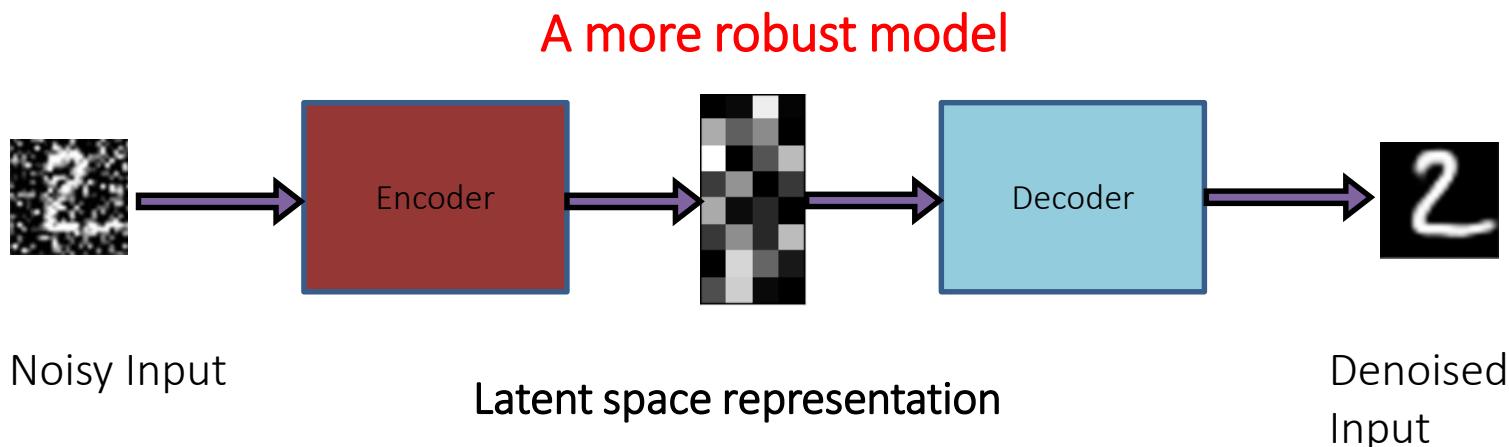
- Bottleneck
- Variational Autoencoders

<https://jaan.io/what-is-variational-autoencoder-vae-tutorial/>

# Denoising Autoencoders

## Intuition:

- We still aim to encode the input and to NOT mimic the identity function.
- We try to undo the effect of *corruption* process stochastically applied to the input.

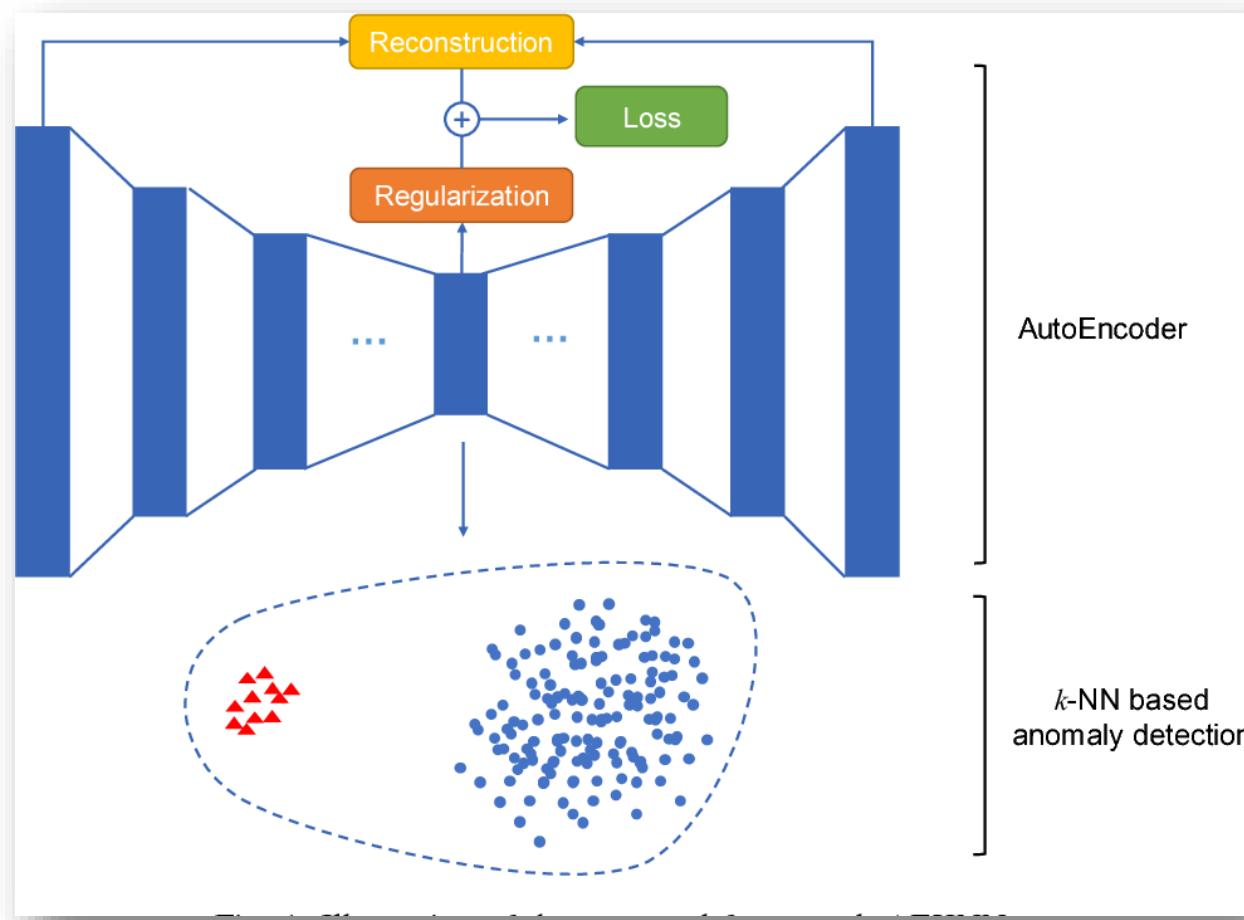


# An Anomaly Detection Framework Based on Autoencoder and Nearest Neighbor

Jia Guo, Guannan Liu, +1 author Junjie Gavin Wu ·

DOI: 10.1109/icsssm.2018.8464983

Published in 15th International Conference on Service Systems... 2018 ·



<https://www.dataversity.net/fraud-detection-using-a-neural-autoencoder/>

# Generations of effects, synthesizers, etc.



Traditionally, a computer turns sound waves directly into numbers. To mix a bass and flute, you would add the numbers at each point along each sound wave. And that would sound like this - a bass playing with a flute:



NSynth also turns sound into numbers. But these numbers don't represent the sound wave itself. They represent a more abstract description of the sound. When you mix a bass and flute using these numbers, it sounds more like a new, hybrid bass-flute sound:



You also get interesting results playing non-instrument sounds through NSynth, like a cow's moo. Since NSynth is only trained on instruments, it tries to recreate the sound using only what it knows about how instruments behave:



# LSTM Model Architecture for Rare Event Time Series Forecasting

by Jason Brownlee on November 2, 2018 in Deep Learning for Time Series



Time series forecasting with LSTMs directly has shown little success.

Welcome to Machine Learning Mastery!



Hi, I'm **Jason Brownlee**, PhD.  
I write tutorials to help  
developers (*like you*) get results  
with machine learning.

[Read More](#)

Uber Engineering

Blog ▾ Research ▾ Locations ▾

[Uber Data](#)

## Engineering Extreme Event Forecasting at Uber with Recurrent Neural Networks

Nikolay Laptev, Slawek Smyl, and Santhosh Shanmugam

June 9, 2017

Sign up for Uber Engineering updates:

Your email address

Subscribe

Completed Trips Test

Number of Trips (scaled)

Time (Days)

Christmas Day

New Year's Day

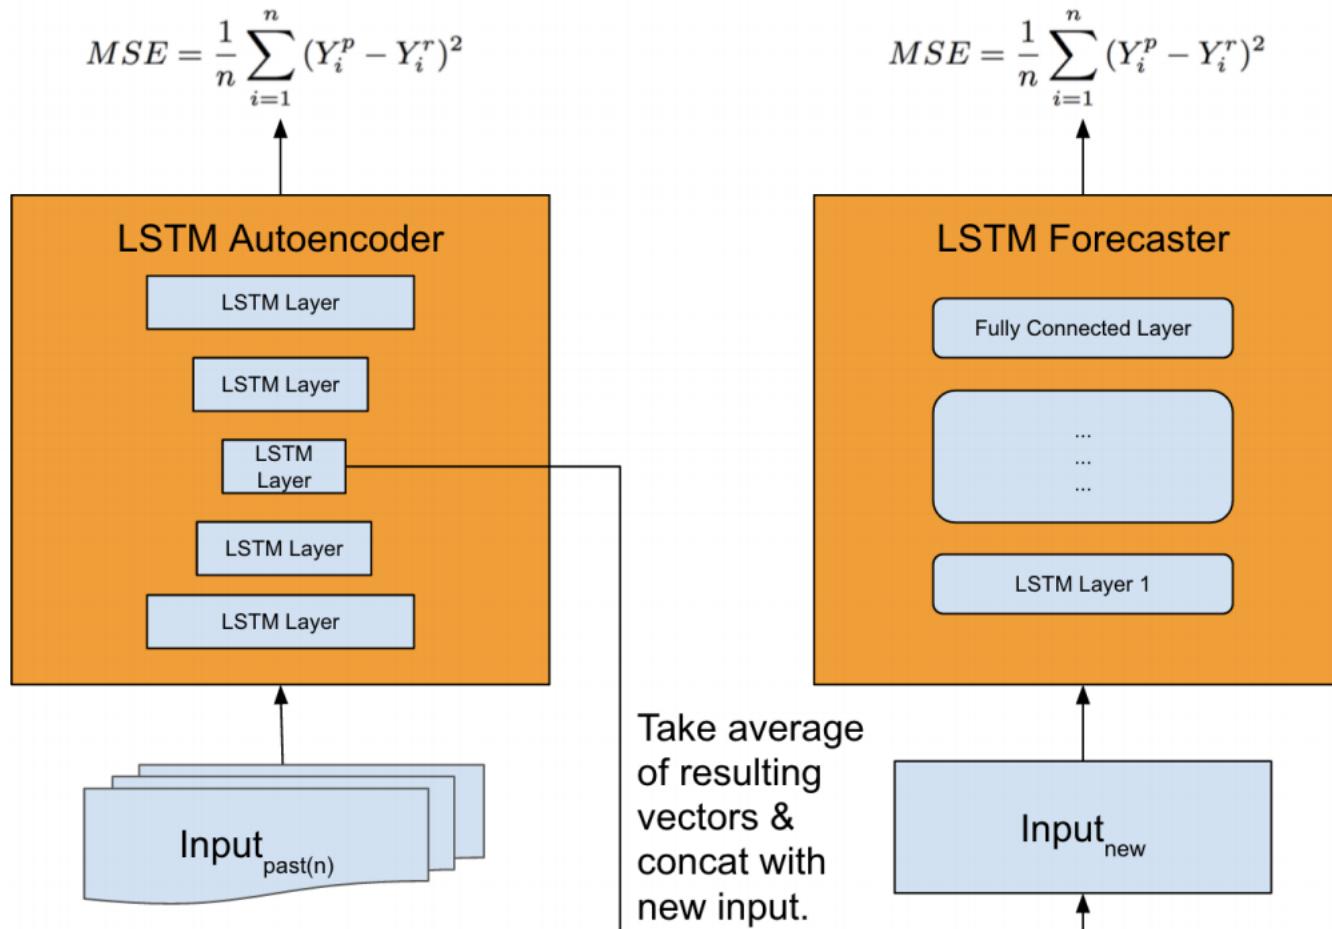
Popular Articles

[Uber's Big Data Platform: 100+ Petabytes with Minute Latency](#)  
October 17, 2018

[Meet Michelangelo: Uber's Machine Learning Platform](#)

# Overview of Feature Extraction Model and Forecast Model

Taken from “Time-series Extreme Event Forecasting with Neural Networks at Uber.”



# Deep Reinforcement Learning

## *Some Refs*

- **Introduction to RL by David Silver**  
[http://www.cs.ucl.ac.uk/staff/d.silver/web/Teaching\\_files/intro\\_RL.pdf](http://www.cs.ucl.ac.uk/staff/d.silver/web/Teaching_files/intro_RL.pdf)  
<https://www.youtube.com/watch?v=2pWv7GOvuf0&list=PL7-jPKtc4r78-wCZcQn5lqyuWhBZ8fOxT>
- **Lex Fridman. MIT 6.S094: Deep Reinforcement Learning**  
<https://www.youtube.com/watch?v=MQ6pP65o7OM>
- **Harini Suresh.MIT 6.S191 Lecture 6: Deep Reinforcement Learning**  
<https://www.youtube.com/watch?v=xWe58WGWWmlk&t=362s>
- <https://www.youtube.com/watch?v=lvoHnicueoE>

# Modelling sequences

- Applications areas are broad:
  - Signals: video, speech, music, sensors,....
  - Time series: financial series, log messages,..
  - Sequence of characters: Natural Language Processing, MIDI music...
  - Sequences of observations, actions, rewards: Reinforcement Learning (robots, bots, autonomous vehicles, dialog systems..)

# Simple Examples: NLP & MUSIC

## Text Prediction/Generation with Keras using LSTM: Long Short Term Memory networks

In this example we will work with the book: Alice's Adventures in Wonderland by Lewis Carroll.

We are going to learn the dependencies between characters and the conditional probabilities of characters in sequences so that we can in turn generate wholly new and original sequences of characters.



Adapted from:

[Text Generation With LSTM Recurrent Neural Networks](#)

By Jason Brownlee

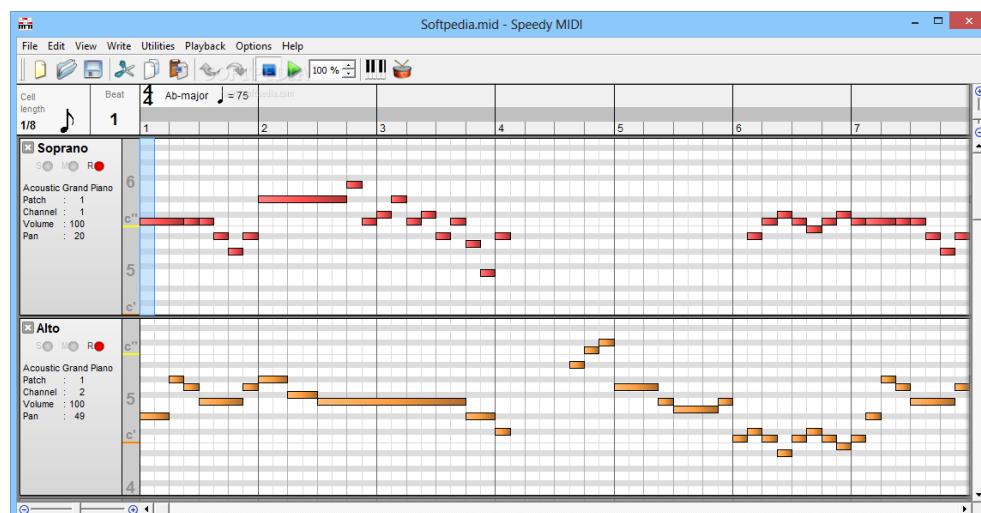
### See example using the IMDB movie review sentiment prediction.

"IMDB dataset", a set of 50,000 highly-polarized reviews from the Internet Movie Database. They are split into 25,000 reviews for training and 25,000 reviews for testing, each set consisting in 50% negative and 50% positive reviews.

Just like the MNIST dataset, the IMDB dataset comes packaged with Keras. It has already been preprocessed: the reviews (sequences of words) have been turned into sequences of integers, where each integer stands for a specific word in a dictionary.



## MIDI Music Generation



<https://blog.floydhub.com/gpt2/>

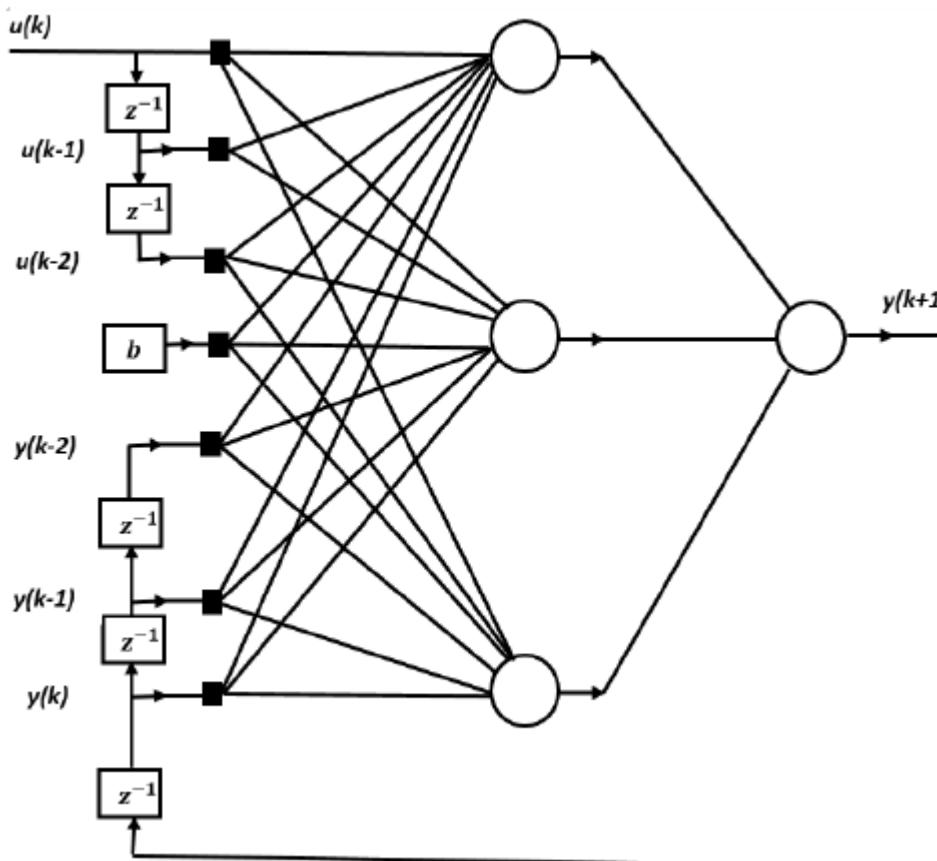
29 APRIL 2019 / DEEP LEARNING

# How to Build OpenAI's GPT-2: "The AI That's Too Dangerous to Release"



# Memoryless models for sequences

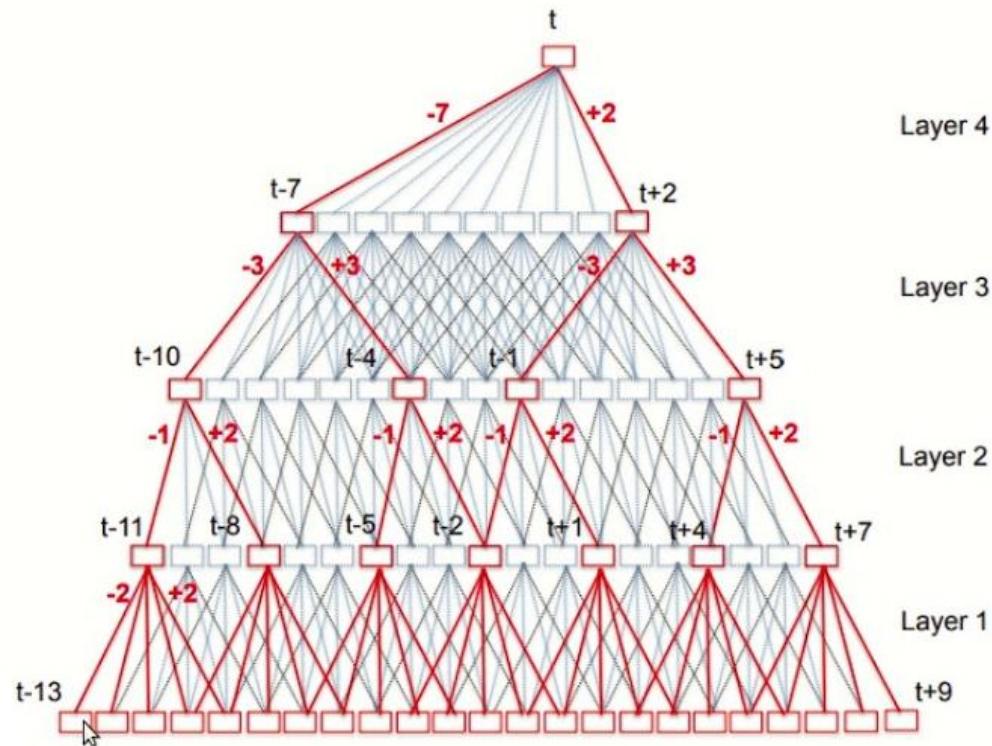
- NARX model: Nonlinear Auto Regressive Model with external inputs.
- Another name: TDNN (time-delay neural network)



# TDNN : Time Delay Neural Network

## 2.1 TDNN basics

1. Input patches consist of fixed number of feature vectors as input to TDNN
2. The TDNN processes them with fixed local temporal context
3. The context width increases as we go to upper layers

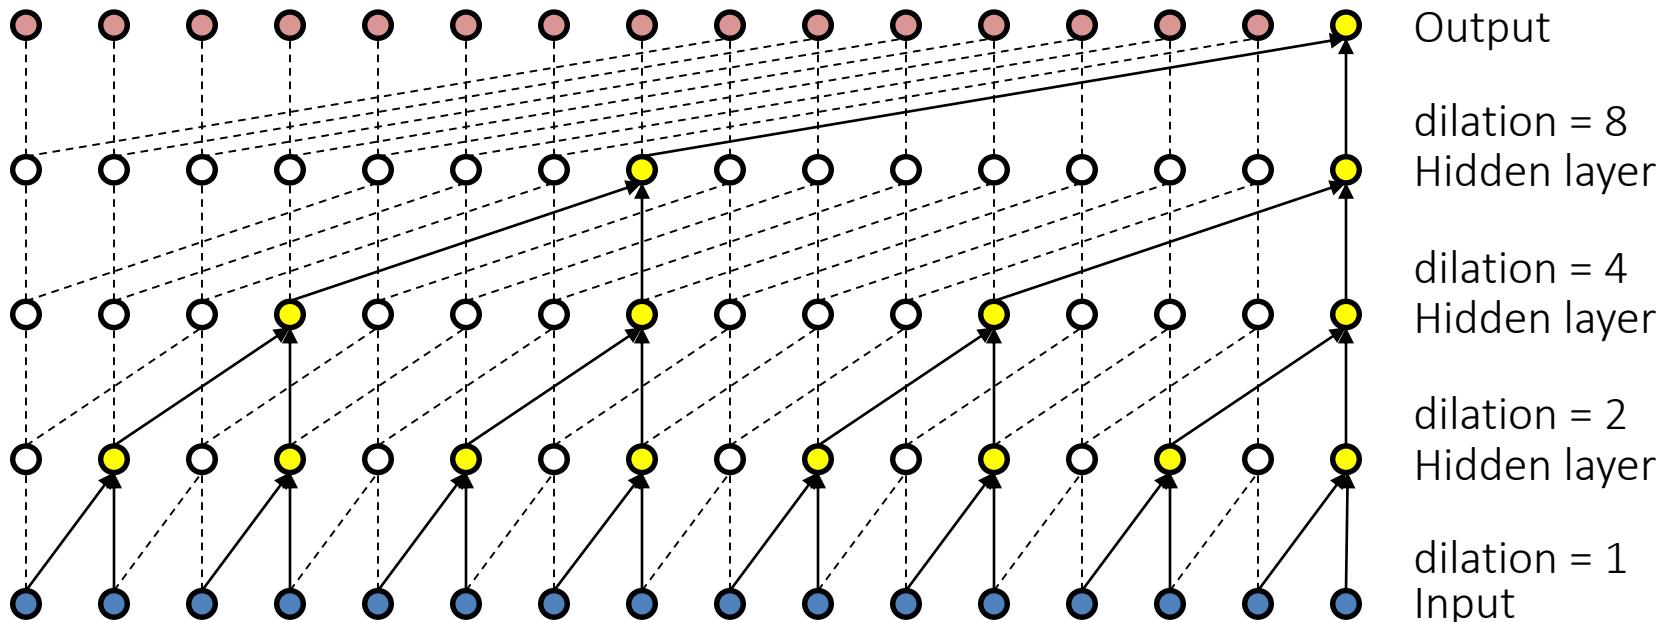


<https://www.youtube.com/watch?v=hwSVs7muapQ>

# WaveNet: A Generative Model for Raw Audio



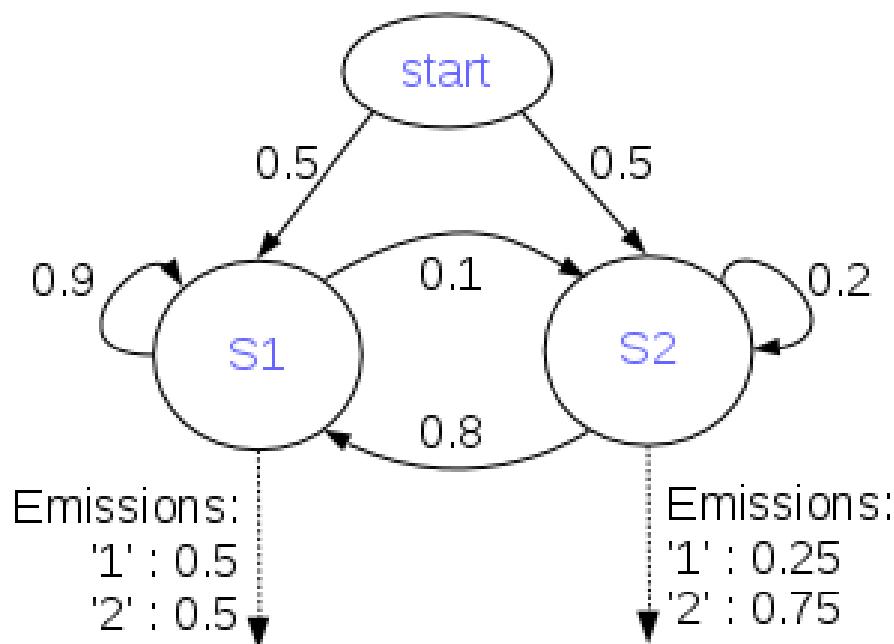
WaveNet models the conditional probability distribution  $p(x_t|x_1, \dots, x_{t-1})$  with a stack of dilated causal convolutions.



<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>

# Dynamic Generative Models: as for example Hidden Markov Models

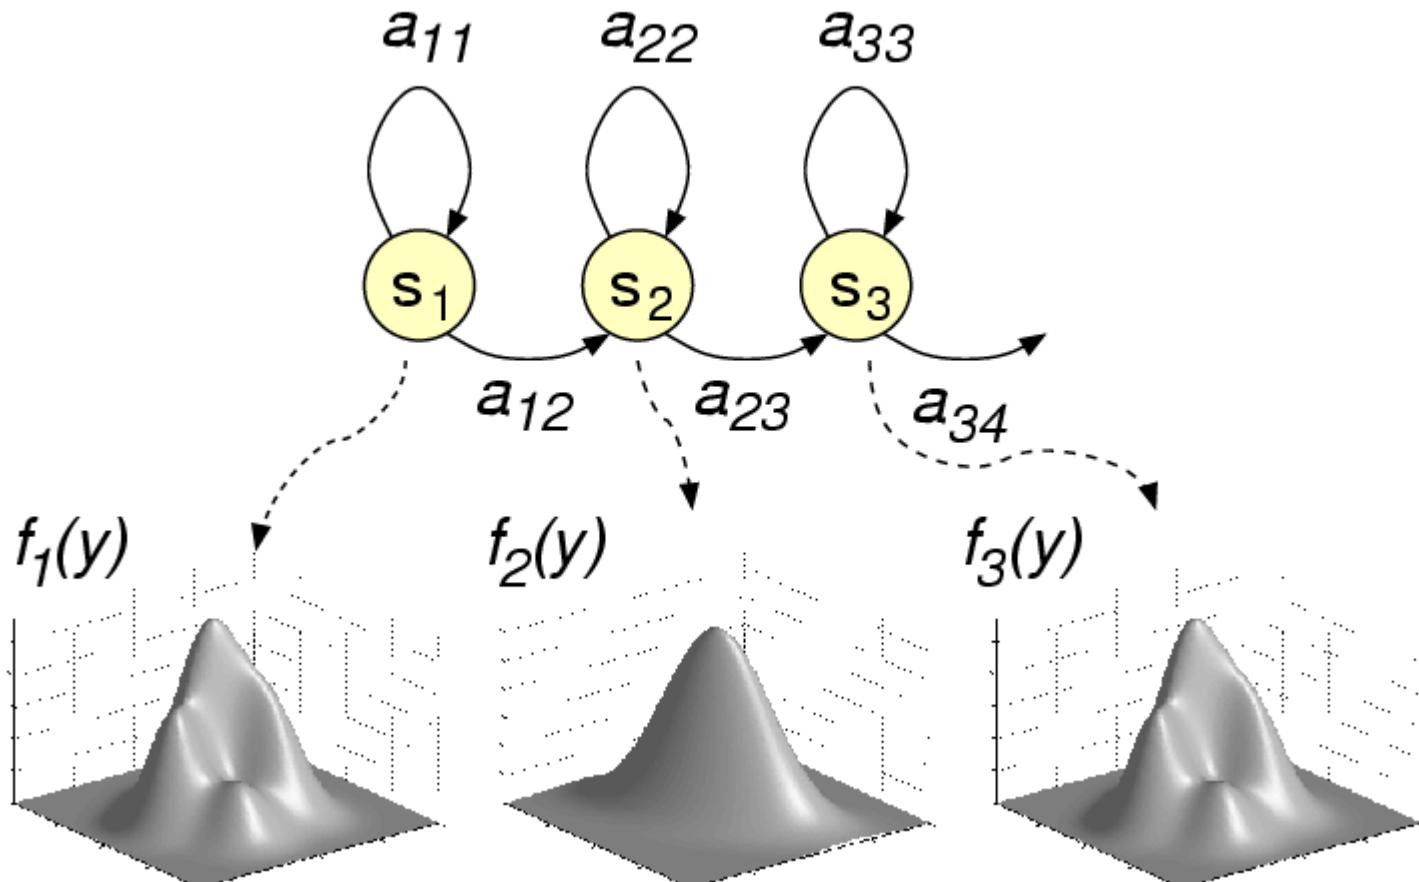
- Hidden Markov Models have a discrete one-of-N hidden state.
- Transitions between states are stochastic and controlled by a transition matrix.
- The outputs produced by a state are stochastic.



From: Geoffrey Hinton  
CSC2535 2013: Advanced Machine Learning

# Traditional Automatic Speech Recognition : HMMs

## Hidden Markov Models



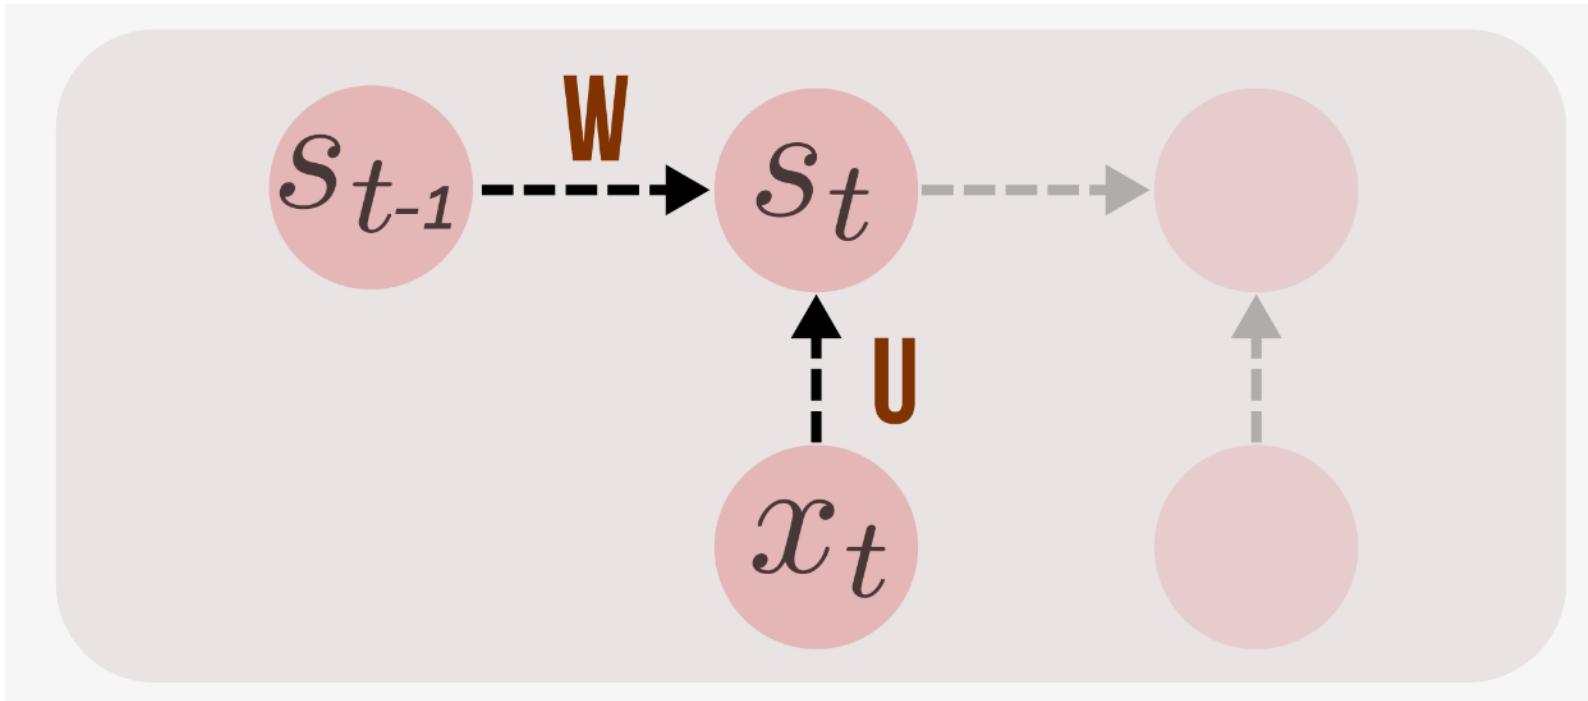
## A fundamental limitation of HMMs

- Consider what happens when a hidden Markov model generates data.
  - At each time step it must select one of its hidden states.
  - So with  $N$  hidden states it can only remember  $\log(N)$  bits about what it generated so far.

From: Geoffrey Hinton  
CSC2535 2013: Advanced Machine Learning

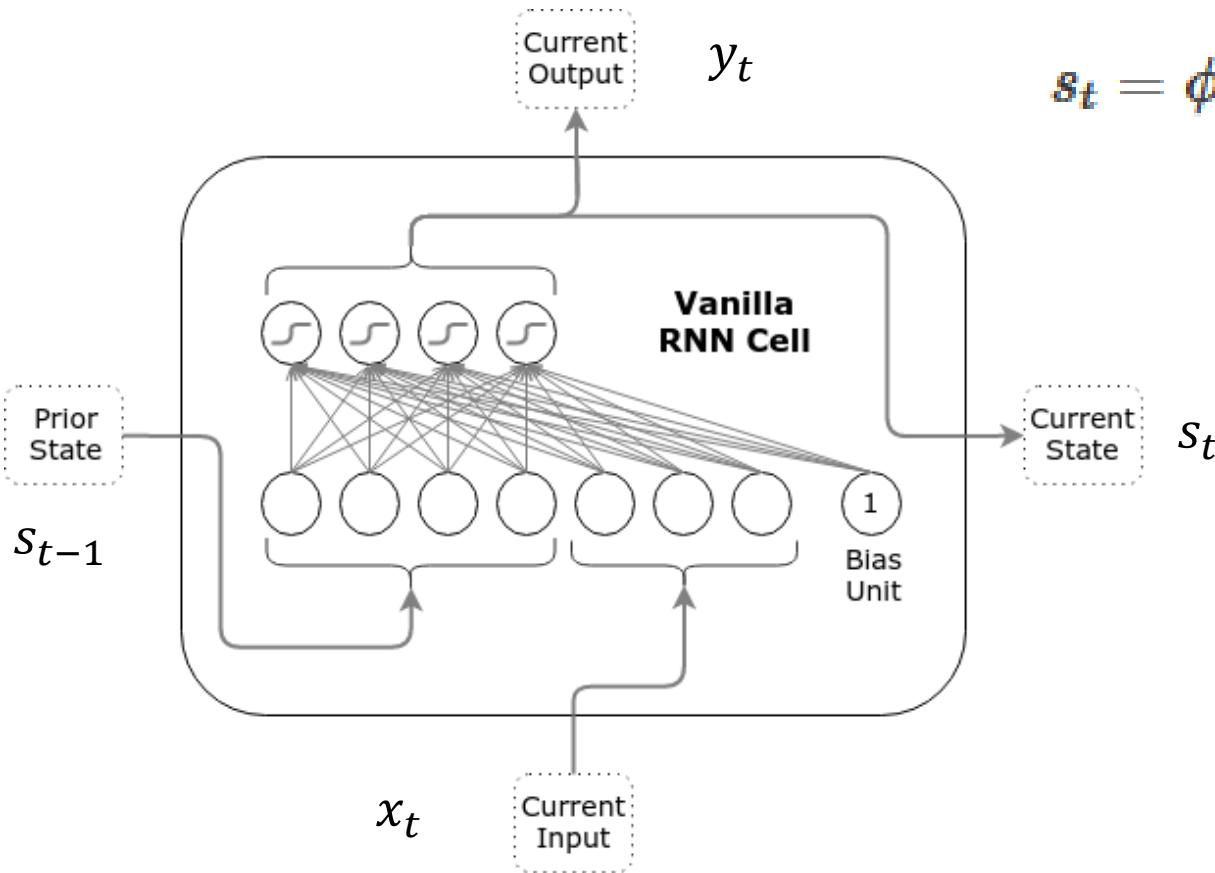
# Recurrent Neural Networks (RNN)

RNNs combine the **input vector** with their **state vector** with a **fixed** (but learned) **function** to produce a new **state vector**.



$$s_t = \tanh(Ux_t + Ws_{t-1}),$$

$$s_t = \phi(Ws_{t-1} + Ux_t + b)$$



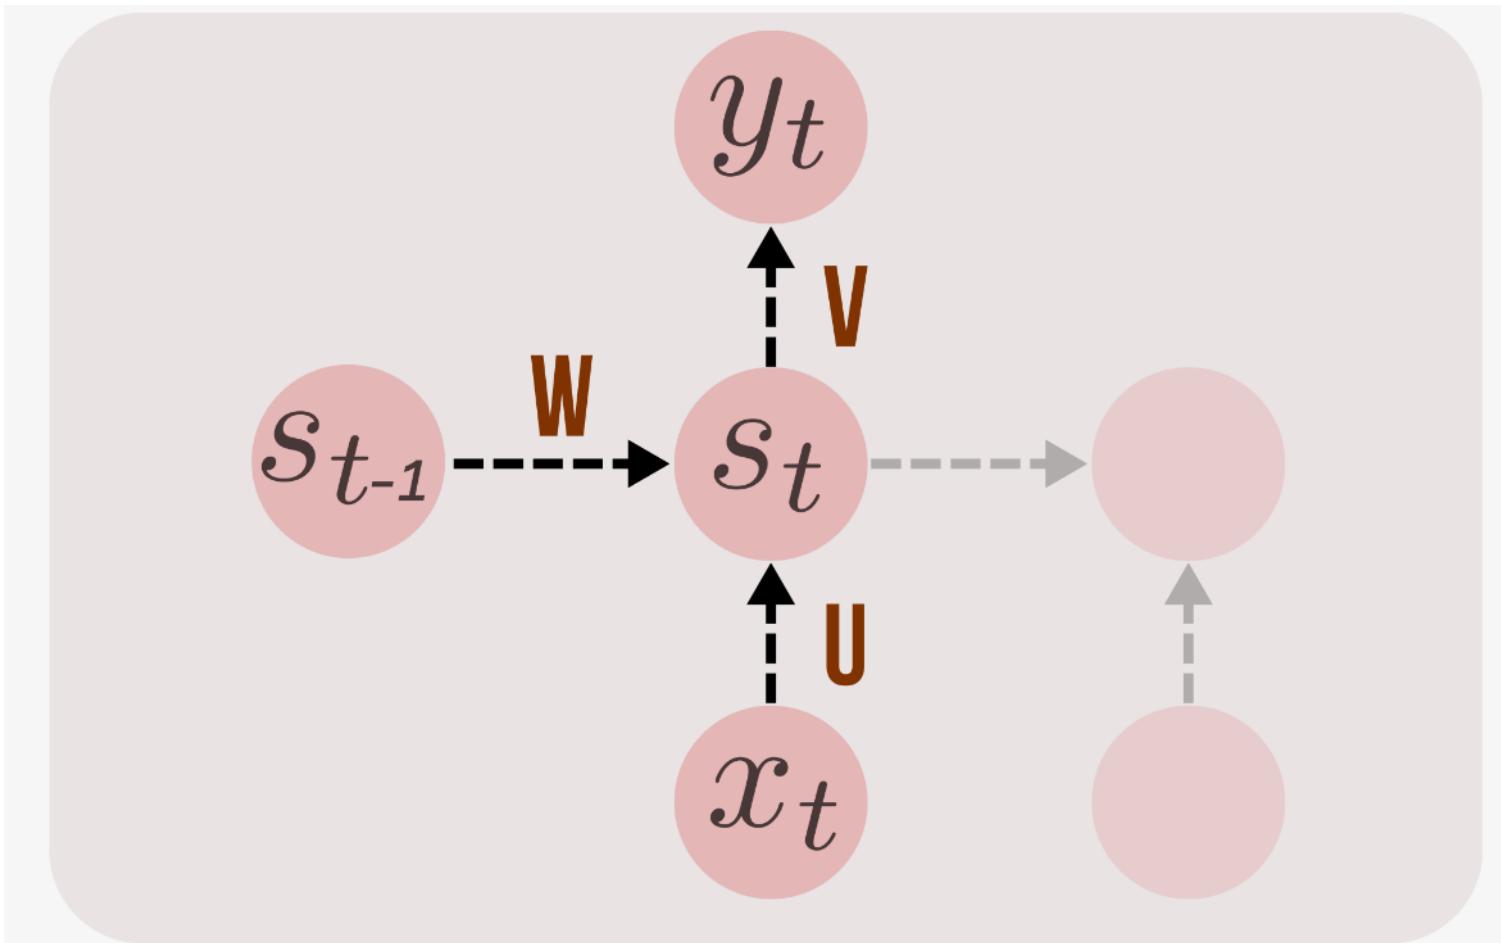
where:

- $\phi$  is the activation function (e.g., sigmoid, tanh, ReLU),
- $s_t \in \mathbb{R}^n$  is the current state (and current output),
- $s_{t-1} \in \mathbb{R}^n$  is the prior state,
- $x_t \in \mathbb{R}^m$  is the current input,
- $W \in \mathbb{R}^{n \times n}$ ,  $U \in \mathbb{R}^{m \times n}$ , and  $b \in \mathbb{R}^n$  are the weights and biases, and
- $n$  and  $m$  are the state and input sizes.

# OUTPUT $y_t$

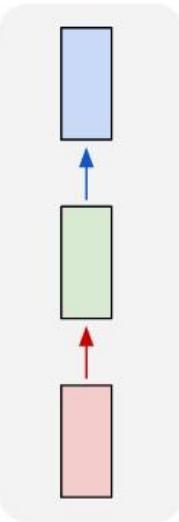
$\text{logits} = \text{tf.matmul}(\text{state}, V) + b$

$\text{predictions} = \text{tf.nn.softmax}(\text{logits})$

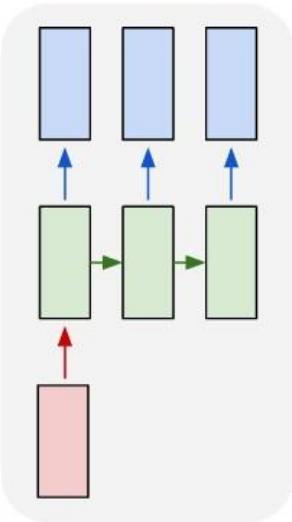


# Recurrent Neural Networks (RNN) concepts

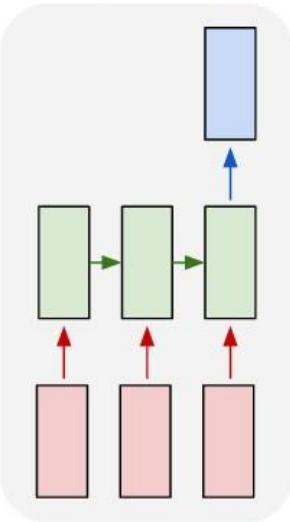
one to one



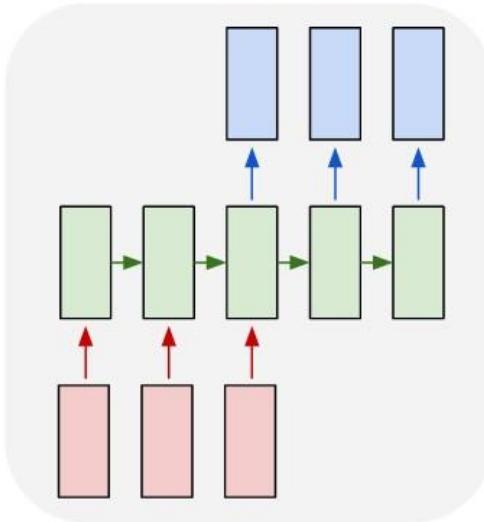
one to many



many to one



many to many



many to many

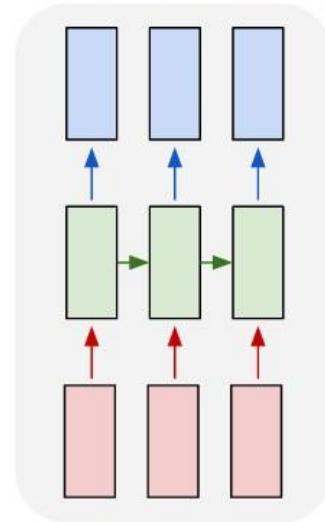


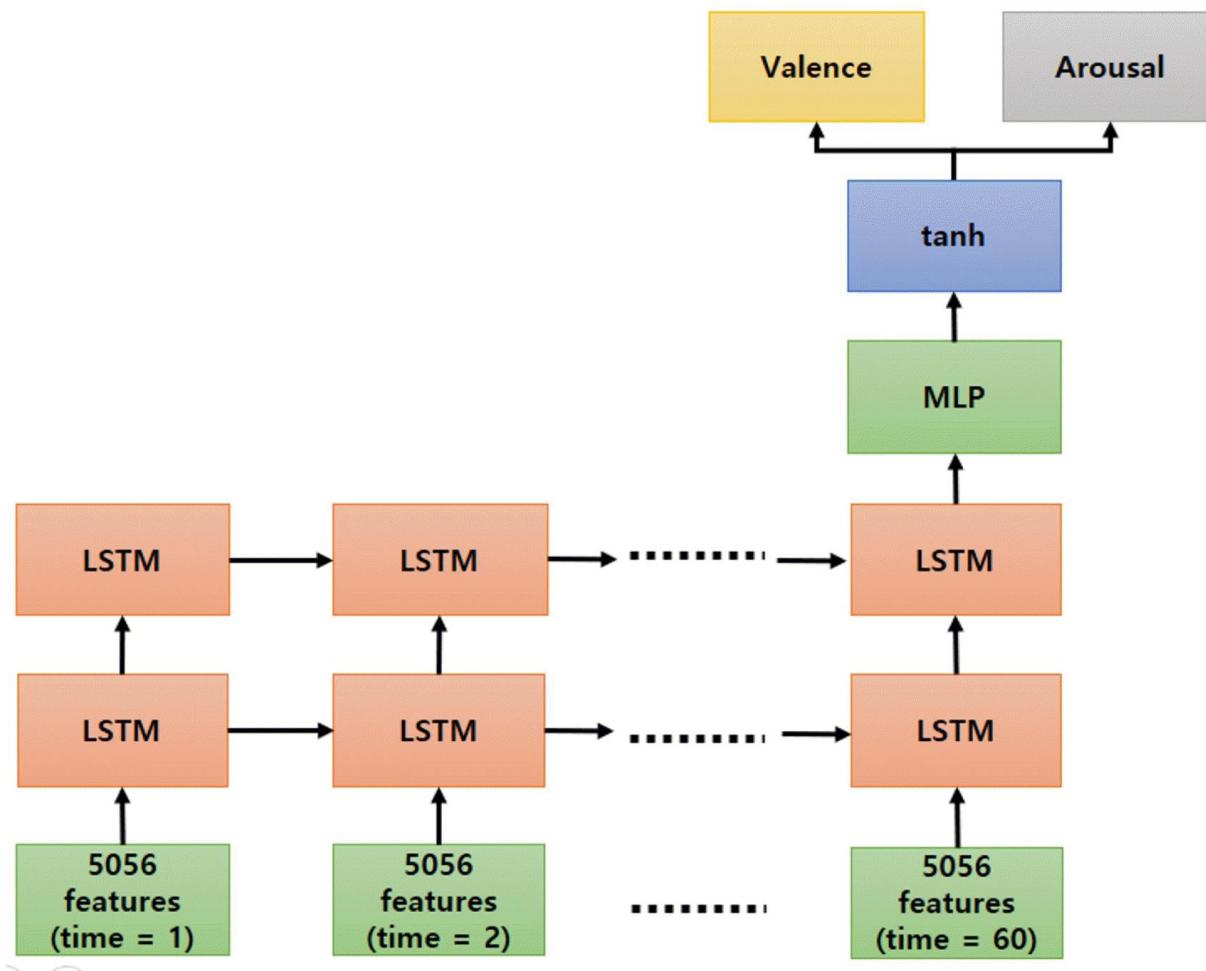
Image  
description  
using text

Emotion  
Recognition  
  
Language  
Model

Machine  
Translation

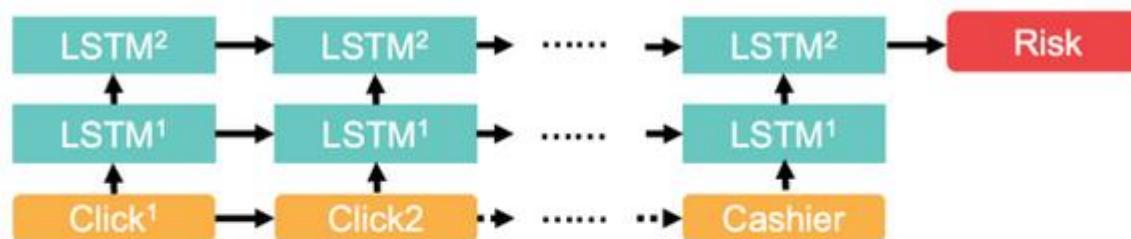
Video frame  
labelling  
  
Phoneme  
Recognition

# Emotion recognition



Deep-Learning Based Model for Emotional Evaluation of Video Clips  
Byoungjun Kim, Joonwhoan Lee

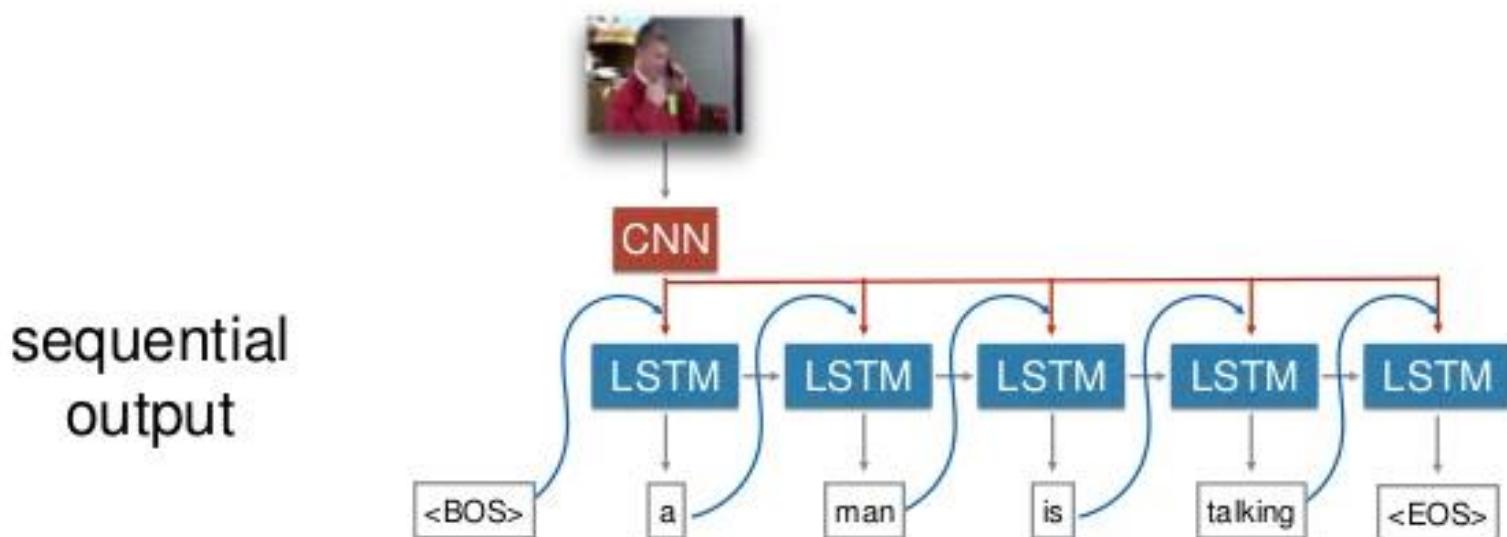
# Deep Learning (LSTM) to predict fraud/authentic user-order



Clicks to detect Fraus , Source : <http://ecmlpkdd2017.ijs.si/papers/paperID69.pdf>

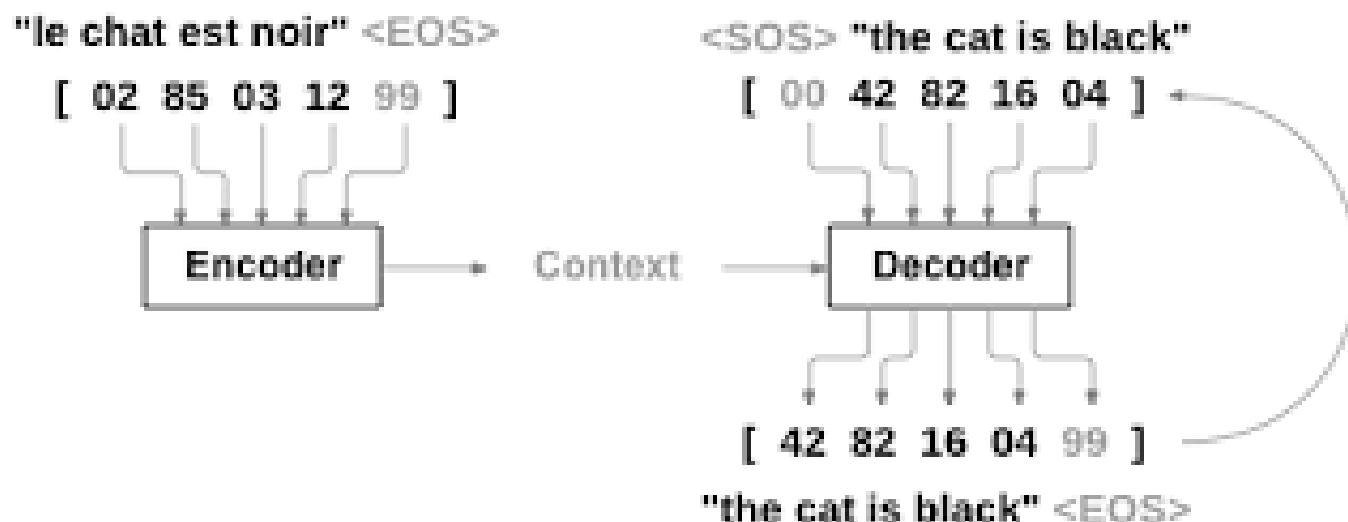
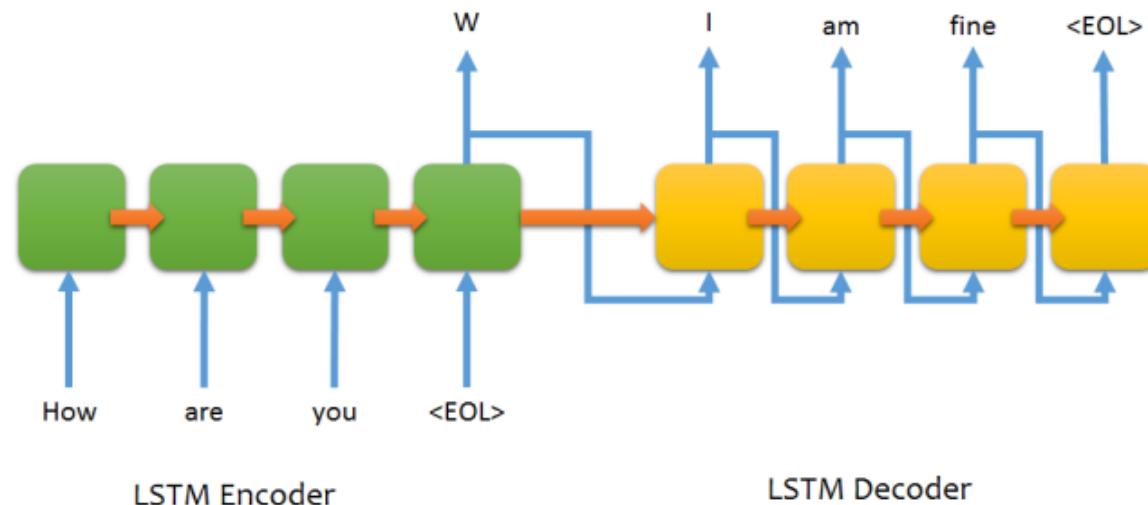
<https://towardsdatascience.com/deep-learning-vs-deep-reinforcement-learning-algorithms-in-retail-industry-ii-9c17c83ecf2f>

# Image Description



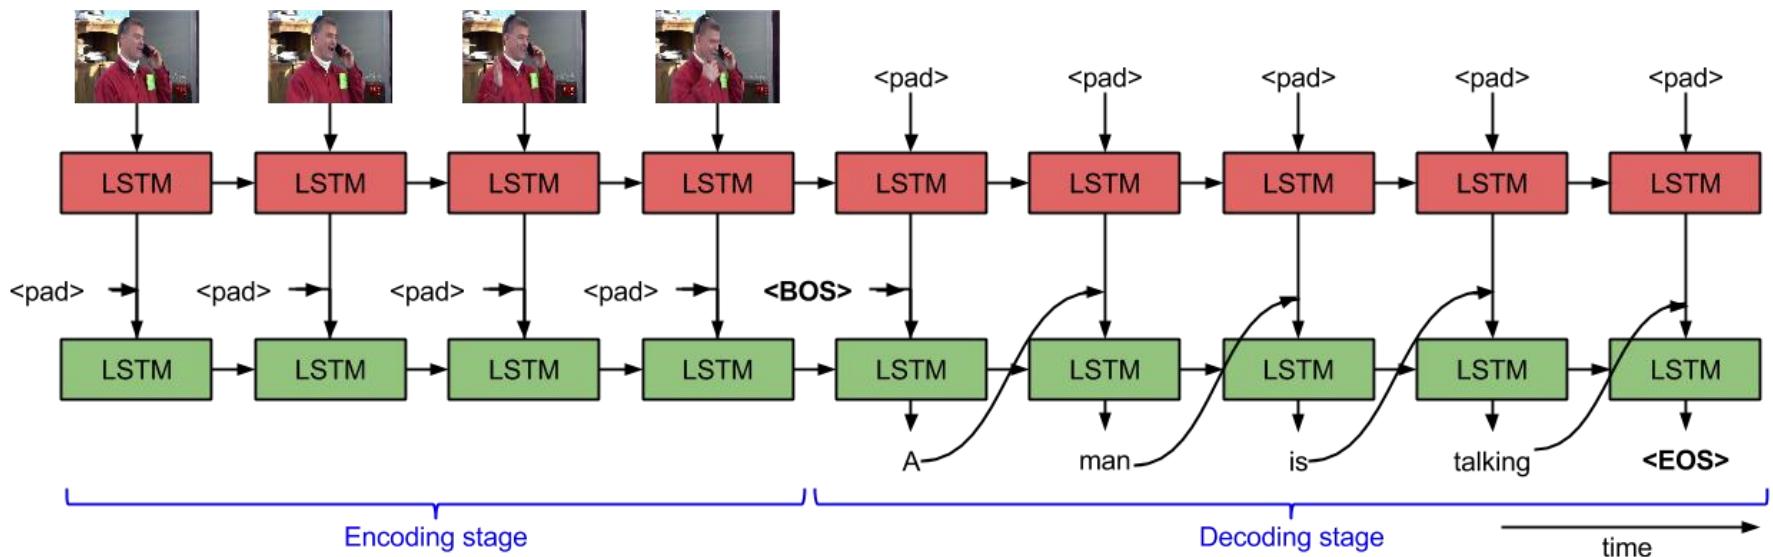
- Jeff Donahue, CVPR Caffe Tutorial, June 6, 2015

# Machine translation, chatbots



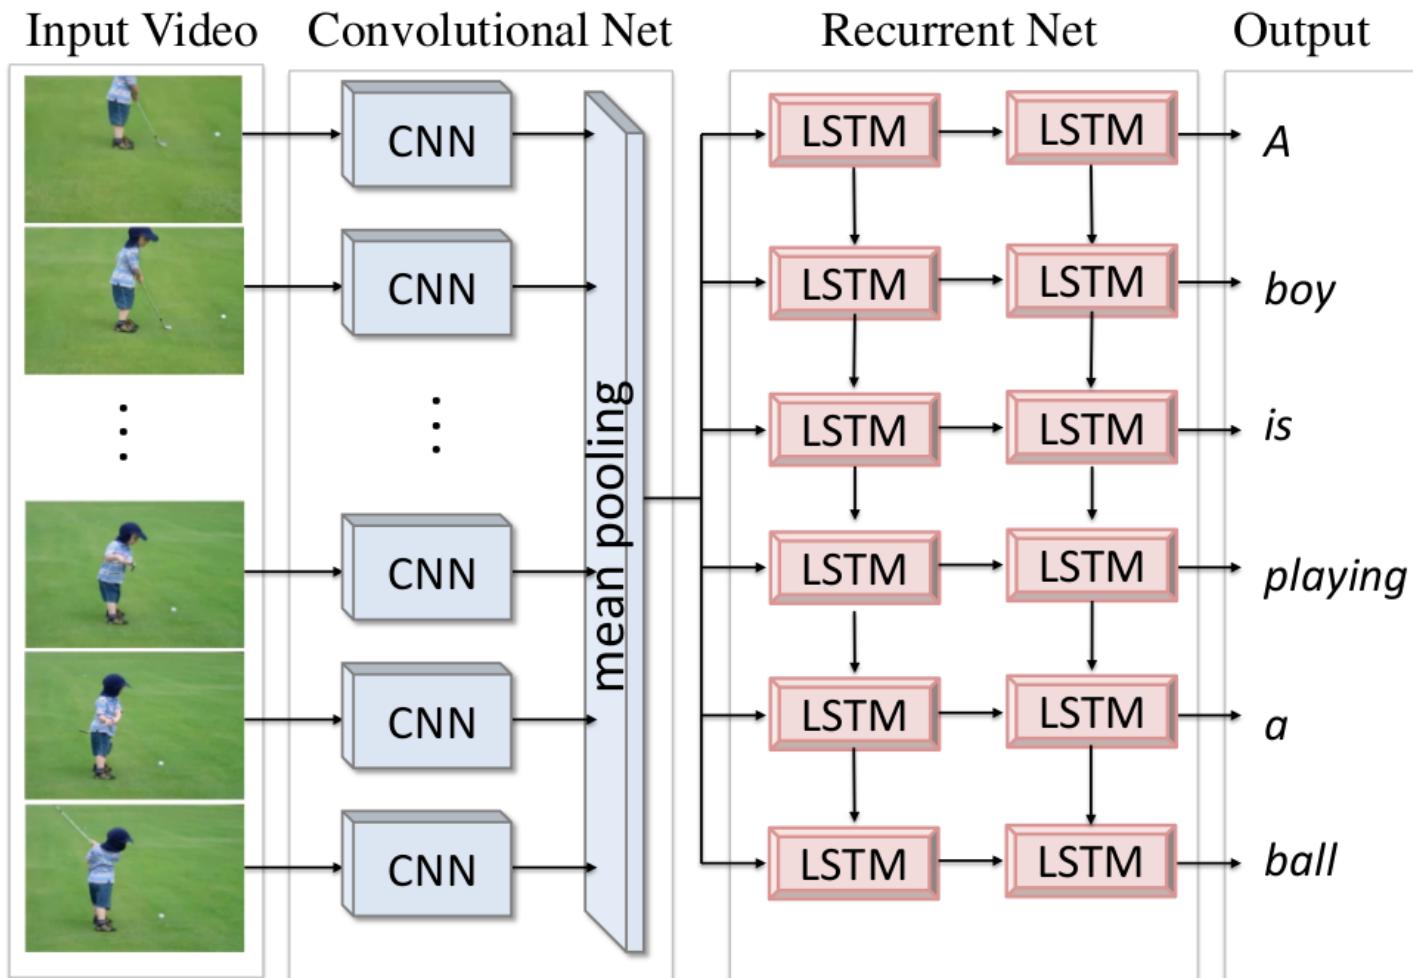
# Sequence to Sequence - Video to Text

Subhashini Venugopalan, et al.



# Translating Videos to Natural Language Using Deep Recurrent Neural Networks

Subhashini Venugopalan, et al.



[Home](#)[Why use Keras](#)[GETTING STARTED](#)[Guide to the Sequential model](#)[Guide to the Functional API](#)[FAQ](#)[MODELS](#)[About Keras models](#)

- **go\_backwards**: Boolean (default False). If True, process the input sequence backwards and return the reversed sequence.
- **stateful**: Boolean (default False). If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
- **unroll**: Boolean (default False). If True, the network will be unrolled, else a symbolic loop will be used. Unrolling can speed-up a RNN, although it tends to be more memory-intensive. Unrolling is only suitable for short sequences.
- **reset\_after**: GRU convention (whether to apply reset gate after or before matrix multiplication). False = "before" (default), True = "after" (CuDNN compatible).

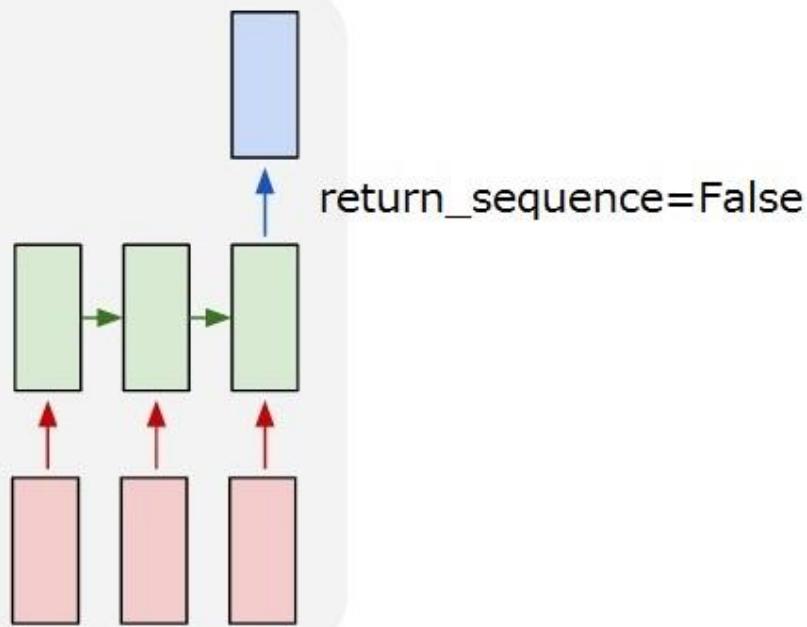
## References

- [Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation](#)
- [On the Properties of Neural Machine Translation: Encoder-Decoder Approaches](#)

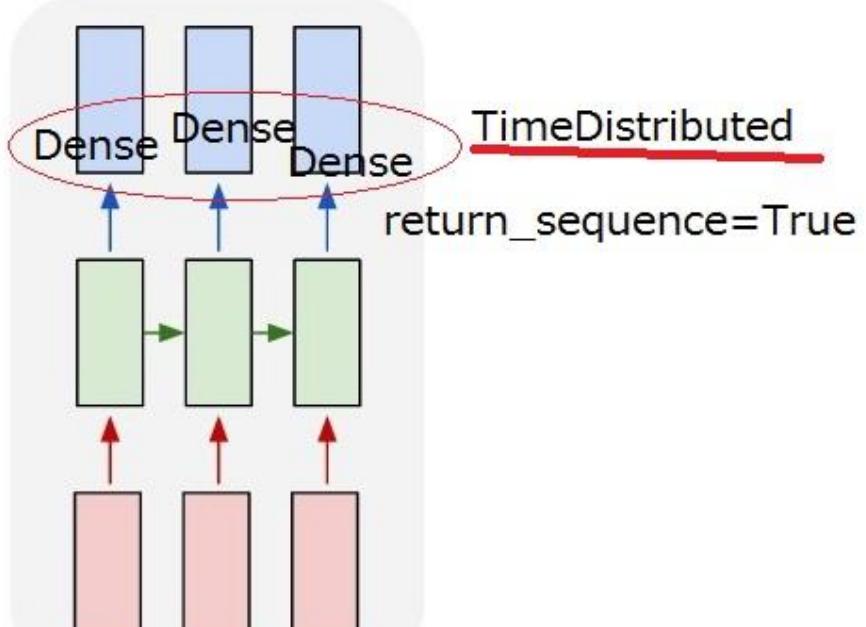
the inputs.

- **recurrent\_dropout**: Float between 0 and 1. Fraction of the units to drop for the linear transformation of the recurrent state.
- **implementation**: Implementation mode, either 1 or 2. Mode 1 will structure its operations as a larger number of smaller dot products and additions, whereas mode 2 will batch them into fewer, larger operations. These modes will have different performance profiles on different hardware and for different applications.
- **return\_sequences**: Boolean. Whether to return the last output in the output sequence, or the full sequence.
- **return\_state**: Boolean. Whether to return the last state in addition to the output. The returned elements of the states list are the hidden state and the cell state, respectively.
- **go\_backwards**: Boolean (default False). If True, process the input sequence backwards and return the reversed sequence.

## many to one



## many to many (2)



# CNN Long Short-Term Memory Networks

by Jason Brownlee on August 21, 2017 in Long Short-Term Memory Networks

[Tweet](#) [Share](#) [Share](#)

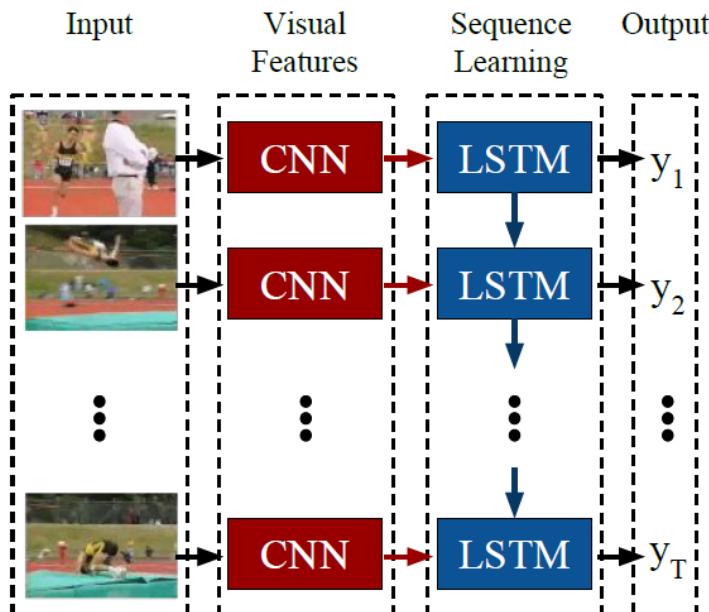
## Gentle introduction to CNN LSTM recurrent neural networks with example Python code.

Welcome to Machine Learning Mastery!



Hi, I'm **Jason Brownlee**, PhD.  
I write tutorials to help  
developers (*like you*) get results  
with machine learning.

[Read More](#)



```
1 model = Sequential()
2 # define CNN model
3 model.add(TimeDistributed(Conv2D(...)))
4 model.add(TimeDistributed(MaxPooling2D(...)))
5 model.add(TimeDistributed(Flatten()))
6 # define LSTM model
7 model.add(LSTM(...))
8 model.add(Dense(...))
```

# RNN essentially describe programs:

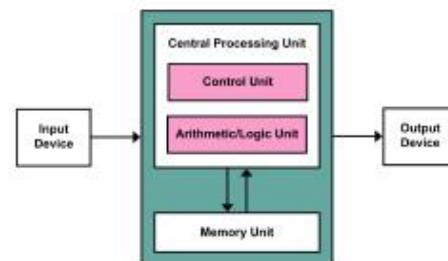
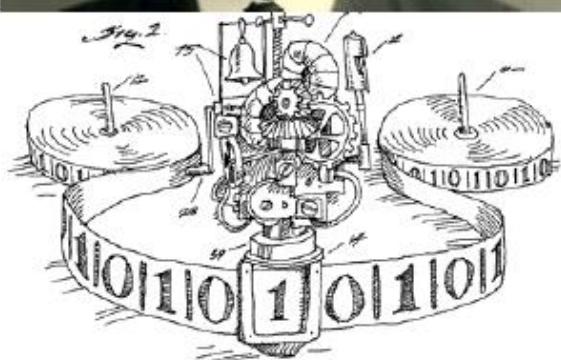
inputs + some internal variables.

In fact, it is known that RNNs are Turing-Complete in the sense that they can simulate arbitrary programs (with proper weights).

# Neural Turing Machines

Can neural nets learn programs?

Alex Graves  
Greg Wayne  
Ivo Danihelka



...path to AI??

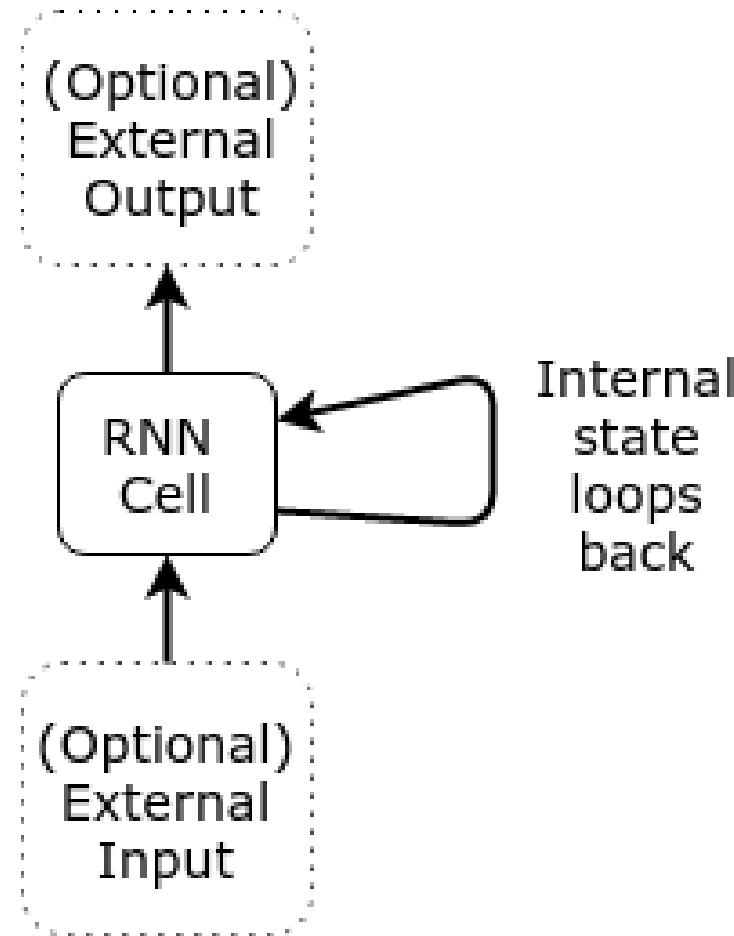
.... But “*forget I said anything.*”

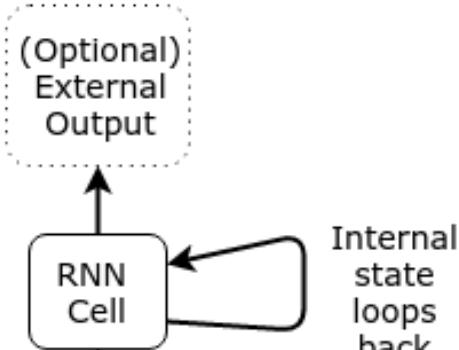
Andrej Karpathy

<http://karpathy.github.io/2015/05/21/rnn-effectiveness>

# Written Memories: Understanding, Deriving and Extending the LSTM

<http://r2rt.com/written-memories-understanding-deriving-and-extending-the-lstm.html>





(Optional)  
External  
Input

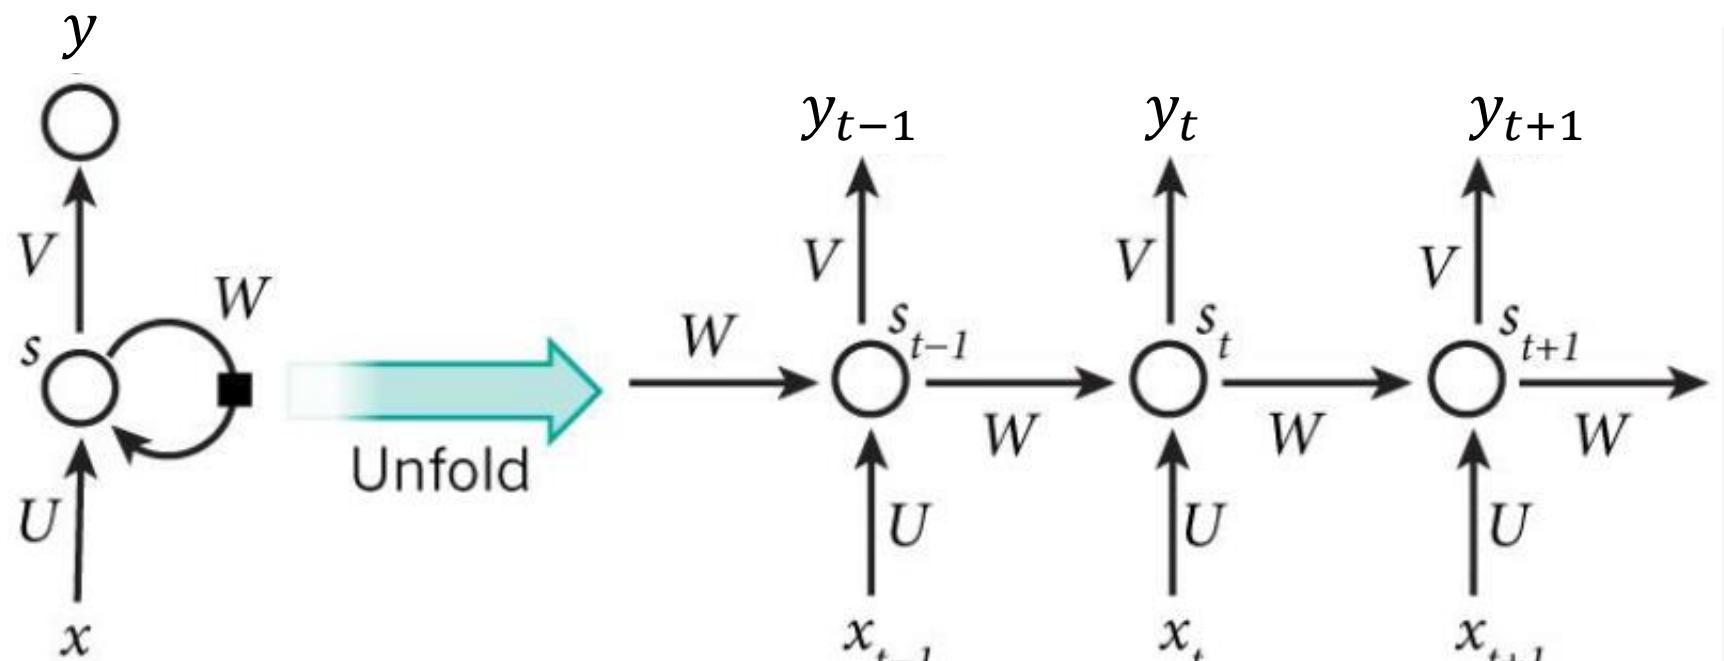
(Optional)  
External  
Output

Here is the algebraic description of the RNN cell:

$$\begin{pmatrix} s_t \\ o_t \end{pmatrix} = f \begin{pmatrix} s_{t-1} \\ x_t \end{pmatrix}$$

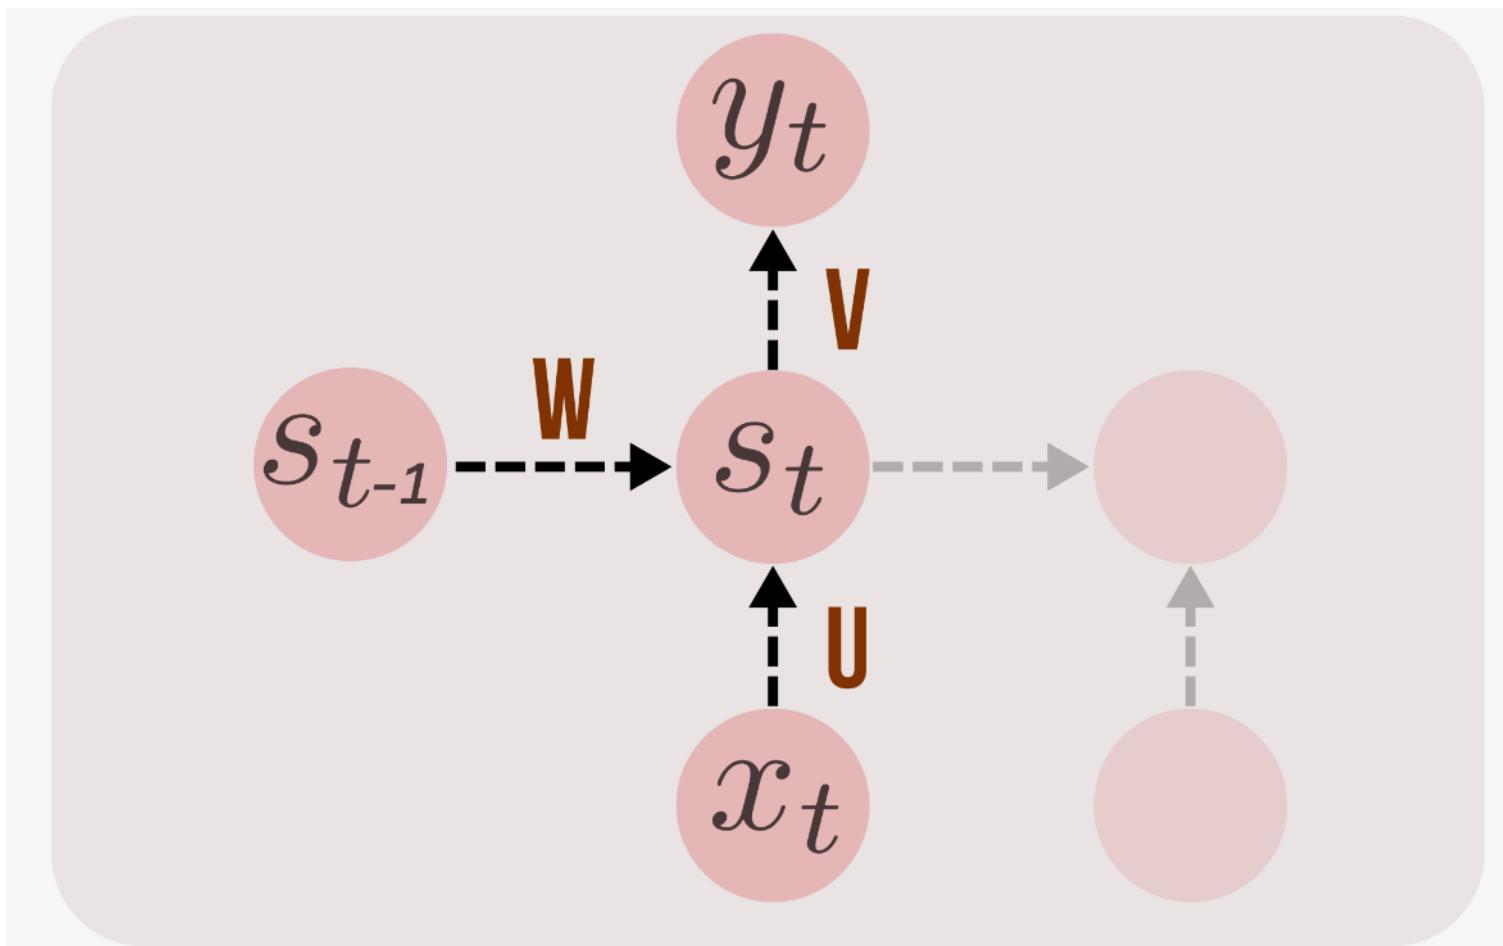
where:

- $s_t$  and  $s_{t-1}$  are our current and prior states,
- $o_t$  is our (possibly empty) current output,
- $x_t$  is our (possibly empty) current input, and
- $f$  is our recurrent function.



# TRAINING RNN: W, U, V ?

- BPTT: *Backpropagation Through Time*
- Truncated



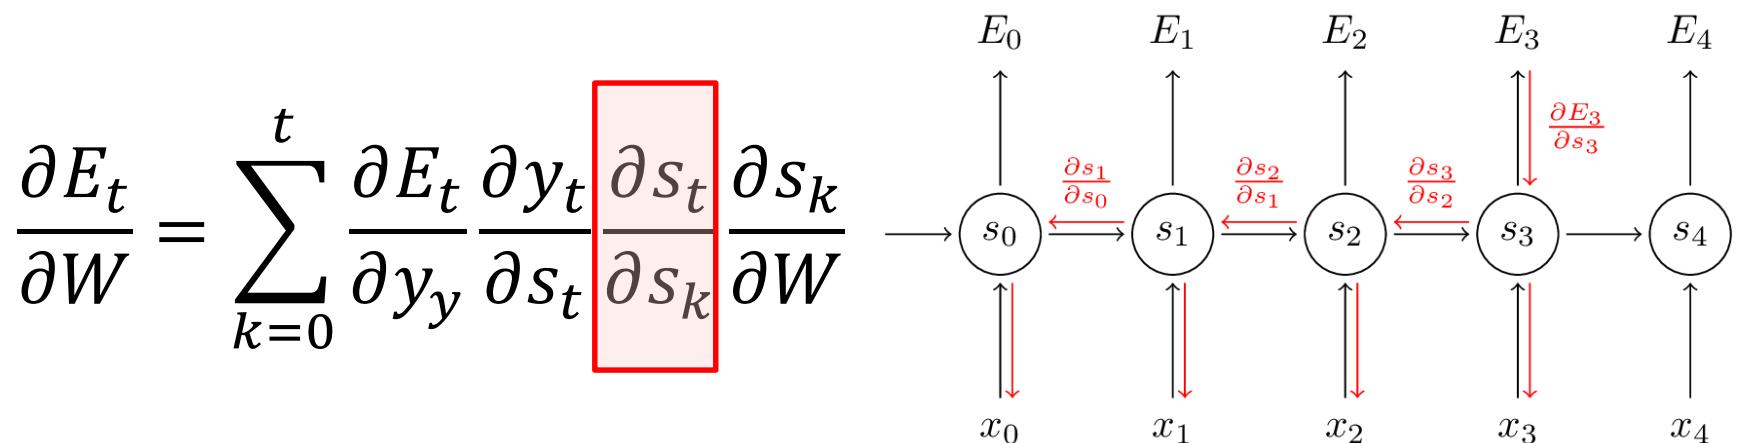
# *Backpropagation Through Time (BPTT)*

- Because the parameters are shared by all time steps in the network
- The gradient at each output depends not only on the calculations of the current time step, but also the previous time steps.

<http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-2-implementing-a-language-model-rnn-with-python-numpy-and-theano/>

# Backpropagation Through Time (BPTT)

- But the recurrent net can be seen as a (very deep) feedforward net with **shared weights**
  - The forward pass builds up a stack of the activities of all the units at each time step.
  - The backward pass peels activities off the stack to compute the error derivatives at each time step.
  - After the backward pass we add together the derivatives at all the different times for each weight



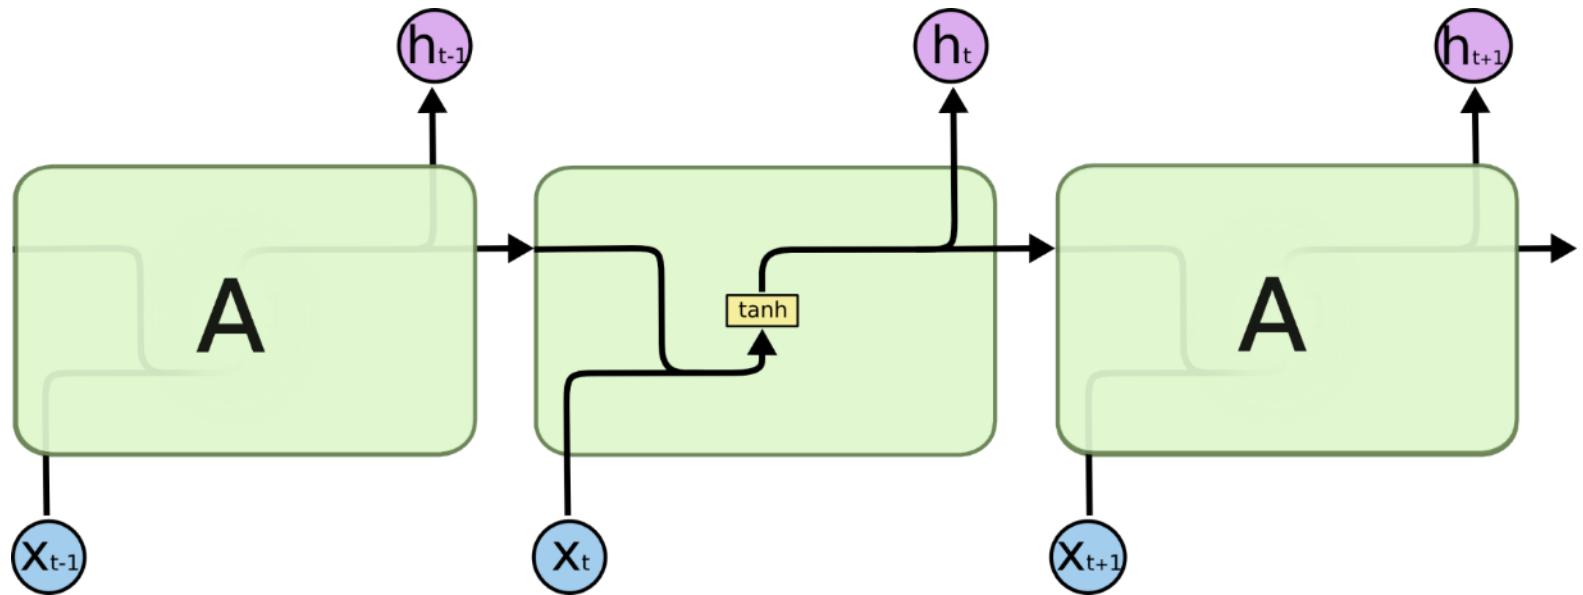
# Basic concepts: LSTM & GRU

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

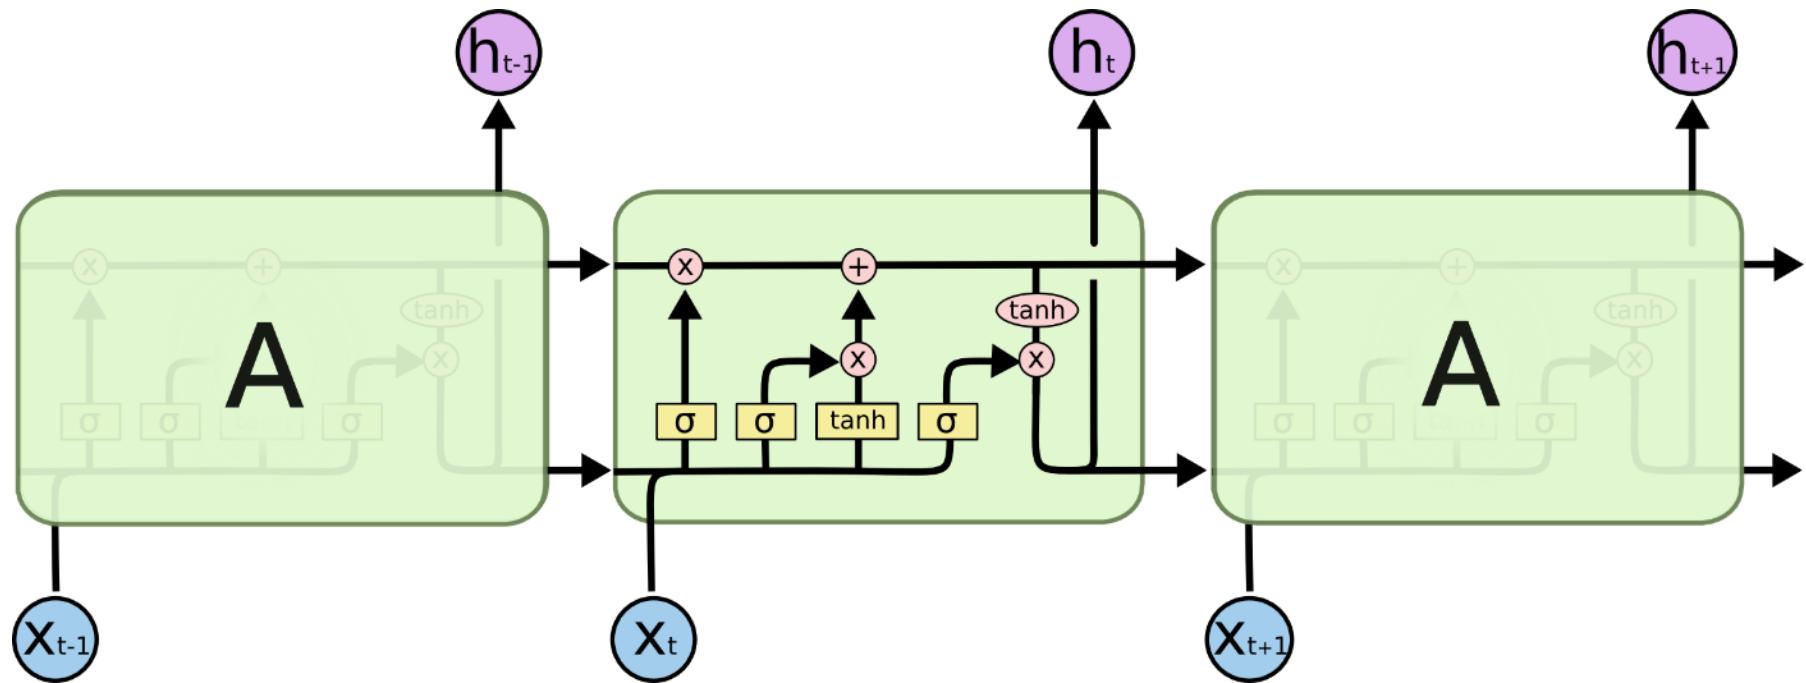
- **Vanishing gradients**: a problem?  
In an RNN trained on long sequences (e.g. 100 time steps) the gradients can easily **explode** or **vanish**.
- Are RNNs capable of handling “**long-term dependencies**?”

Long Short Term Memory networks – LSTMs  
Hochreiter and Schmidhuber (1997)

*NOTATION: states =  $h_t$*

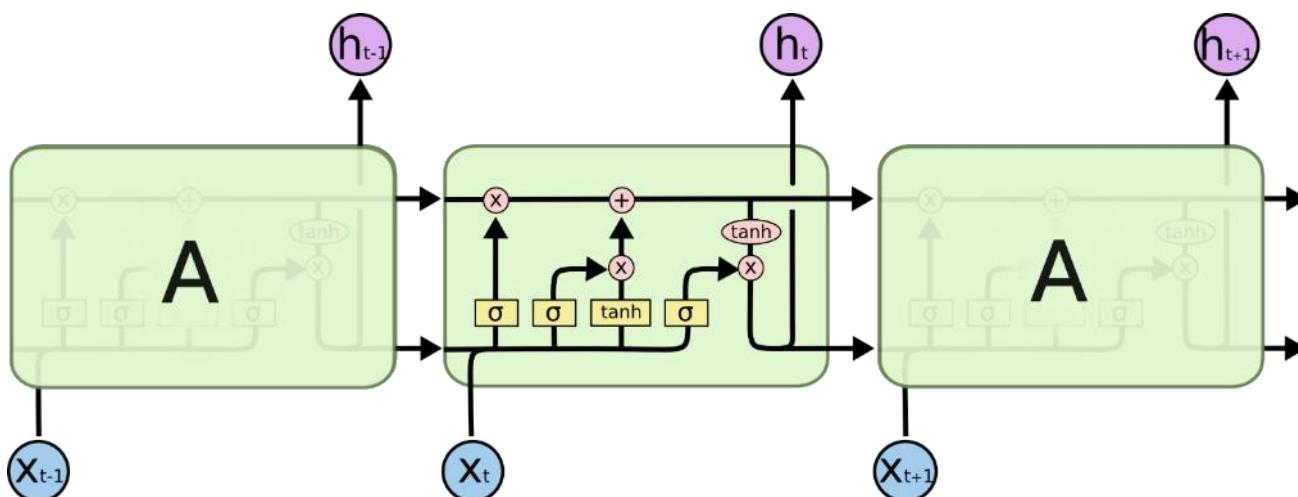
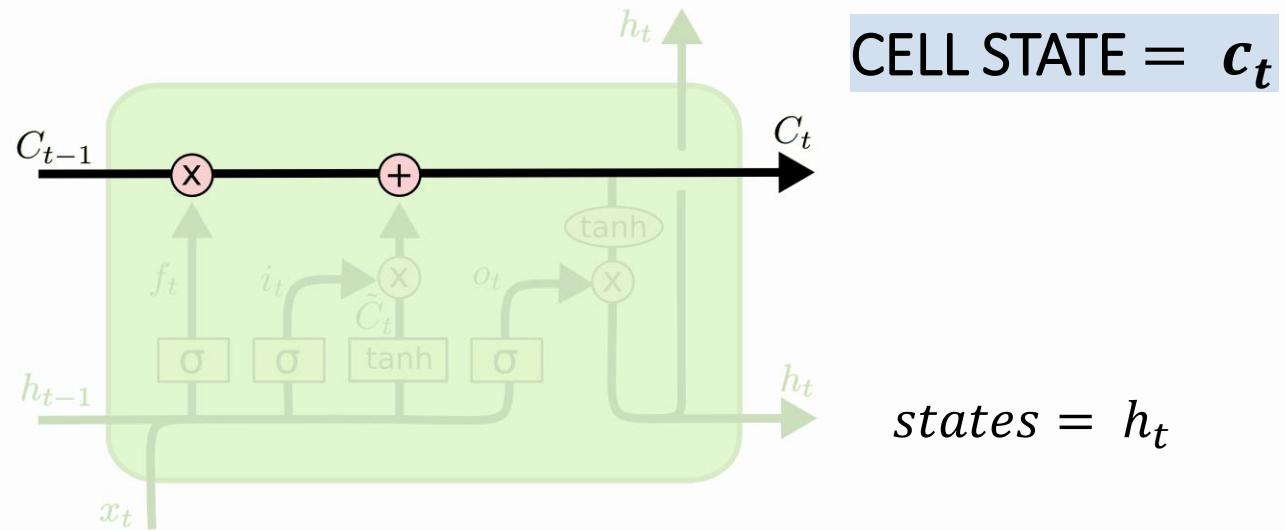


The repeating module in a standard RNN contains a single layer



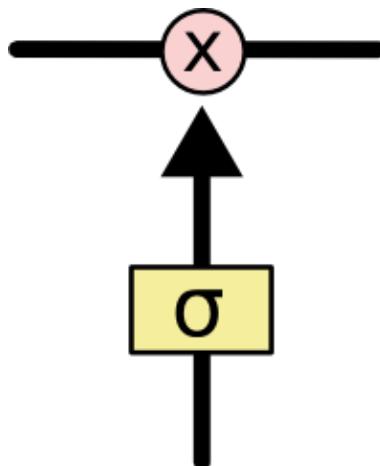
The repeating module in LSTM contains four interacting layers

# Core Idea Behind LSTM



# Core Idea Behind LSTM

- The LSTM does have the ability to **remove or add information to the cell state**, carefully regulated by structures called *gates*.
- **Gates** are a way to optionally let information through.
- They are composed out of a sigmoid neural net layer and a pointwise multiplication operation.



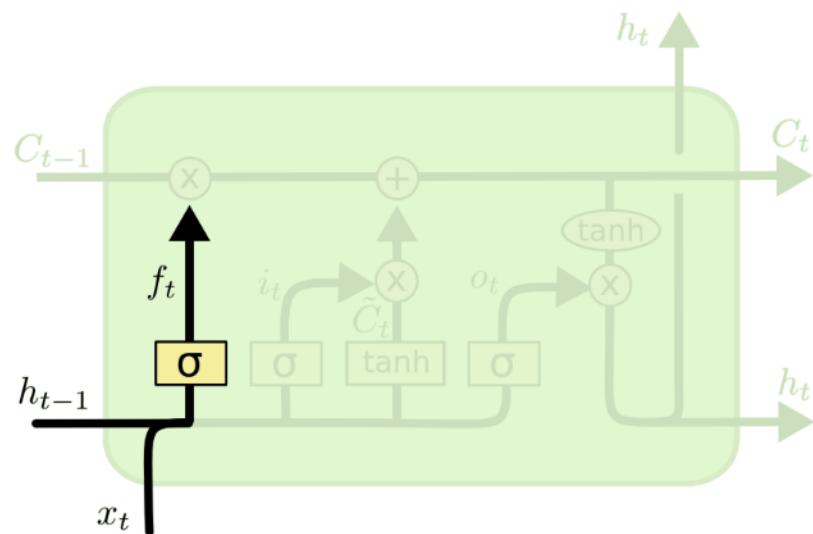
A value of zero means “let nothing through,” while a value of one means “let everything through!”

An LSTM has three of these gates, **to protect and control the cell state**

# Step-by-Step LSTM Walk Through

**Forget gate:** The first step in our LSTM is to decide what information we're going to throw away from the cell state.

*When we “see” something new relevant we decide forget....*

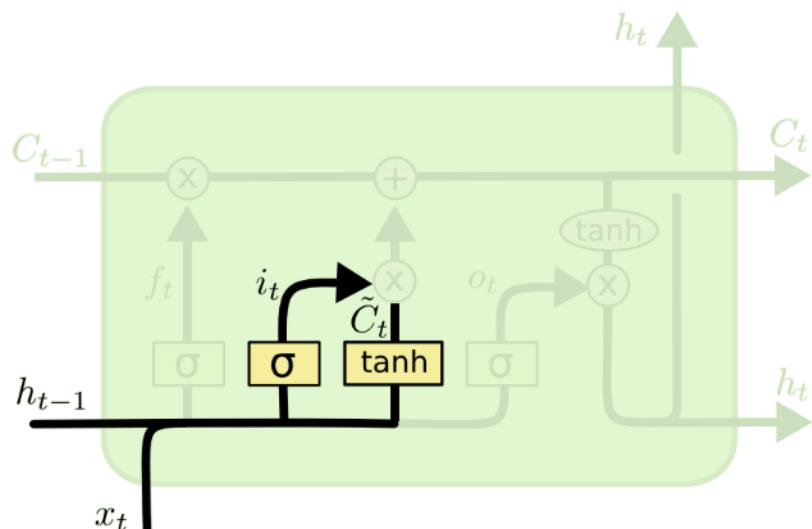


$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

# Step-by-Step LSTM Walk Through

The next step is to decide what new information we're going to store in the cell state. This has two parts.

- First, a sigmoid layer called the “**Input gate** layer” decides which values we'll update.
- Next, a **tanh layer** creates a vector of new candidate values,  $\tilde{C}_t$ , that could be added to the state.



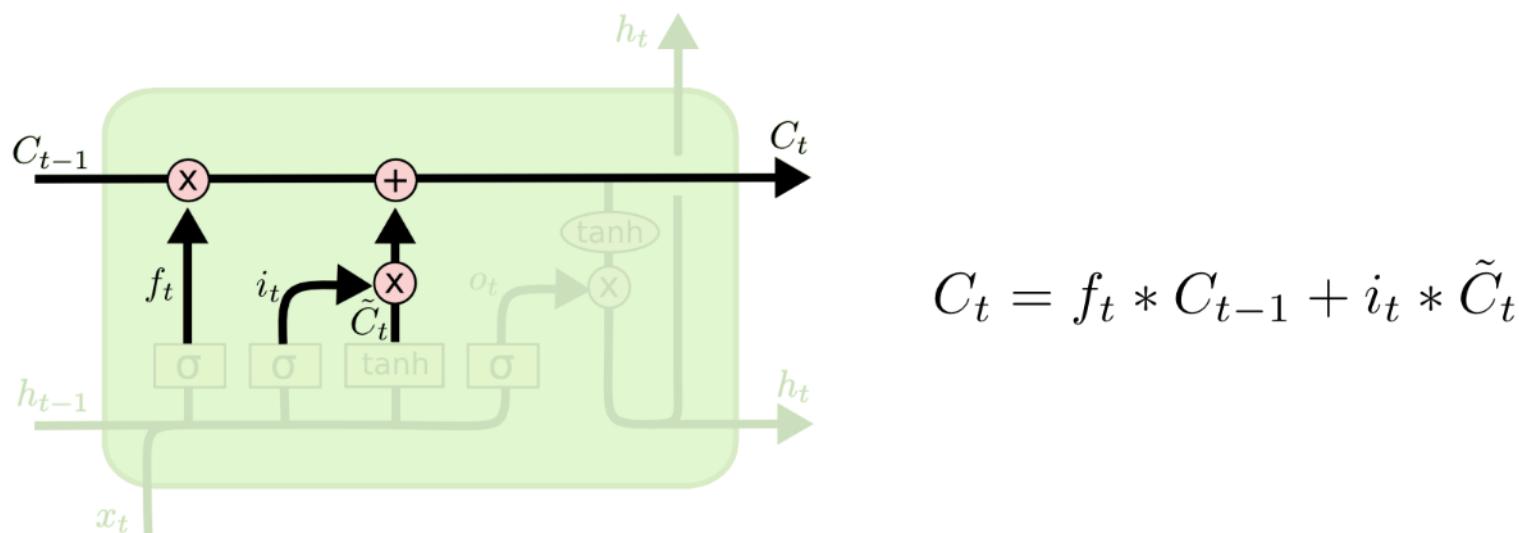
$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

# Step-by-Step LSTM Walk Through

It's now time to update the old cell state,  $C_{t-1}$ , into the new cell state  $C_t$

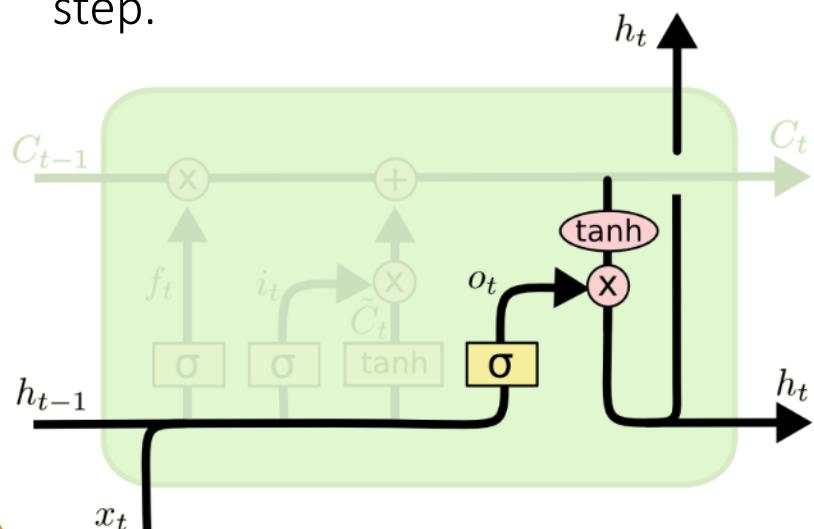
- The previous steps already decided what to do, we just need to actually do it.
- We multiply the old state by  $f_t$  (forget gate) then we add  $i_t * \tilde{C}_t$



# Step-by-Step LSTM Walk Through

Finally, we need to decide **what we're going to output**  $h_t$ . This output will be based on our cell state, but will be a filtered version.

- First, we pass the previous hidden state and the current input into a sigmoid function.
- Then we pass the newly modified cell state to the tanh function. We multiply the tanh output with the sigmoid output to decide what information the hidden state should carry. The output is the hidden state.
- The new cell state and the new hidden is then carried over to the next time step.



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh (C_t)$$

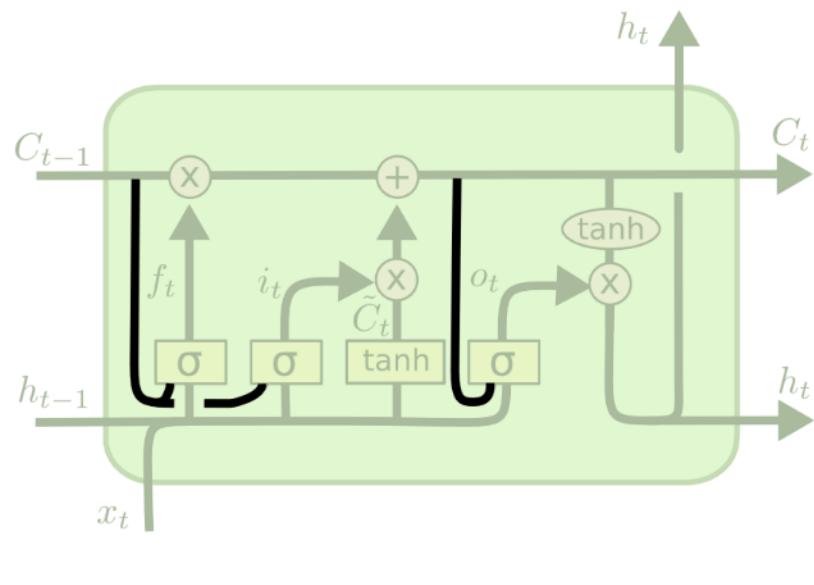
You can also see:

<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

# Variants on Long Short Term Memory

One popular LSTM variant, introduced by Gers & Schmidhuber (2000), is adding “**peephole connections**.”

- This means that we let the gate layers look at the cell state.



$$f_t = \sigma (W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$
$$i_t = \sigma (W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$
$$o_t = \sigma (W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$

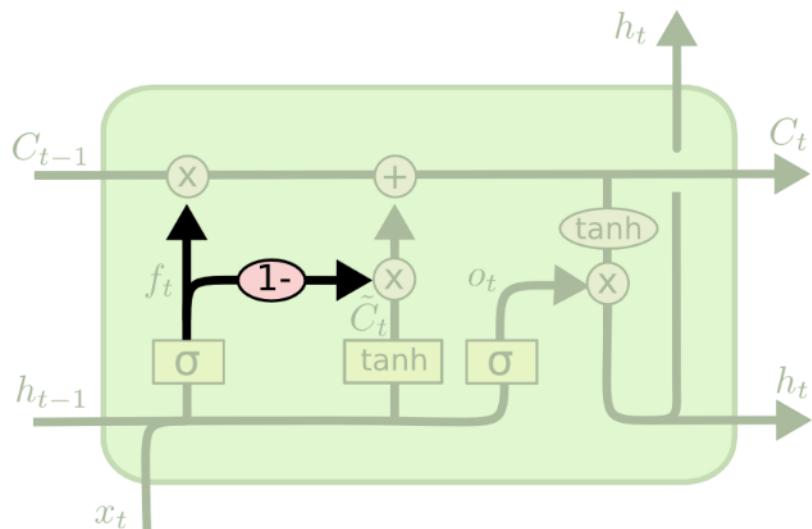


*The above diagram adds peepholes to all the gates, but many papers will give some peepholes and not others.*

# Variants on Long Short Term Memory

Another variation is to use **coupled forget and input gates**.

- We only input new values to the state when we forget something older



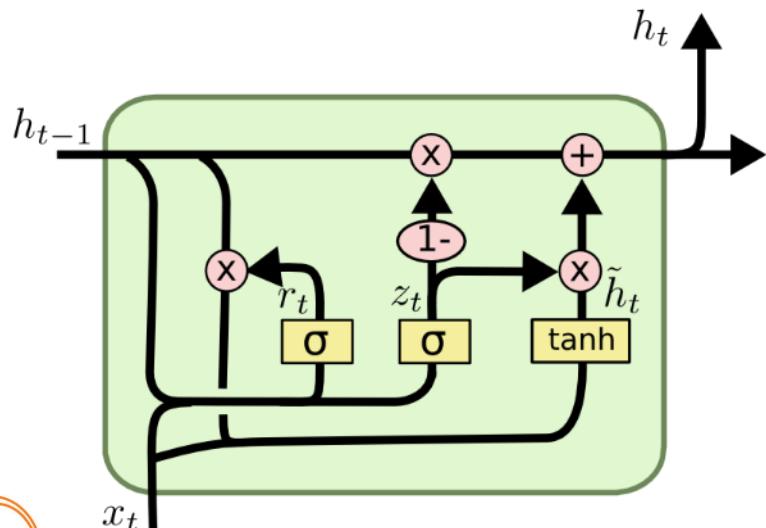
$$C_t = f_t * C_{t-1} + (1 - f_t) * \tilde{C}_t$$

# Variants on Long Short Term Memory

Gated Recurrent Unit, or **GRU**, introduced by Cho, et al. (2014).

- It **combines the forget and input gates** into a single “update gate.”
- It also **merges the cell state and hidden state**, and makes some other changes.

The resulting model is simpler than standard LSTM models, and has been growing increasingly popular.



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

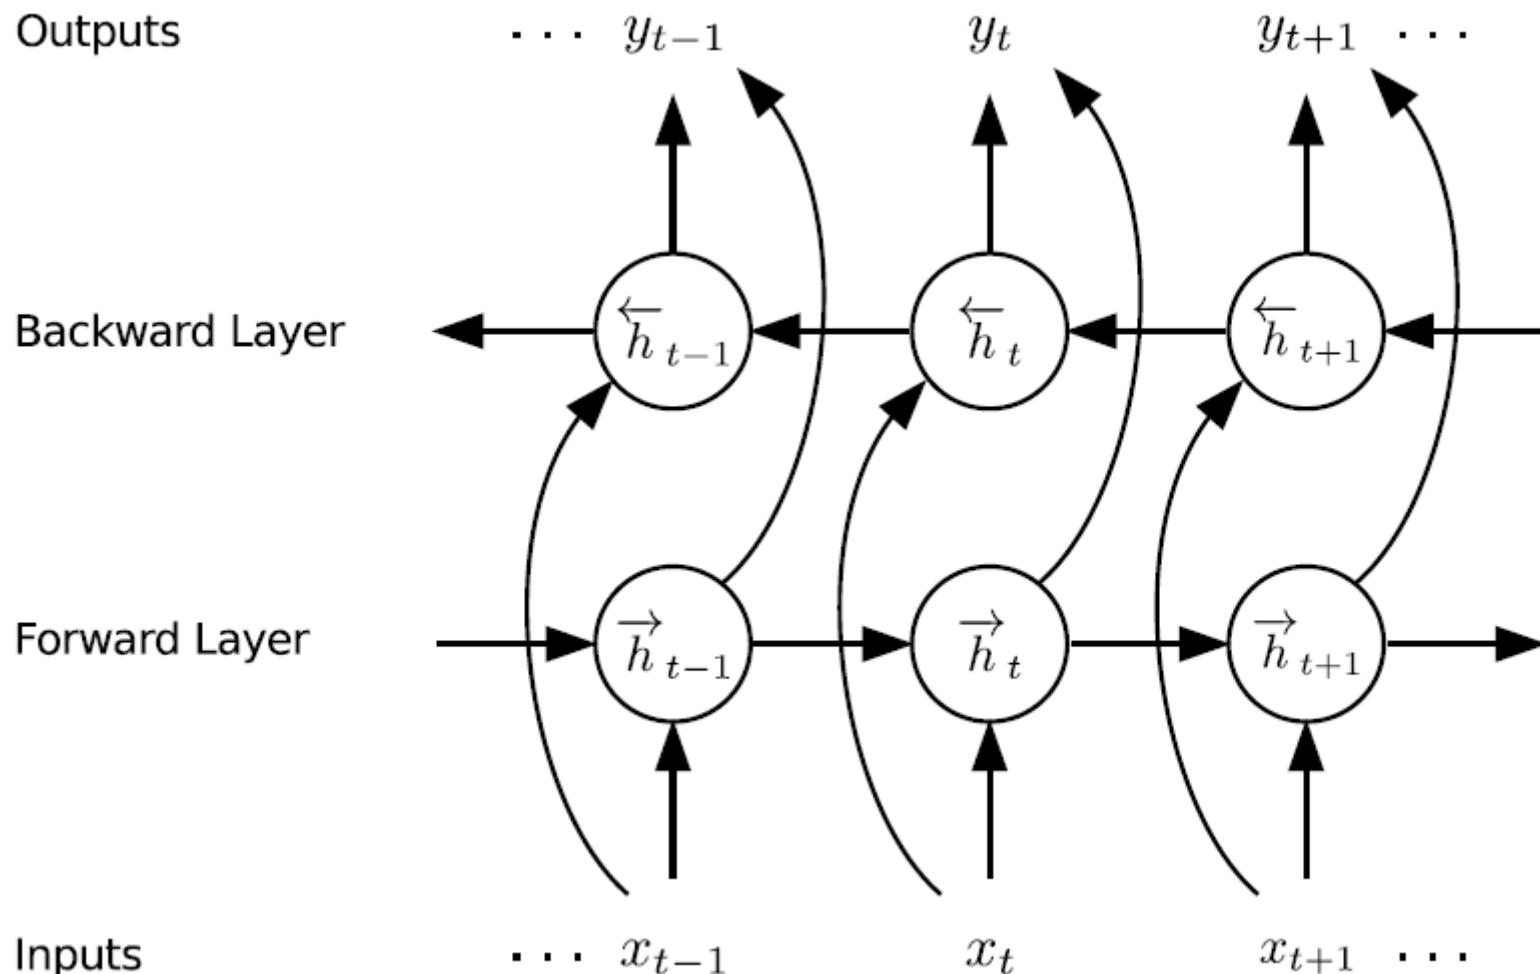
$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# Variants on Long Short Term Memory

These are only a few of the most notable LSTM variants.

See for example: **BLSTM Bi-directional LSTM**



# Variants on Long Short Term Memory

**Which of these variants is best?** Do the differences matter?

- Greff, et al. (2015) do a nice comparison of popular variants, finding that they're all about the same.
- Jozefowicz, et al. (2015) tested more than ten thousand RNN architectures, finding some that worked better than LSTMs on certain tasks.

# Simple Examples: NLP & MUSIC

## Text Prediction/Generation with Keras using LSTM: Long Short Term Memory networks

In this example we will work with the book: Alice's Adventures in Wonderland by Lewis Carroll.

We are going to learn the dependencies between characters and the conditional probabilities of characters in sequences so that we can in turn generate wholly new and original sequences of characters.



Adapted from:

[Text Generation With LSTM Recurrent Neural Networks](#)

By Jason Brownlee

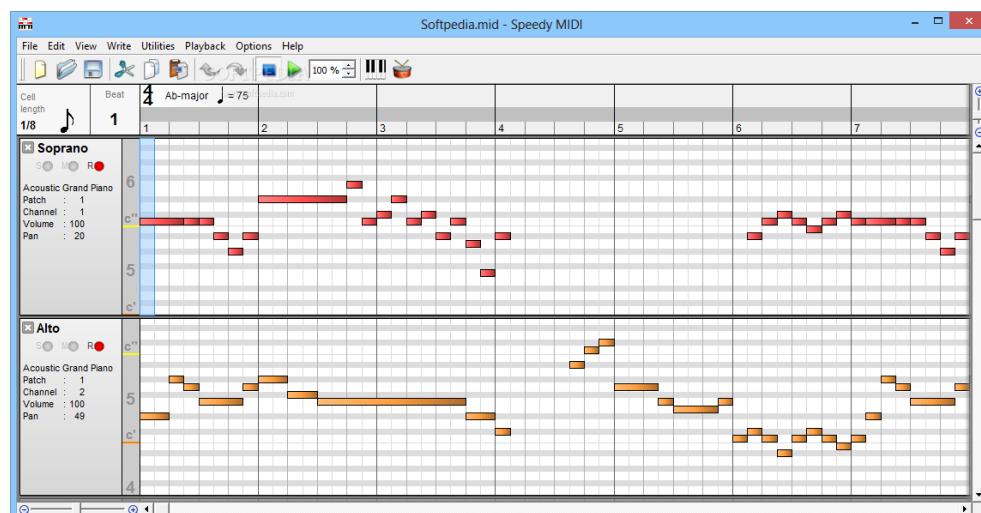
### See example using the IMDB movie review sentiment prediction.

"IMDB dataset", a set of 50,000 highly-polarized reviews from the Internet Movie Database. They are split into 25,000 reviews for training and 25,000 reviews for testing, each set consisting in 50% negative and 50% positive reviews.

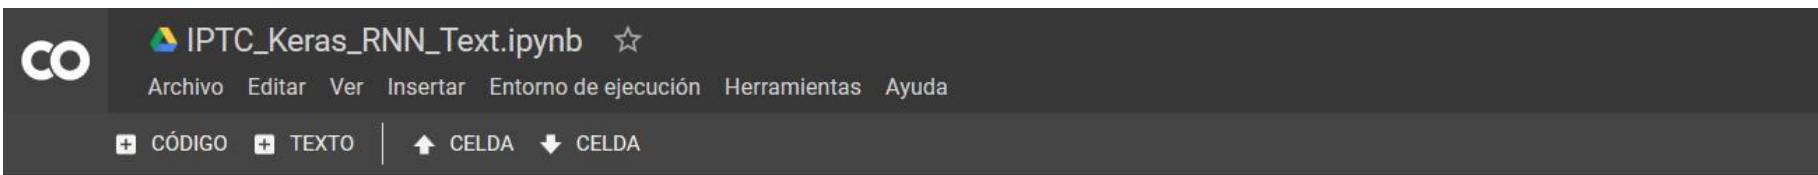
Just like the MNIST dataset, the IMDB dataset comes packaged with Keras. It has already been preprocessed: the reviews (sequences of words) have been turned into sequences of integers, where each integer stands for a specific word in a dictionary.



## MIDI Music Generation



# Simple Examples (I): NLP & MUSIC

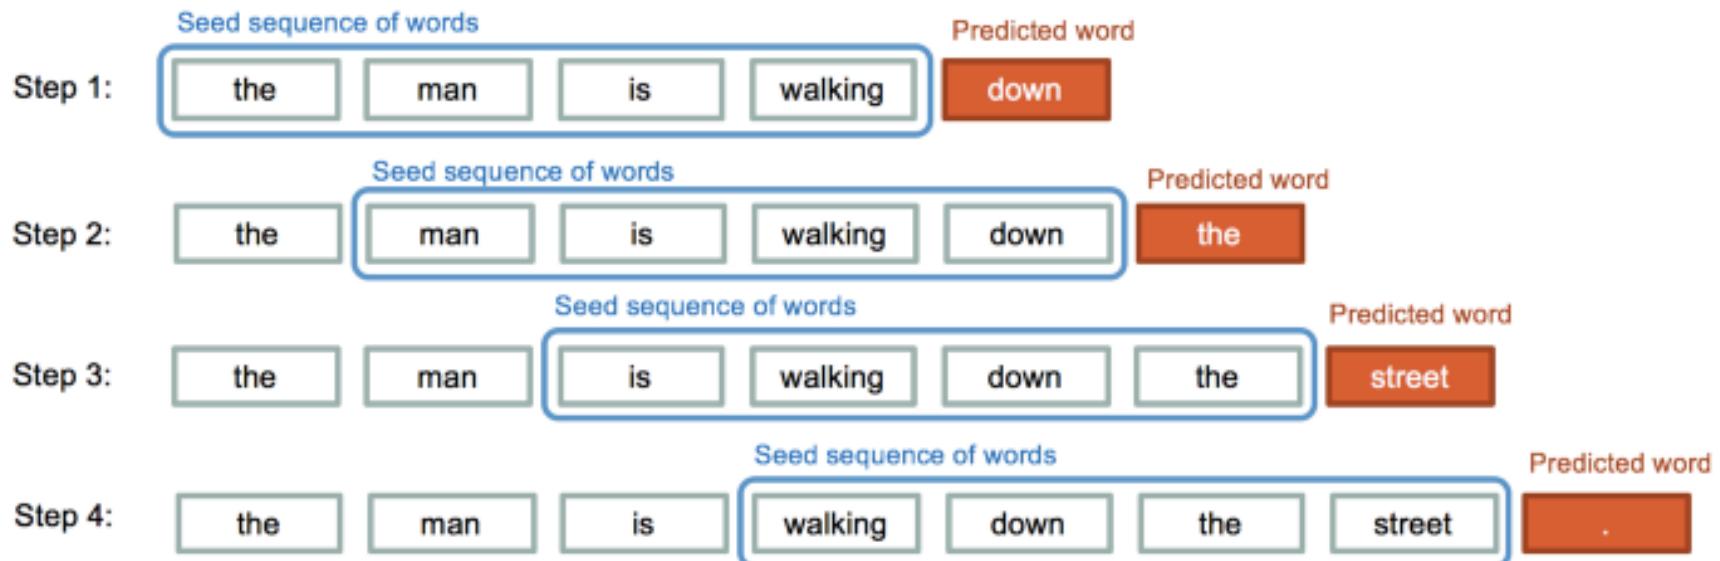


The screenshot shows the Google Colab interface. At the top, there's a file menu with "IPTC\_Keras\_RNN\_Text.ipynb" and a star icon. Below the menu are standard options: Archivo, Editar, Ver, Insertar, Entorno de ejecución, Herramientas, Ayuda. Underneath the menu bar are buttons for "+ CÓDIGO" and "+ TEXTO", and arrows for "CELDA". The main content area has a header "Text Prediction/Generation with Keras using LSTM: Long Short Term Memory networks". A text block below it says: "In this example we will work with the book: Alice's Adventures in Wonderland by Lewis Carroll. We are going to learn the dependencies between characters and the conditional probabilities of characters in sequences so that we can in turn generate wholly new and original sequences of characters." Overlaid on the bottom right of the text block is a yellow box containing the text "Google Colab: Reads a shared file in a Drive". Below the text block is a 3D rendering of Alice and the White Rabbit from Disney's Alice in Wonderland.

Text Prediction/Generation with Keras using **LSTM: Long Short Term Memory networks**

In this example we will work with the book: Alice's Adventures in Wonderland by Lewis Carroll. We are going to learn the dependencies between characters and the conditional probabilities of characters in sequences so that we can in turn generate wholly new and original sequences of characters.

Google Colab: Reads a shared file in a Drive



<https://medium.com/@david.campion/text-generation-using-bidirectional-lstm-and-doc2vec-models-1-3-8979eb65cb3a>

# OpenAI GPT-2 writes alternate endings for Game of Thrones



Chintan Trivedi in Towards Data Science

[Follow](#)

May 23 · 12 min read ★

I trained the GPT-2 language model on GRRM's book series "*A Song of Ice and Fire*" and let it complete the HBO show's storyline. Can it do better than HBO's season 8 train-wreck?

<https://towardsdatascience.com/openai-gpt-2-writes-alternate-endings-for-game-of-thrones-c9be75cd2425>

# Simple Examples (II): NLP & MUSIC

The screenshot shows the Google Colab interface. At the top, there's a toolbar with 'Archivo', 'Editar', 'Ver', 'Insertar', 'Entorno de ejecución', 'Herramientas', and 'Ayuda'. Below the toolbar, there are buttons for 'CÓDIGO' and 'TEXTO', and arrows for 'CELDA'. A status bar at the bottom says 'are designed to remember information for long periods'. On the left, there's a sidebar with a 'Music generation' section.

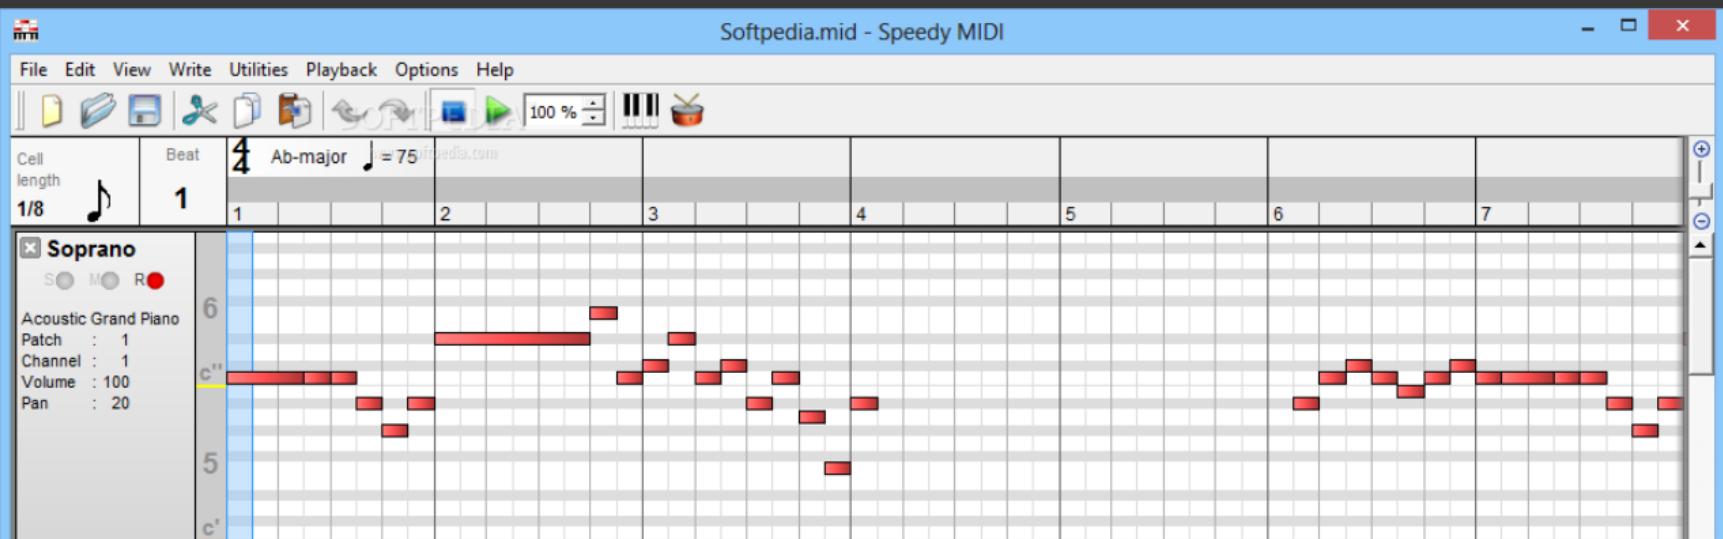
Google Colab:

- mounts your google drive

## Music generation

To train our model, we will use [MIDI files](#). MIDI files contain **information about a music composition**, not the music itself. They contain information about notation, pitch, velocity, vibrato, panning, and clock signals (which set tempo).

There are some synthesizers which are able to transform this composition information into a real audio track. Our model will learn from these **MIDI files** and will be able to generate new ones.



# Simple Examples (III): NLP & MUSIC

## Sentiment Analysis

### See example using the IMDB movie review sentiment prediction.

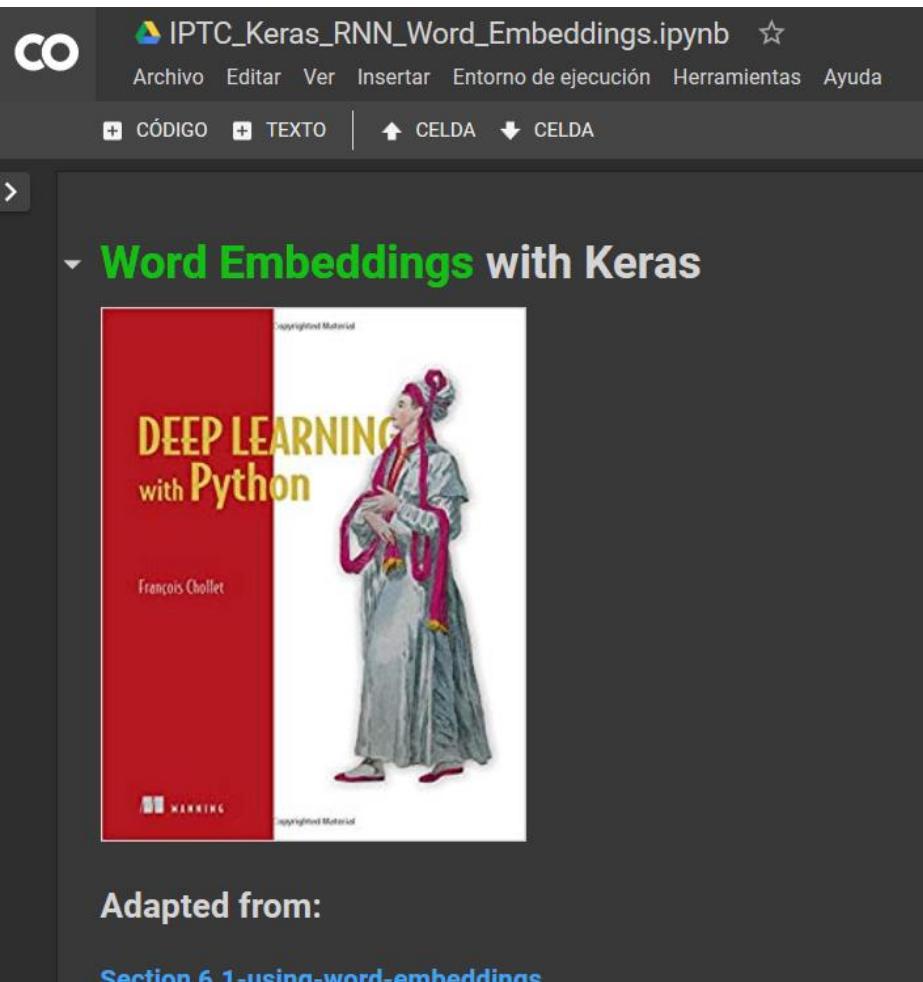
"IMDB dataset", a set of 50,000 highly-polarized reviews from the Internet Movie Database. It contains 25,000 reviews for training and 25,000 reviews for testing, each set containing 500 different movies. The "review" column contains the movie review text, and the "sentiment" column contains the label ("positive" or "negative").

Just like the MNIST dataset, the IMDB dataset comes packaged with Keras. The movie reviews (sequences of words) have been turned into sequences of integer word indices in a dictionary.

Let's quickly prepare the data.

- We will restrict the movie reviews to the top 10,000 most common words
- and cut the reviews after only 20 words.

Our network will simply learn 8-dimensional embeddings for each of the 10,000 words, turn the input sequences (3D float tensor), flatten the tensor to 2D, and train a single Dense layer on top for classification.



The screenshot shows a Jupyter Notebook interface. The title bar says 'IPTC\_Keras\_RNN\_Word\_EMBEDDINGS.ipynb'. The menu bar includes Archivo, Editar, Ver, Insertar, Entorno de ejecución, Herramientas, and Ayuda. Below the menu is a toolbar with CÓDIGO, TEXTO, CELDA, and CELDA. The main area has a section titled 'Word Embeddings with Keras' which is expanded. To the right of the section title is an image of the book 'Deep Learning with Python' by François Chollet. The book cover is red with yellow text and features a drawing of a person in traditional Chinese attire.

Adapted from:

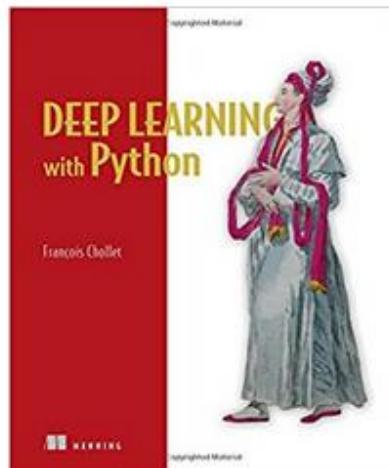
[Section 6.1-using-word-embeddings](#)

# Embedding:

Indices by themselves, carry no **semantic** meaning

[\*\*IPTC\\_Keras\\_RNN\\_Word\\_EMBEDDINGS.ipynb\*\*](#)

## Word Embeddings with Keras

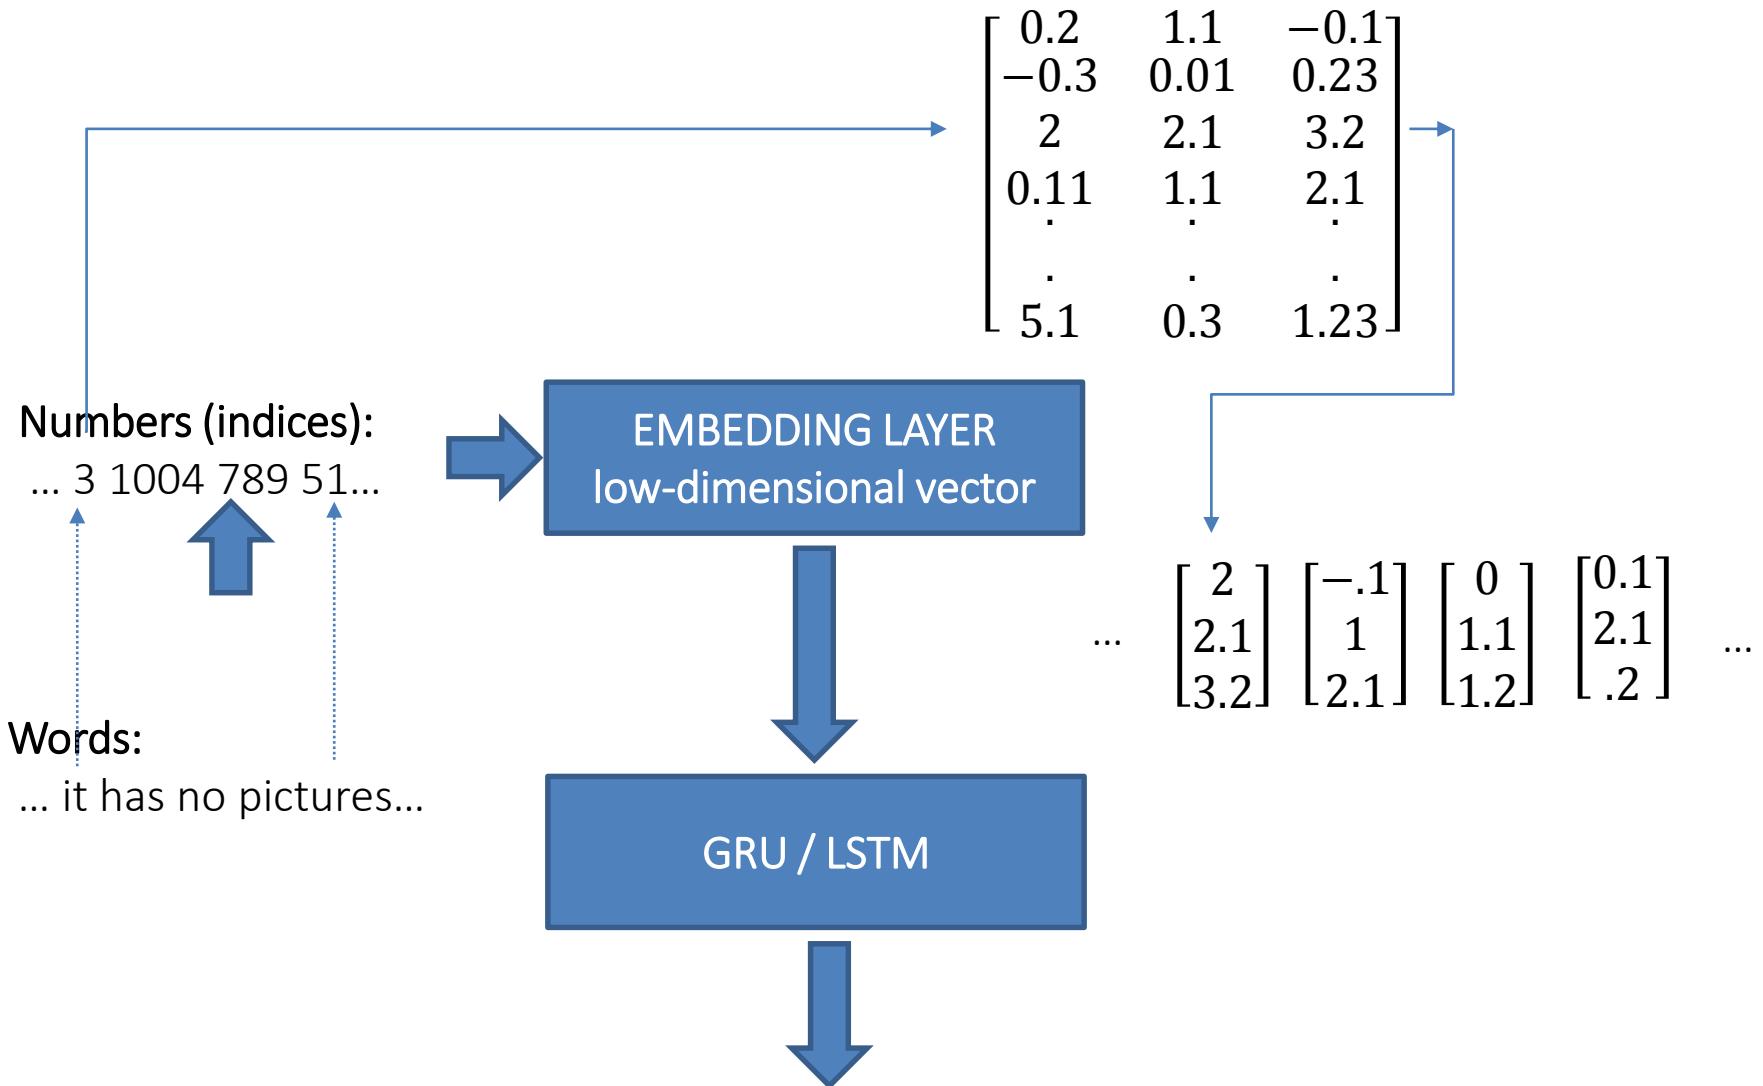


Adapted from:

[6.1-using-word-embeddings](#)

By François Chollet

K Keras

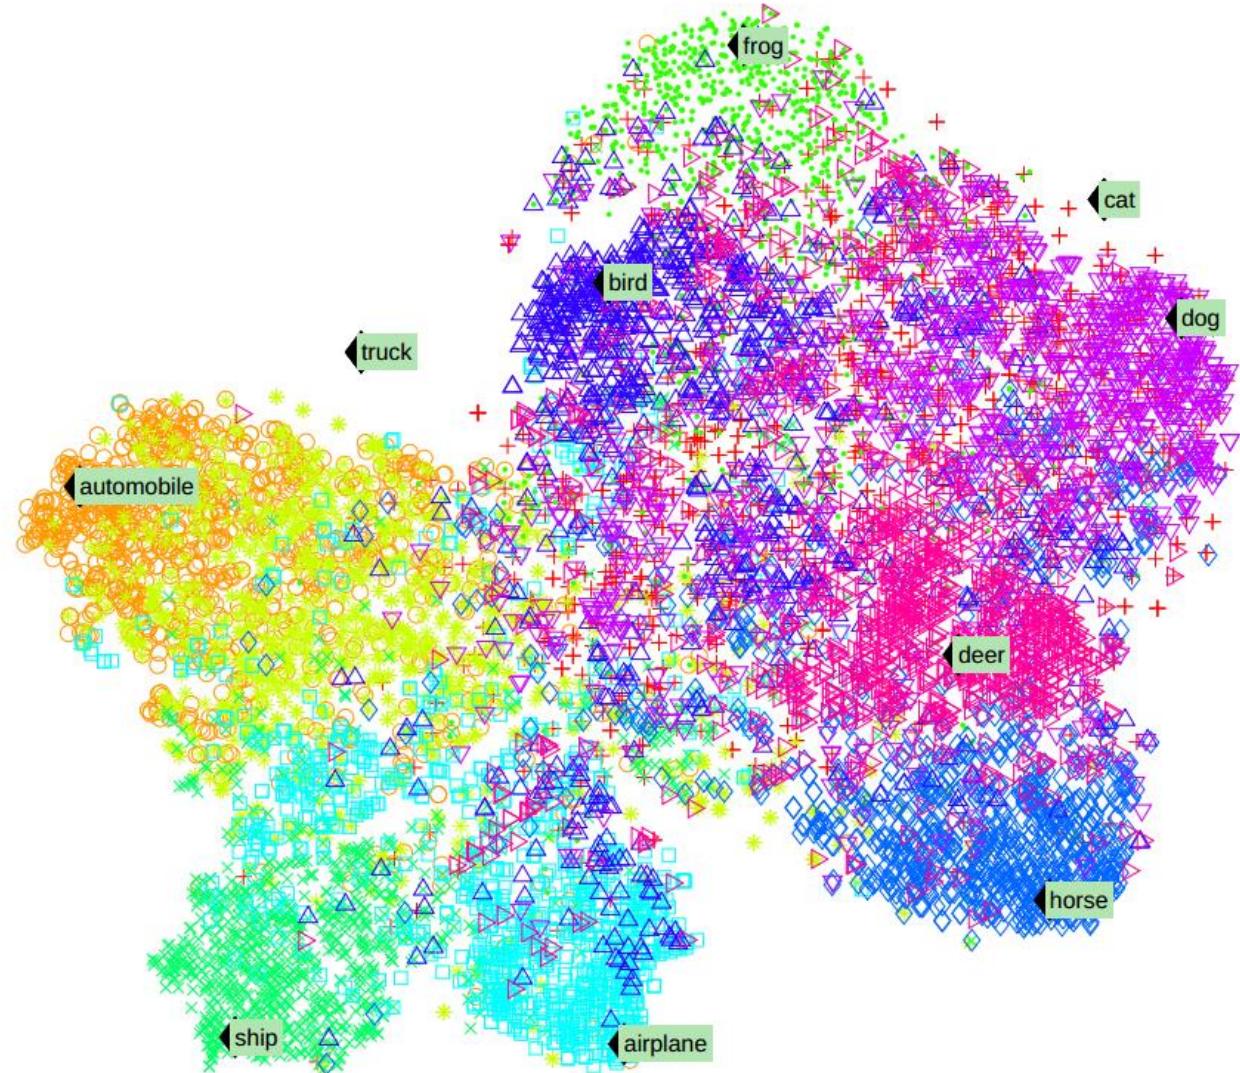


<http://stackoverflow.com/questions/40184537/what-does-embedding-do-in-tensorflow>

# Embeddings

- Indices by themselves, carry no semantic meaning
- This is where embedding comes in; more commonly known as *word vector* or *word embedding*.

- + cat
- automobile
- \* truck
- frog
- ✗ ship
- airplane
- ◊ horse
- △ bird
- ▽ dog
- ▷ deer



<http://colah.github.io/posts/2014-07-NLP-RNNs-Representations/>



See also:

**Word2vec** <https://www.tensorflow.org/tutorials/word2vec>

Vector space models (VSMs) represent (embed) words in a continuous vector space where **semantically similar words** are mapped to nearby points ('are embedded nearby each other')

# word2vec

(WATER - WET) + FIRE = FLAMES

(PARIS - FRANCE) + ITALY = ROME

(WINTER - COLD) + SUMMER = WARM

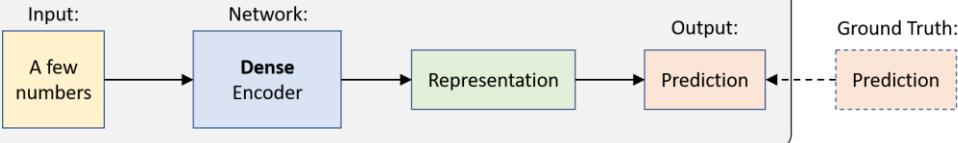
(MINOTAUR - MAZE) + DRAGON = SIMCITY

<https://ronxin.github.io/wevi/>

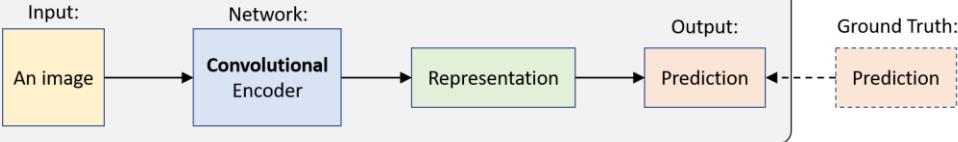
# Deep learning basics

## Supervised Learning

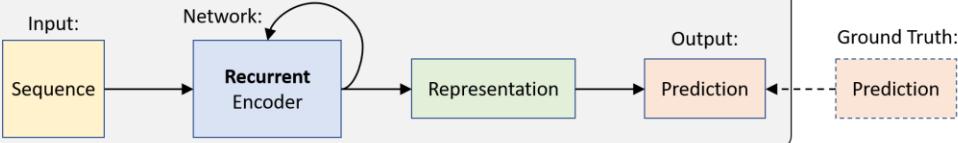
### 1. Feed Forward Neural Networks



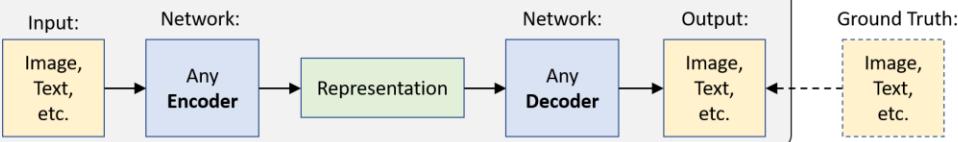
### 2. Convolutional Neural Networks



### 3. Recurrent Neural Networks

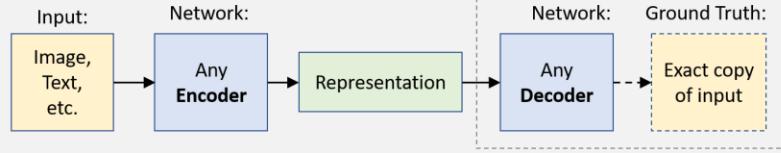


### 4. Encoder-Decoder Architectures

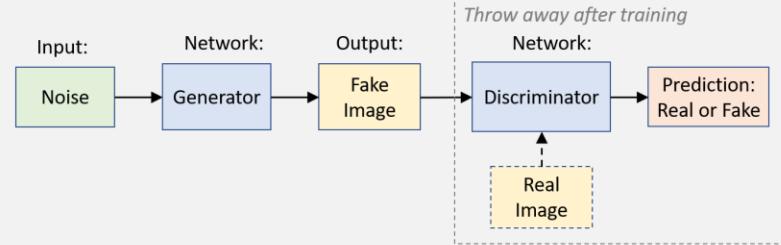


## Unsupervised Learning

### 5. Autoencoder

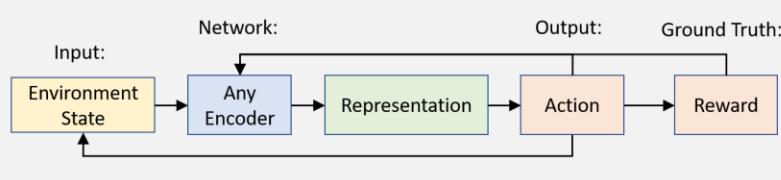


### 6. Generative Adversarial Networks



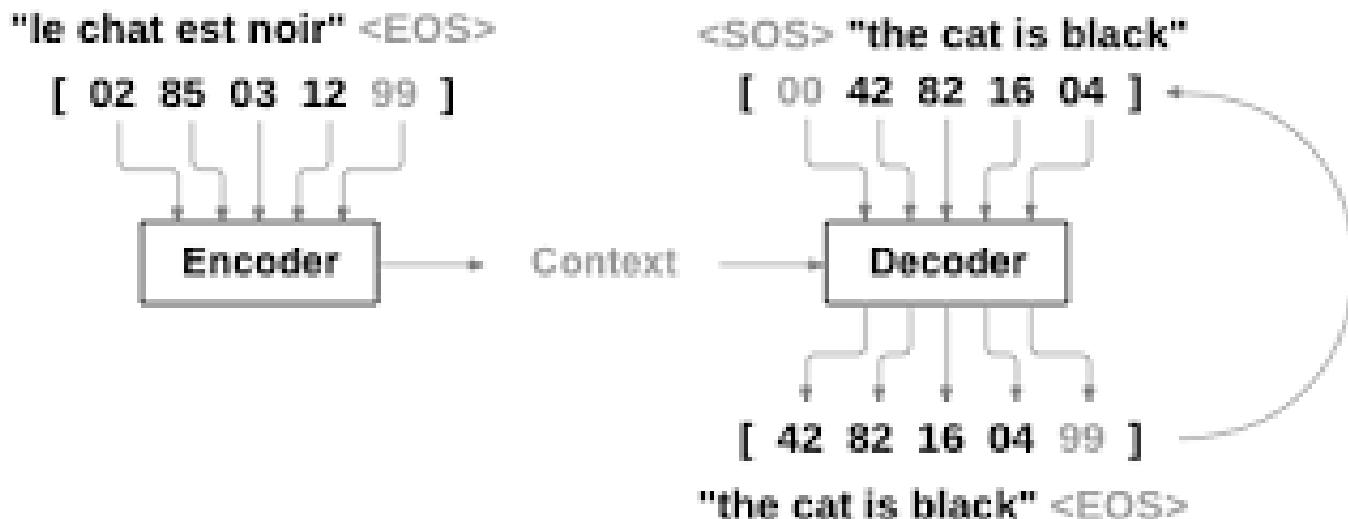
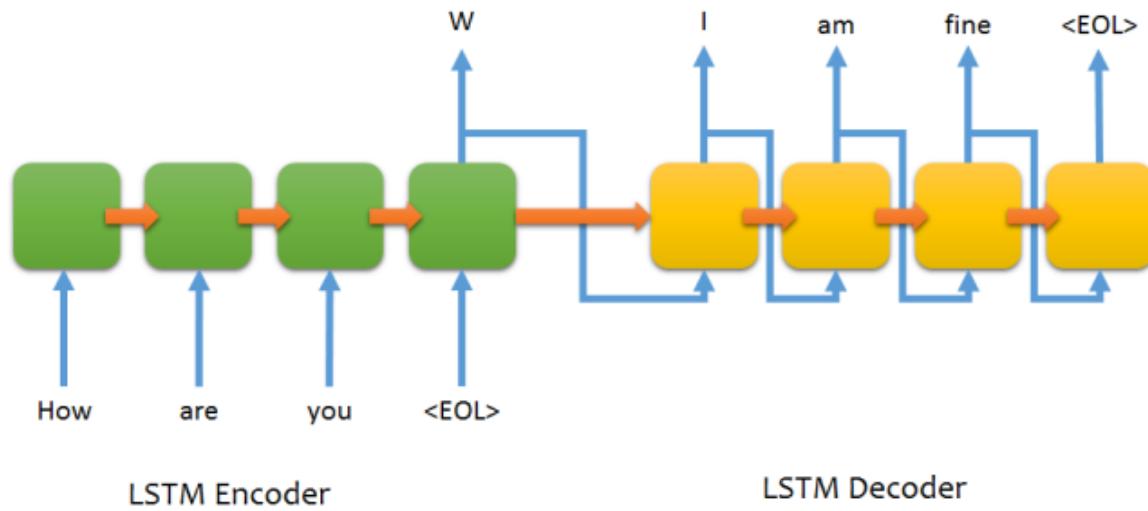
## Reinforcement Learning

### 7. Networks for Learning Actions, Values, and Policies

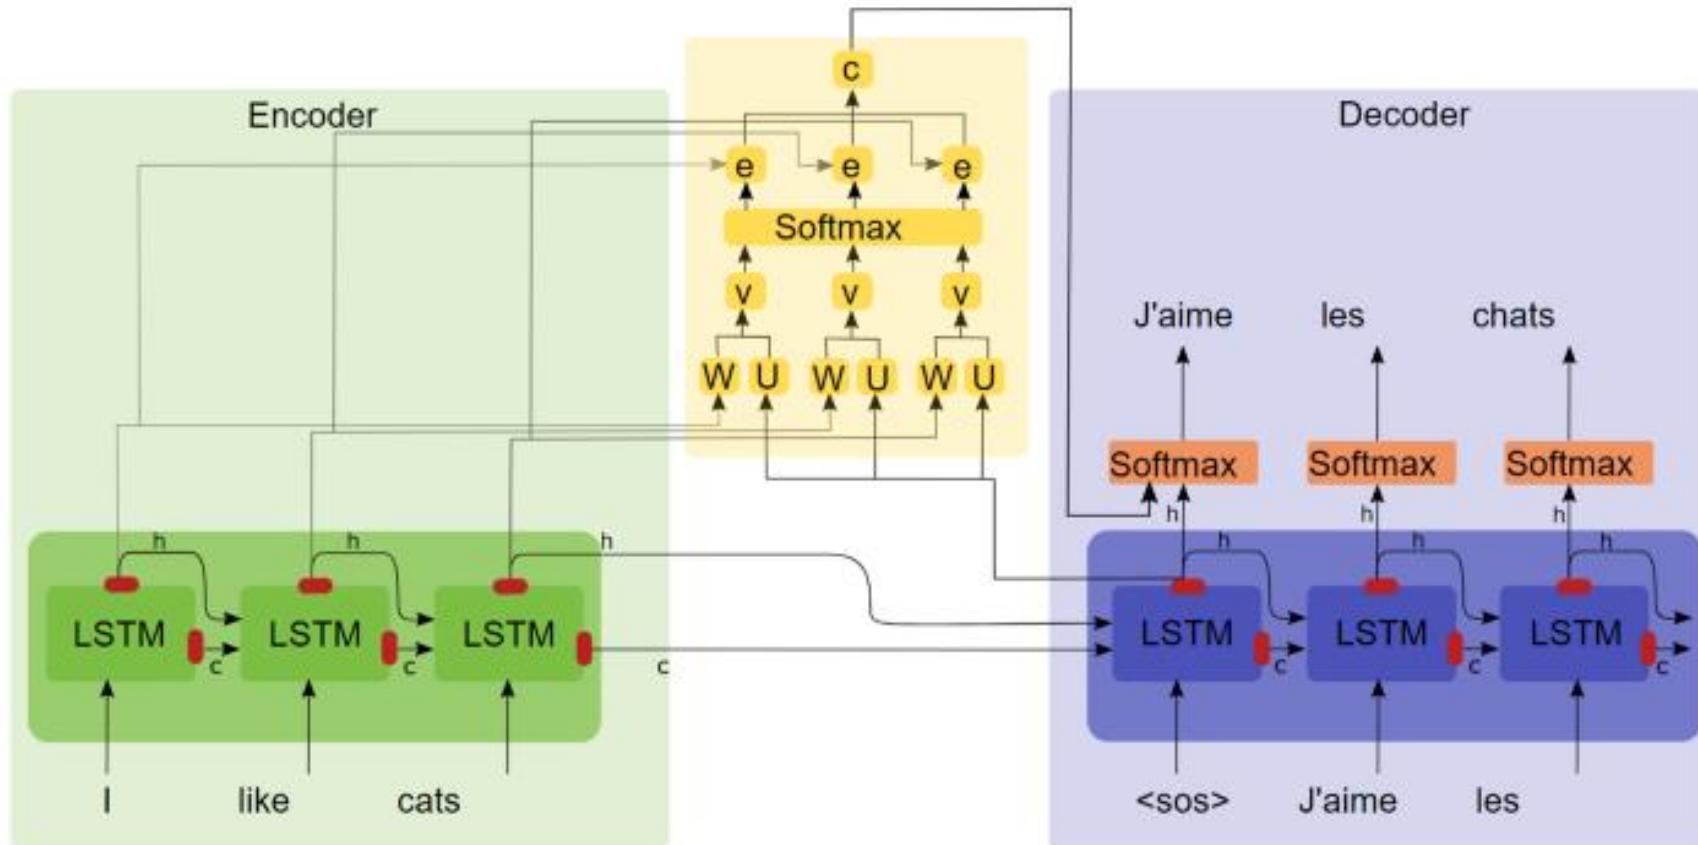


[https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial\\_deep\\_learning\\_basics/deep\\_learning\\_basics.ipynb](https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial_deep_learning_basics/deep_learning_basics.ipynb)

# Machine translation, chatbots,...

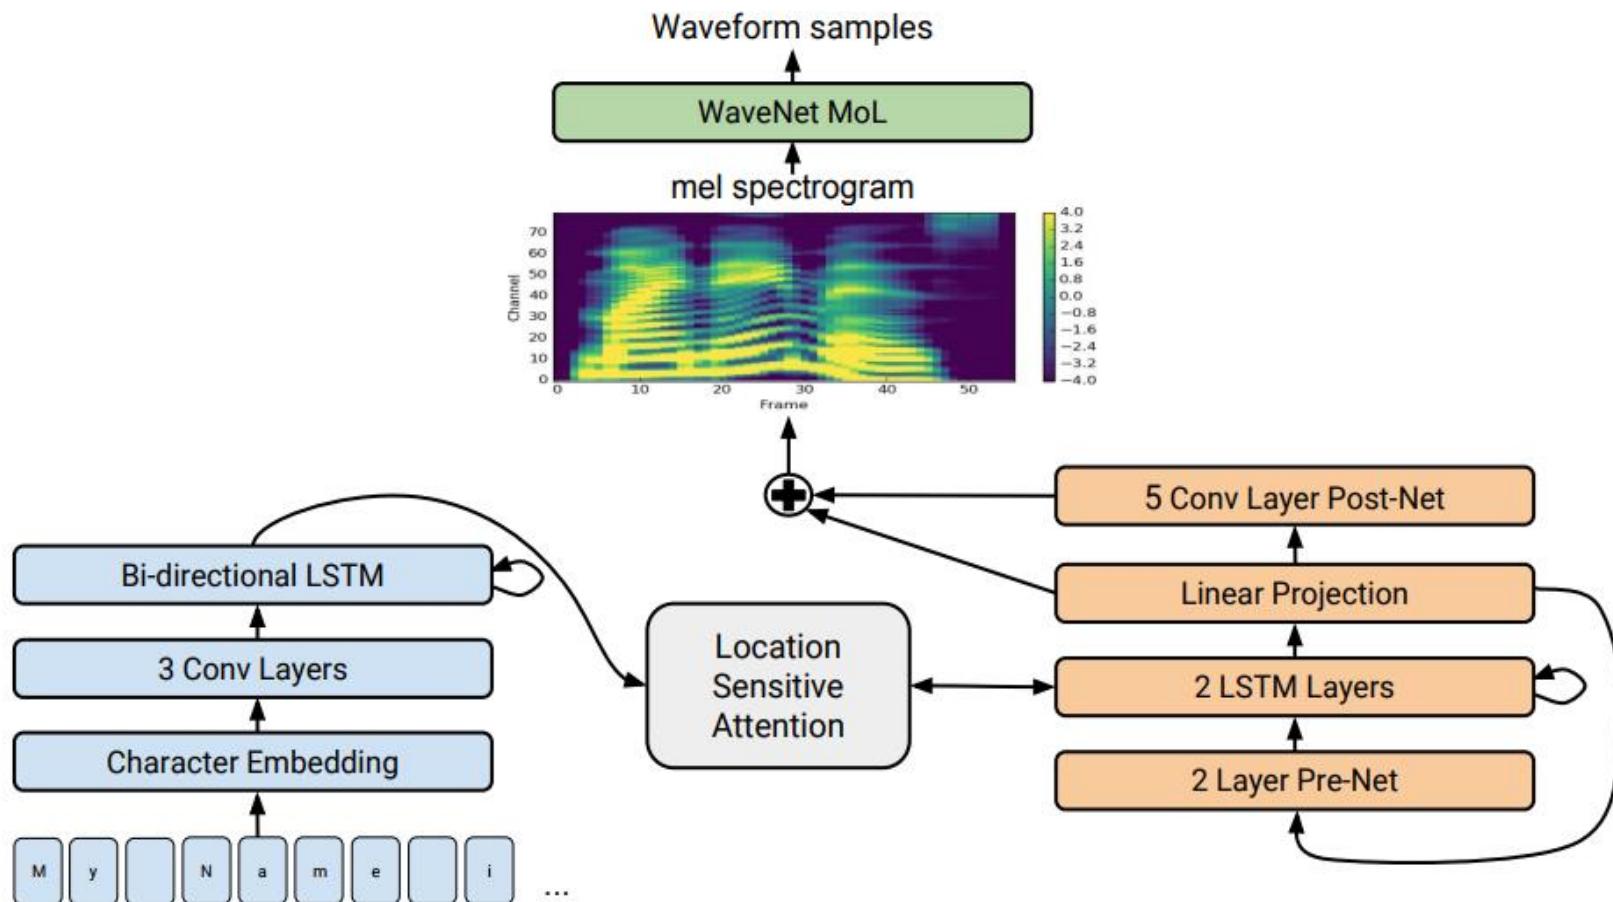


# ATTENTION



<https://towardsdatascience.com/light-on-math-ml-attention-with-keras-dc8dbc1fad39>

# END-TO-END Text-to-Speech

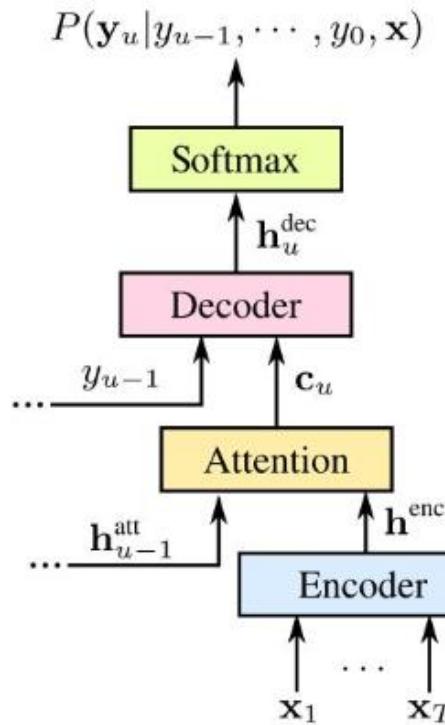


[Tacotron 2: Generating Human-like Speech from Text](#)

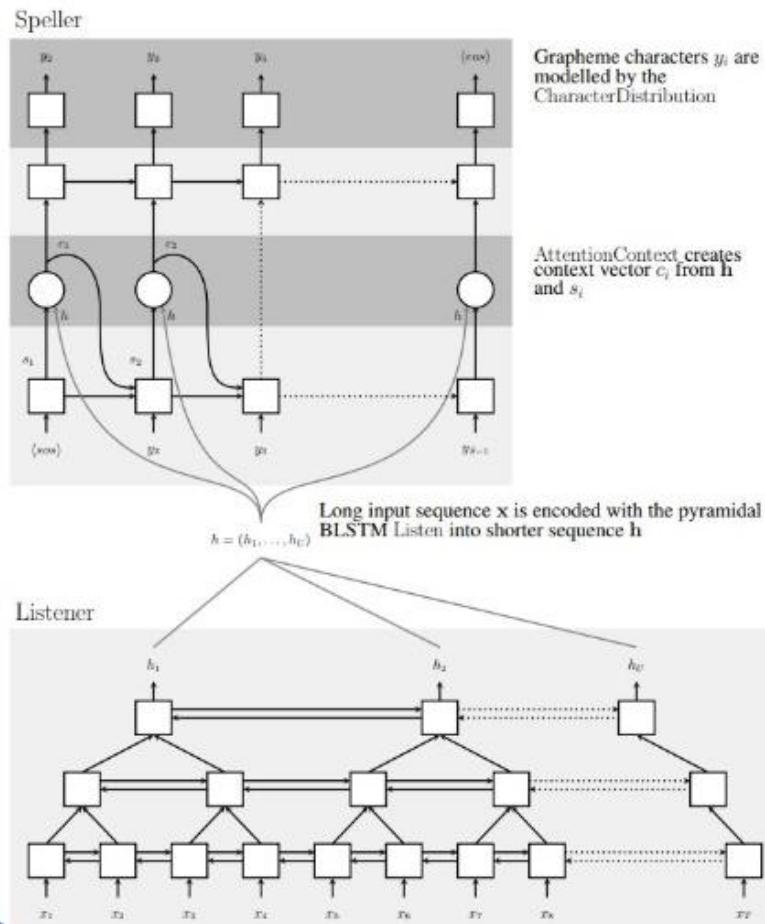
<https://ai.googleblog.com/2017/12/tacotron-2-generating-human-like-speech.html>

# END-TO-END Automatic Speech Recognition

## Attention-based Models



Reproduced from [Chan et al., 2015]



Google

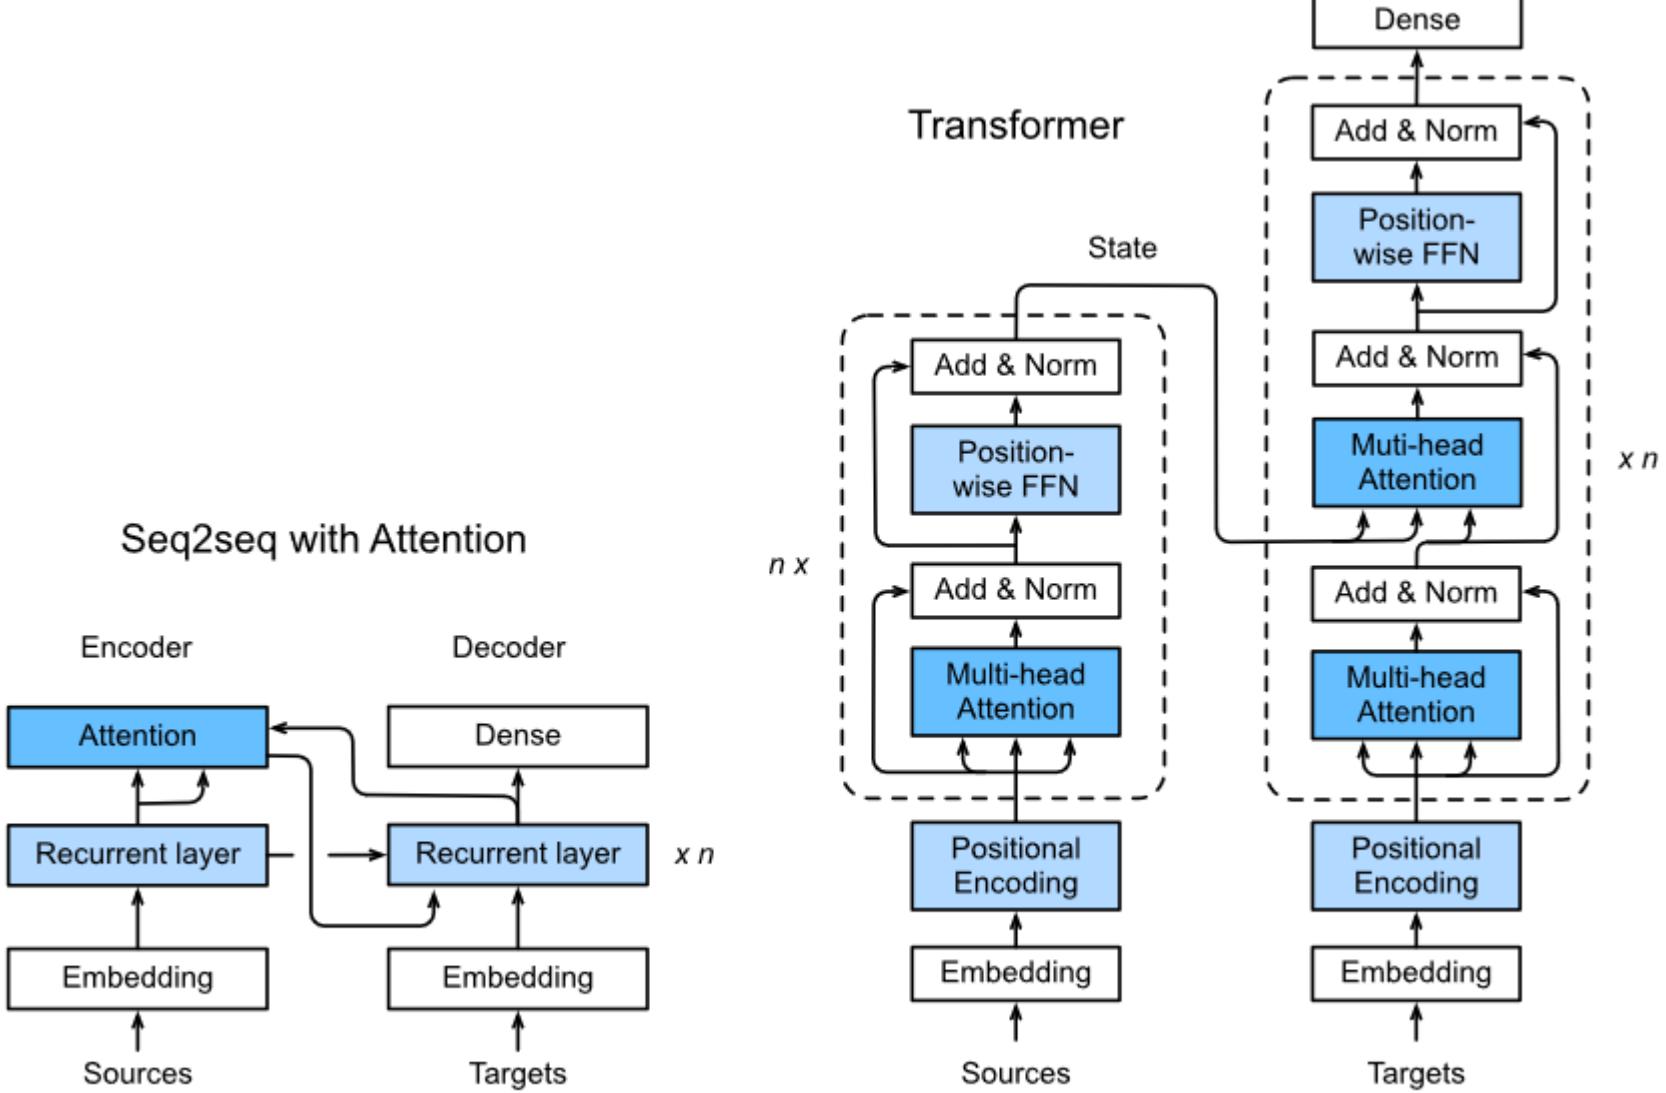
[http://isclsp2018.org/images/T4\\_Towards%20end-to-end%20speech%20recognition.pdf](http://isclsp2018.org/images/T4_Towards%20end-to-end%20speech%20recognition.pdf)

# Attention is all you need

Md Mehrab Tanjim

<http://jalammar.github.io/illustrated-transformer/>

<https://nlp.stanford.edu/seminar/details/lkaiser.pdf>



[https://www.d2l.ai/\\_images/transformer.svg](https://www.d2l.ai/_images/transformer.svg)

# References (I)

<https://theneuralspective.com/2016/10/04/05-recurrent-neural-networks-rnn-part-1-basic-rnn-char-rnn/>

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

<http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-2-implementing-a-language-model-rnn-with-python-numpy-and-theano/>

<https://github.com/suriyadeepan/rnn-from-scratch>

<http://stats.stackexchange.com/questions/241985/understanding-lstm-units-vs-cells>

<http://monik.in/a-noobs-guide-to-implementing-rnn-lstm-using-tensorflow/>

<https://medium.com/@erikhallstrm/hello-world-rnn-83cd7105b767>

<http://suriyadeepan.github.io/2017-01-07-unfolding-rnn/>

[http://archive.eetindia.co.in/www.eetindia.co.in/VIDEO\\_DETAILS\\_700001601.HTM](http://archive.eetindia.co.in/www.eetindia.co.in/VIDEO_DETAILS_700001601.HTM)

# References (II)

<http://r2rt.com/styles-of-truncated-backpropagation.html>

<http://akkikiki.github.io/assets/LSTM+and+GRU.html>

[http://campuspress.yale.edu/yw355/deep\\_learning/](http://campuspress.yale.edu/yw355/deep_learning/)

<http://datascience.stackexchange.com/questions/12964/what-is-the-meaning-of-the-number-of-units-in-the-lstm-cell>

<http://stackoverflow.com/questions/37901047/what-is-num-units-in-tensorflow-basiclstmcell>

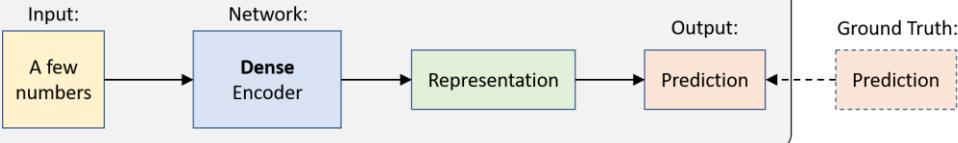
## [Transformer] Attention Is All You Need | AISC Foundational

[https://www.youtube.com/watch?v=S0KakHcj\\_rs](https://www.youtube.com/watch?v=S0KakHcj_rs)

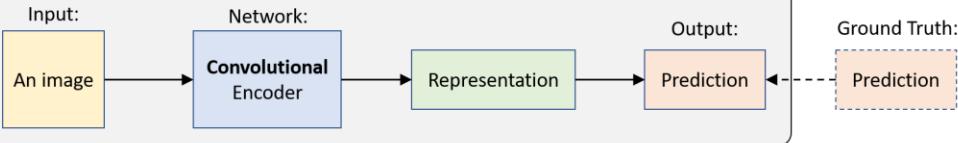
# Deep learning basics

## Supervised Learning

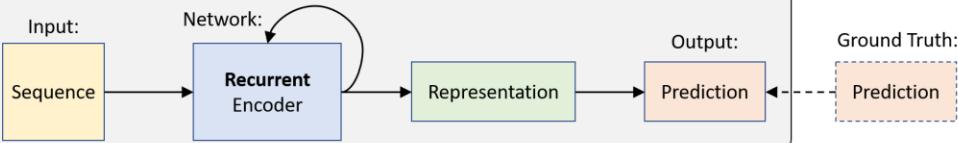
### 1. Feed Forward Neural Networks



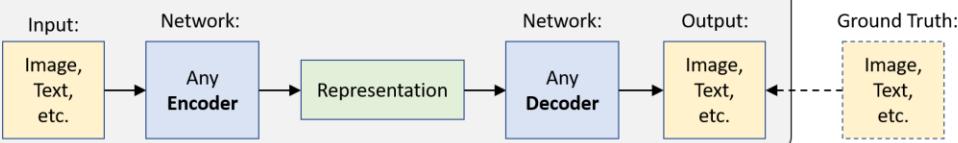
### 2. Convolutional Neural Networks



### 3. Recurrent Neural Networks

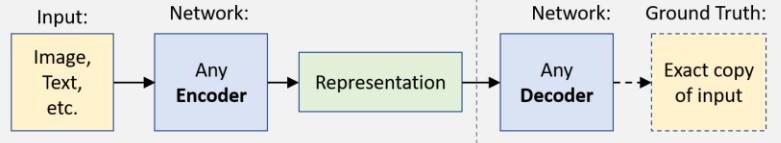


### 4. Encoder-Decoder Architectures

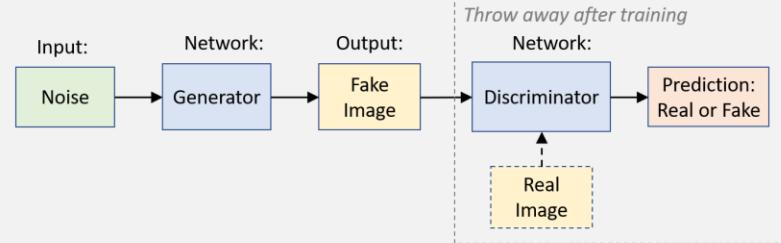


## Unsupervised Learning

### 5. Autoencoder

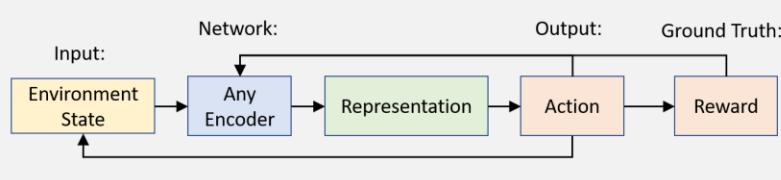


### 6. Generative Adversarial Networks



## Reinforcement Learning

### 7. Networks for Learning Actions, Values, and Policies

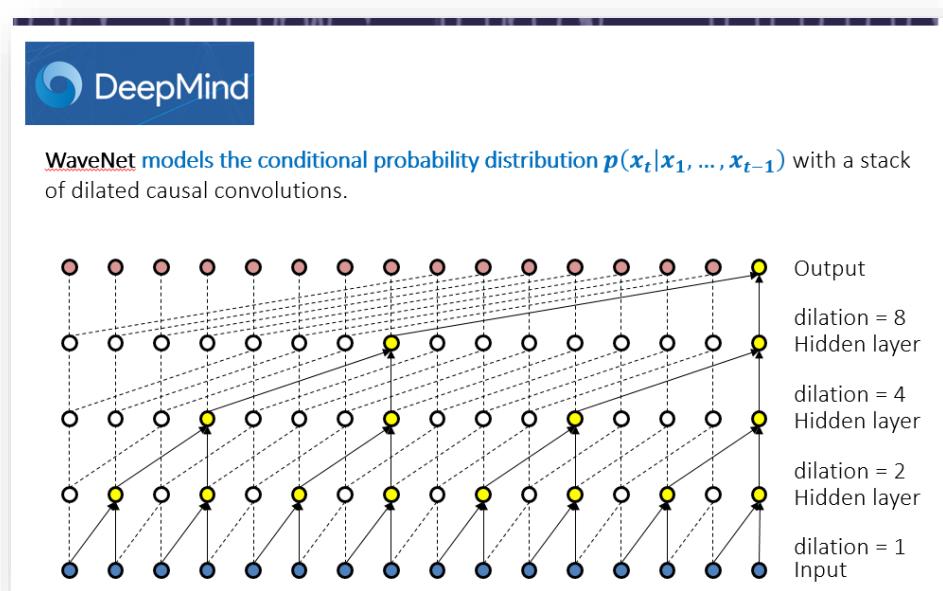


[https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial\\_deep\\_learning\\_basics/deep\\_learning\\_basics.ipynb](https://github.com/lexfridman/mit-deep-learning/blob/master/tutorial_deep_learning_basics/deep_learning_basics.ipynb)

... there is much more to explore!

Remember that RNN are GENERATIVE models !!!

- Generative Adversarial Networks : GANs
- WaveNet ...



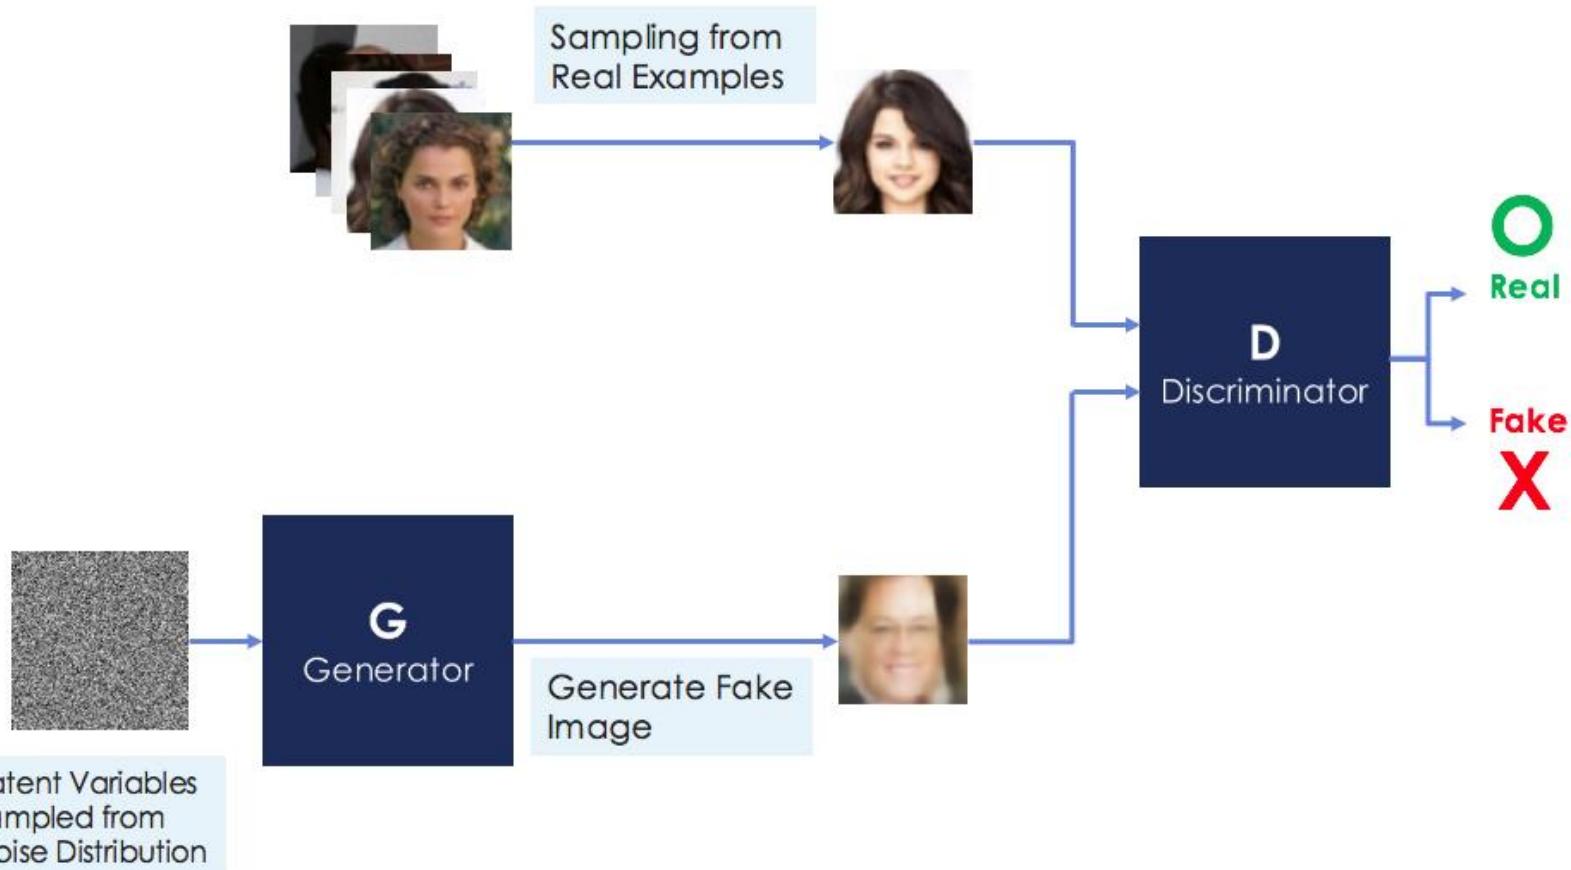
# Generative Adversarial Networks

## GANs

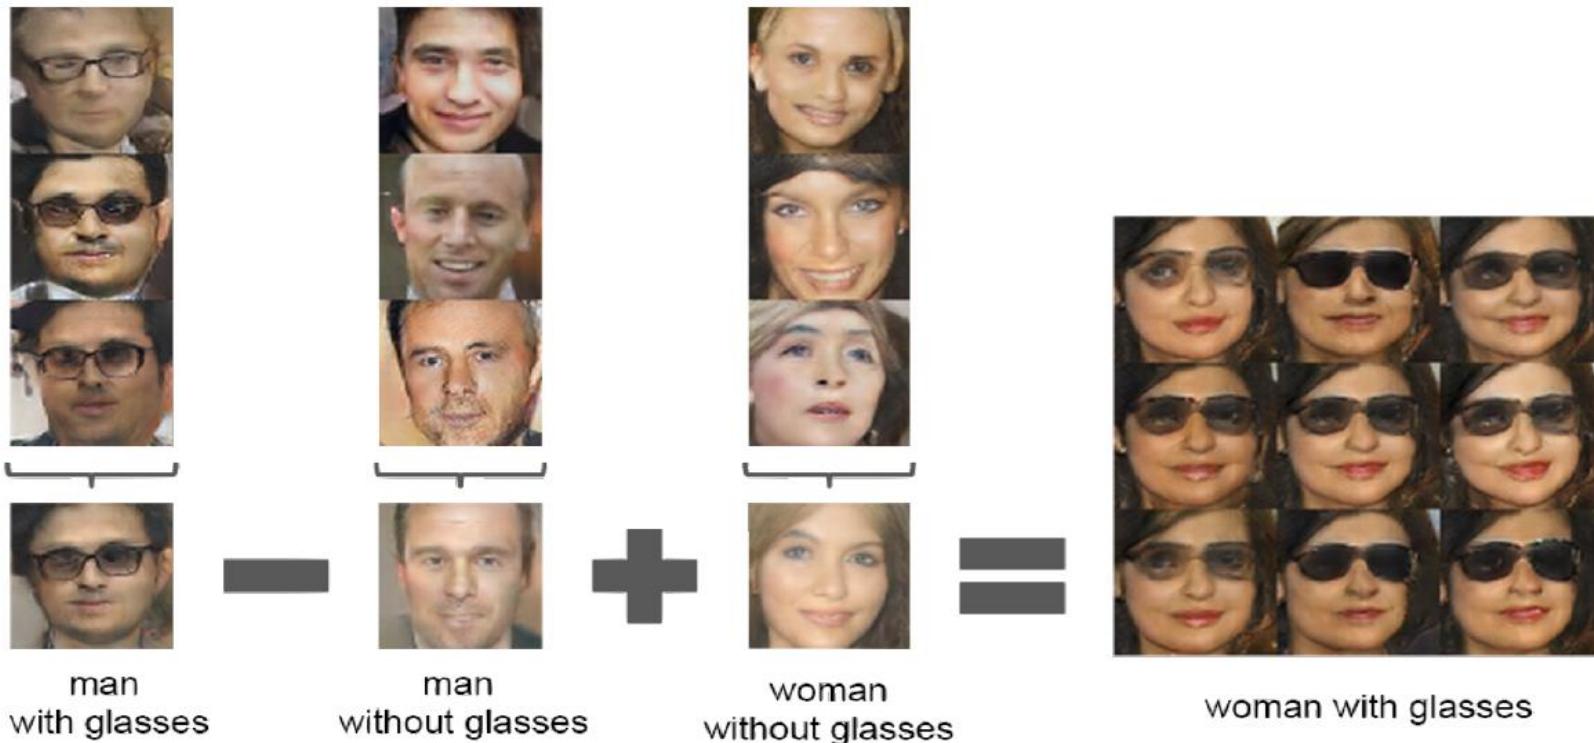
<https://youtu.be/2edOMMREazo>

<https://towardsdatascience.com/notes-on-artificial-intelligence-ai-machine-learning-ml-and-deep-learning-dl-for-56e51a2071c2>

# Generative Adversarial Networks(GAN)



# Latent vectors capture interesting patterns...



Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks." arXiv:1511.06434 (2015).

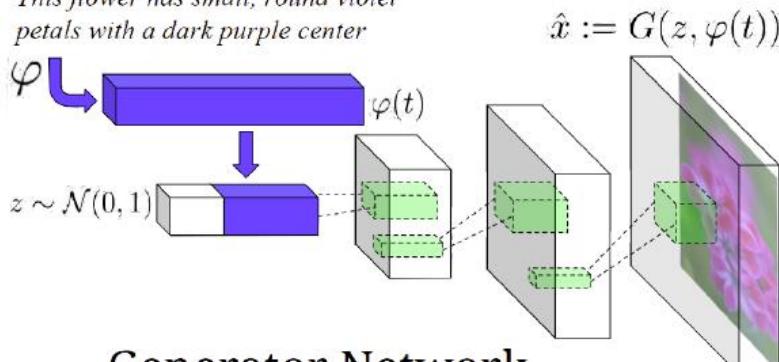


<https://youtu.be/2edOMMREazo>

<https://towardsdatascience.com/notes-on-artificial-intelligence-ai-machine-learning-ml-and-deep-learning-dl-for-56e51a2071c2>

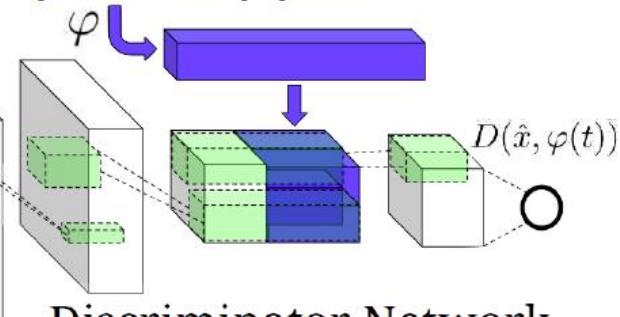
# Text-to-Image Synthesis

*This flower has small, round violet petals with a dark purple center*



Generator Network

*This flower has small, round violet petals with a dark purple center*



Discriminator Network

Figure 2 in the original paper.

Positive Example:

Real Image, Right Text

Negative Examples:

Real Image, Wrong Text

Fake Image, Right Text

Reed, S., Akata, Z., Yan, X., Logeswaran, L., Schiele, B., & Lee, H. "Generative adversarial text to image synthesis". ICML (2016).

- Image (Audio) Super Resolution
- Image-to-Image
- Visually-Aware Fashion Recommendation and DesignnWith GANs (Kang et al., 2017)
- Fraud / anomaly detection ,...

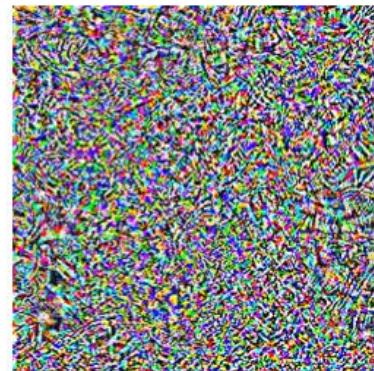
# Some refs on GANs

- <https://tryolabs.com/blog/2016/12/06/major-advancements-deep-learning-2016/>
- <https://blog.waya.ai/introduction-to-gans-a-boxing-match-b-w-neural-nets-b4e5319cc935#.6l7zh8u50>
- [https://en.wikipedia.org/wiki/Generative\\_adversarial\\_networks](https://en.wikipedia.org/wiki/Generative_adversarial_networks)
- <http://blog.aylien.com/introduction-generative-adversarial-networks-code-tensorflow/>
- <https://github.com/soumith/ganhacks>

“pig”



$$+ 0.005 \times$$



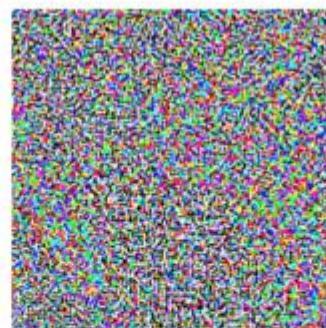
“airliner”



=



$$+ .007 \times$$



=

$x$   
“panda”  
57.7% confidence

$\text{sign}(\nabla_x J(\theta, x, y))$   
“nematode”  
8.2% confidence

$x +$   
 $\epsilon \text{sign}(\nabla_x J(\theta, x, y))$   
“gibbon”  
99.3 % confidence

<https://buzzrobot.com/4-ways-to-easily-fool-your-deep-neural-net-dca49463bd0>

## ACM NEWS

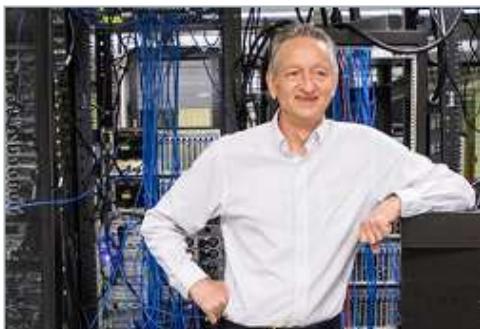
# Artificial Intelligence Pioneer Says We Need to Start Over

By Axios

September 18, 2017

[Comments](#)

VIEW AS: SHARE:



In 1986, Geoffrey Hinton co-authored a paper that, three decades later, is central to the explosion of artificial intelligence (AI). But Hinton says his breakthrough method should be dispensed with, and a new path to AI found.

Speaking with Axios on the sidelines of an AI conference in Toronto on Wednesday, Hinton, a professor emeritus at the University of Toronto and a Google researcher, said he is now "deeply suspicious" of [back-propagation](#), the workhorse method

**SIGN IN** for Full Access

User Name

Password

» [Forgot Password?](#)

» [Create an ACM Web Account](#)

**SIGN IN**

**MORE NEWS & OPINIONS**

[Researchers Combine Technologies for New Digital](#)



# Deep Learning est mort. Vive Differentiable Programming!

**Yann LeCun**, Director of Facebook AI Research and the inventor of convolutional neural networks, took it.

<https://www.facebook.com/yann.lecun/posts/10155003011462143>

## Age of AI Talk: ``Deep Learning est Mort! Vive Differentiable Programming”

Differentiable Programming Framework for Deep Learning and Machine Intelligence.



Lisha li [Follow](#)

Feb 13, 2018 · 10 min read

<https://techburst.io/deep-learning-est-mort-vive-differentiable-programming-5060d3c55074>

<https://medium.com/amplify-partners/age-of-ai-talk-deep-learning-est-morte-vive-differentiable-programming-6b1a1c9800d8>

# Demystifying Differentiable Programming: Shift/Reset the Penultimate Backpropagator

FEI WANG, Purdue University, USA

XILUN WU, Purdue University, USA

GREGORY ESSERTEL, Purdue University, USA

JAMES DECKER, Purdue University, USA

TIARK ROMPF, Purdue University, USA

Deep learning has seen tremendous success over the past decade in computer vision, machine translation, and gameplay. This success rests in crucial ways on *gradient-descent optimization* and the ability to “learn” parameters of a neural network by backpropagating observed errors. However, neural network architectures are growing increasingly sophisticated and diverse, which motivates an emerging quest for even more general forms of *differentiable programming*, where arbitrary parameterized computations can be trained by gradient descent. In this paper, we take a fresh look at automatic differentiation (AD) techniques, and especially aim to demystify the *reverse-mode* form of AD that generalizes backpropagation in neural networks.

<https://arxiv.org/abs/1803.10228>

# AGI: The forgotten science

Real AI – Human-level learning and understanding

Artificial General Intelligence (AGI)	Conventional AI
Focus on <b>acquiring</b> knowledge and skills	Focus on <b>having</b> knowledge and skills
Acquisition via learning	Acquisition via programming
General ability, using abstraction and context	Domain specific, rule-based and concrete
Ongoing cumulative, adaptive, grounded, self-directed learning	Relatively fixed abilities. Externally initiated improvements

# Cognitive Computing vs. AI

COGNITIVE COMPUTING		ARTIFICIAL INTELLIGENCE
Machine learning, natural language processing, neural networks, deep learning, sentiment analysis	TECHNOLOGIES	Machine learning, natural language processing, neural networks, deep learning
Simulate human thought processes to assist humans in finding solutions to complex problems	CAPABILITIES	Find patterns in big data to learn and either reveal hidden information or deliver solutions to complex problems
Augment human capabilities	PURPOSE	Automate processes
Customer service, healthcare, industrial sector	INDUSTRIES	Finance, security, healthcare, retail, manufacturing, government



ILLUSTRATION: MAPLUM/GETTY IMAGES

©2019 TECHTARGET. ALL RIGHTS RESERVED. TechTarget

... and more pragmatic ML things...



Towards Data Science

Follow

HOME DATA SCIENCE MACHINE LEARNING PROGRAMMING VISUALIZATION AI JOURNALISM PICKS | CONTRIBUTE

# AutoKeras: The Killer of Google's AutoML



George Seif [Follow](#)  
Jul 31, 2018 · 4 min read

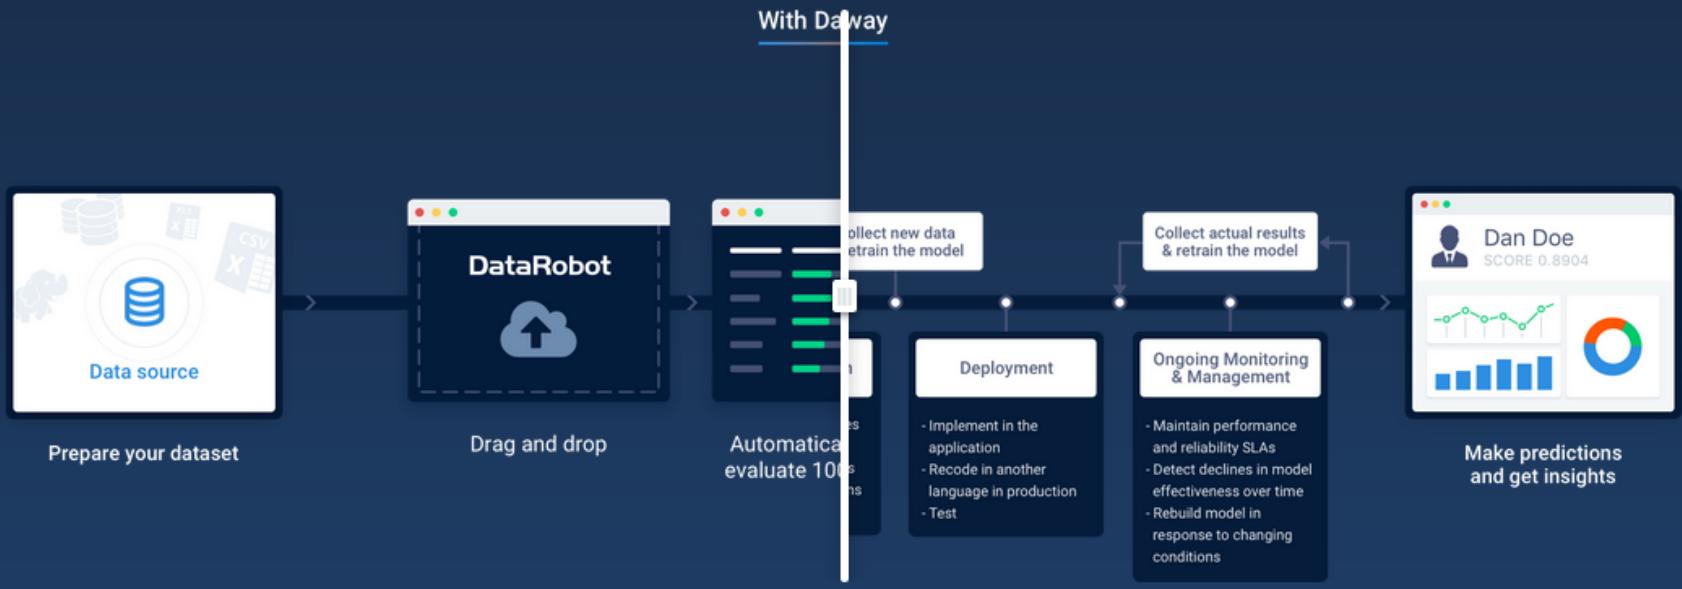


<https://towardsdatascience.com/autokeras-the-killer-of-googles-automl-9e84c552a319>

<https://www.datarobot.com/>

## DataRobot transforms model building

Keeping up with the ever-growing ecosystem of algorithms has never been this easy



# Ethics?

*Engineers at Amazon, Google, and Microsoft are demanding the companies put ethics before profits*

Would You Quit a Tech Project Over Ethical Concerns?  
By KATHY PRETZ 2 November 2018

<http://theinstitute.ieee.org/ieee-roundup/career-and-education/career-guidance/would-you-quit-a-tech-project-over-ethical-concerns>

# Why Schools Are Getting More Serious About Teaching Engineering Students About Ethics?

*Harvard, MIT, and Stanford say **it's more important than ever to recognize the ramifications of AI technology***

By KATHY PRETZ 5 September 2018

<http://theinstitute.ieee.org/ieee-roundup/blogs/blog/why-schools-are-getting-more-serious-about-teaching-engineering-students-about-ethics>

*Engineers at Amazon, Google, and Microsoft are demanding the companies put ethics before profits*

Would You Quit a Tech Project Over Ethical Concerns?  
By KATHY PRETZ 2 November 2018

<http://theinstitute.ieee.org/ieee-roundup/career-and-education/career-guidance/would-you-quit-a-tech-project-over-ethical-concerns>

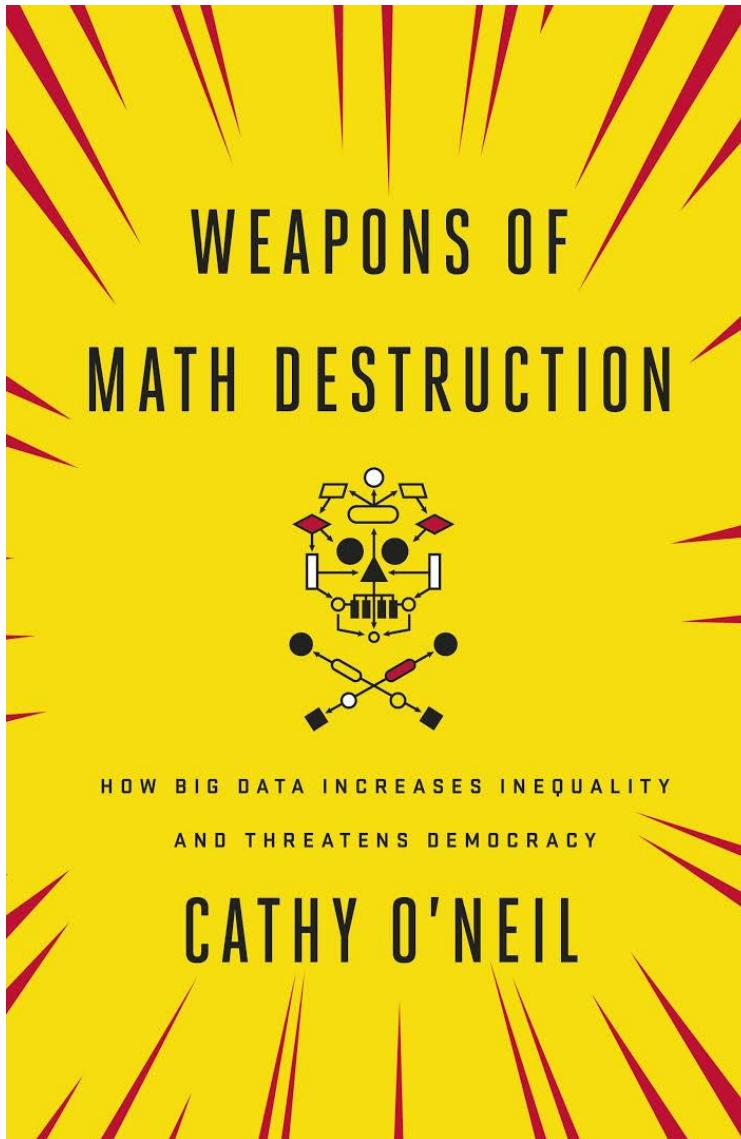
- **Ethics for Machines** : *what type of machines is acceptable for us to make (privacy, use of and relationship with technology,...).*
- **Ethics of Machines** : *how we want machines to be behaving when they are taking decisions autonomously.*

Machine ethics podcast episode with Derek Leben

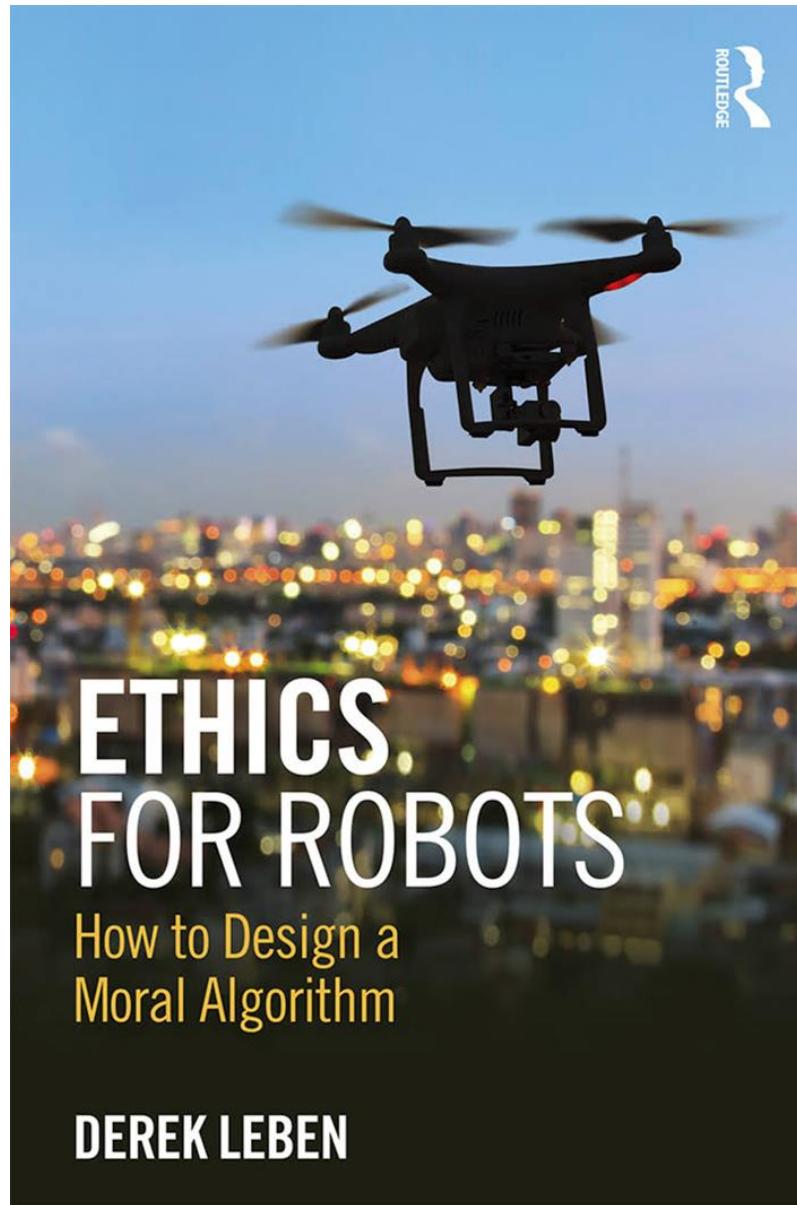
Associate Professor of Philosophy at the University of Pittsburgh, Johnstown

<https://www.youtube.com/watch?v=OjLA6lZEPW4>

# Ethics for Machines



# Ethics of Machines



# ....or just BE GOOD!!!

Inicio Momentos Buscar en Twitter

 **Andrew Ng**   
@AndrewYNg 

I'm glad DeepNude is dead. As a person and as a father, I thought this was one of the most disgusting applications of AI. To the AI Community: You have superpowers, and what you build matters. Please use your powers on worthy projects that move the world forward.

11:06 - 28 jun. 2019

1.984 Retweets 8.092 Me gusta 

164 2,0K 8,1K

 **Muriz** @MurgioMurmani · 28 jun.  
En respuesta a [@AndrewYNg](#)  
I feel like this is just the beginning of AI misuse

3 3 73

# That's all !

I hope this seminar will be useful to you

...we keep waiting for your amazing deep-applications..

## THANK YOU!

[luisalfonso.hernandez@upm.es](mailto:luisalfonso.hernandez@upm.es)